pyAmpliCol at one loop

The recursive one-loop engine against MadLoop, Recola, OpenLoops and NJet

milestone M-LIT-a0 (2026-11-06) · engine `ca602436` on `loop` · AMD EPYC 9754, shared host · finite part, binary64, one pinned thread, $B = 128$

×1/16 faster — ×1/4 — ×1 — ×4 — ×16 slower
run time ÷ fastest provider

Gate pass ×0.25 · fail ×1.4 · open ×1.0 · <5 ms: not gated ×1.4 · time ÷ limit
Stages inputs · numerator · scatter · NinjaRS · masters · glue
missing (why) · no colour: generation, LC · † within estimate · ¶ reference's own (DD) · ‡ user decision
ML MadLoop · Rc Recola · OL OpenLoops

Table 1: **State of the one-loop engine.** How to read it: after Figure 1.

	Process · validation	Generation	Full colour · per point, batched	Leading colour (LC)
pyAmpliCol values: the most recent measured configuration — `loop` · $2 \to 2$ to $2 \to 5$: `ca602436`, native per-point inputs, G1 2026-11-06				
	legs · corrections · loop families; largest deviation of the finite part per reference	generate + plan export, peak memory, plan size; provider generation on this host	median per point at $B = 128$, the gate batch; provider times; gate G1 (time ÷ limit); stage shares	batched, same configuration; LC/FC; no provider ratio
$2 \to 2$	$gg \to gg$ 4 legs · QCD · 24 families (LC 12) Rc·OL·ML ≤ 1.3e−12 16 pts · median 1.1e−14 · †	**5.70 s** ×40 0.57 GiB peak · plan 0.0240 MB **Rc 142 ms** · ML 9.70 s	**1.29 ms** ×0.64 $B = 128$ **ML 2.02 ms** · OL 3.20 ms · Rc 10.0 ms no gate (2→2) numerator 51 %	**0.596 ms** LC/FC 0.46 $B = 128$ · FC 1.29 ms
$2 \to 3$	$u\bar{u} \to d\bar{d}g$ 5 legs · QCD · 42 families (LC 27) ML·OL ≤ 7.9e−13 · Rc 3.3e−10 16 pts · median 4.8e−14 · ¶	**6.73 s** ×17 0.58 GiB peak · plan 0.0300 MB **Rc 406 ms** · ML 9.87 s	**1.72 ms** ×1.1 $B = 128$ **OL 1.52 ms** · ML 3.17 ms · Rc 3.18 ms <5 ms: not gated ×0.87 numerator 38 %	**1.12 ms** LC/FC 0.65 $B = 128$ · FC 1.72 ms
	$u\bar{u} \to u\bar{u}g$ 5 legs · QCD · 68 families (LC 42) OL·ML ≤ 1.2e−11 · Rc 7.2e−11 16 pts · median 1.4e−13 · ¶	**7.07 s** ×15 0.58 GiB peak · plan 0.0360 MB **Rc 462 ms** · ML 10.9 s	**2.64 ms** ×0.90 $B = 128$ **OL 2.92 ms** · Rc 5.55 ms · ML 8.59 ms <5 ms: not gated ×0.69 numerator 42 %	**1.59 ms** LC/FC 0.60 $B = 128$ · FC 2.64 ms
	$u\bar{u} \to e^-e^+g$ 5 legs · QCD · 11 families (LC 7) ML·Rc ≤ 7.5e−14 16 pts · median 6.8e−15	**10.1 s** ×14 0.61 GiB peak · plan 0.0530 MB **Rc 701 ms** · ML 14.0 s	**0.788 ms** ×0.83 $B = 128$ **ML 0.948 ms** · Rc 1.08 ms · OL no library <5 ms: not gated ×0.83 numerator 36 %	**0.446 ms** LC/FC 0.57 $B = 128$ · FC 0.788 ms
	$u\bar{u} \to ggg$ 5 legs · QCD · 90 families (LC 62) ML·Rc ≤ 1.5e−12 · OL 8.4e−11 16 pts · median 4.9e−14 · ¶	**6.62 s** ×10 0.58 GiB peak · plan 0.0490 MB **Rc 634 ms** · ML 17.6 s	**4.34 ms** ×0.35 $B = 128$ **OL 12.4 ms** · Rc 25.3 ms · ML 28.2 ms <5 ms: not gated ×0.27 numerator 53 %	**2.38 ms** LC/FC 0.55 $B = 128$ · FC 4.34 ms
	$gg \to ggg$ 5 legs · QCD · 154 families (LC 77) OL·Rc·ML ≤ 3.7e−7 16 pts · median 5.8e−13 · †	**8.10 s** ×1.2 0.59 GiB peak · plan 0.112 MB **Rc 6.63 s** · ML 47.0 s	**20.1 ms** ×0.21 $B = 128$ **OL 97.0 ms** · ML 247 ms · Rc 351 ms pass ×0.16 numerator 50 %	**7.59 ms** LC/FC 0.38 $B = 128$ · FC 20.1 ms
$2 \to 4$	$u\bar{u} \to e^-e^+gg$ 6 legs · QCD · 69 families (LC 33) Rc·ML ≤ 2.2e−10 16 pts · median 6.0e−14 · †	**9.53 s** ×11 0.62 GiB peak · plan 0.0680 MB **Rc 833 ms** · ML 24.8 s	**4.21 ms** ×0.41 $B = 128$ **Rc 10.2 ms** · ML 27.3 ms · OL no library <5 ms: not gated ×0.32 numerator 54 %	**1.71 ms** LC/FC 0.41 $B = 128$ · FC 4.21 ms
	$gg \to t\bar{t}gg$ 6 legs · QCD · 1 086 families (LC 375) OL·Rc·ML ≤ 8.8e−9 16 pts · median 1.2e−11 · †	**9.98 s** ×0.16 0.64 GiB peak · plan 0.320 MB **Rc 61.4 s** · ML 218 s	**79.8 ms** ×0.17 $B = 128$ **OL 472 ms** · Rc 1.97 s · ML 5.45 s pass ×0.13 numerator 44 %	**28.2 ms** LC/FC 0.35 $B = 128$ · FC 79.8 ms
	$u\bar{u} \to gggg$ 6 legs · QCD · 628 families (LC 375) Rc 8.0e−12 · OL 2.3e−10 16 pts · median 1.7e−12 · †	**9.86 s** ×0.30 0.65 GiB peak · plan 0.382 MB **Rc 32.8 s**	**65.3 ms** ×0.17 $B = 128$ **OL 374 ms** · Rc 908 ms · ML not exported pass ×0.13 numerator 57 %	**25.3 ms** LC/FC 0.39 $B = 128$ · FC 65.3 ms
	$gg \to gggg$ 6 legs · QCD · 1 100 families (LC 550) ML·Rc·OL ≤ 1.5e−7 8 pts · median 2.5e−11 · †	**14.0 s** ×0.0066 0.89 GiB peak · plan 1.23 MB **Rc 35.4 min** · ML 144.6 min	**307 ms** ×0.11 $B = 128$ **OL 2.84 s** · Rc 17.6 s · ML 38.7 s pass ×0.083 numerator 51 %	**126 ms** LC/FC 0.41 $B = 128$ · FC 307 ms

	Process · validation	Generation	Full colour · per point, batched	Leading colour (LC)
pyAmpliCol values: the most recent measured configuration — `loop` · $2 \to 2$ to $2 \to 5$: ca602436, native per-point inputs, G1 2026-11-06				
	legs · corrections · loop families; largest deviation of the finite part per reference	generate + plan export, peak memory, plan size; provider generation on this host	median per point at $B = 128$, the gate batch; provider times; gate G1 (time ÷ limit); stage shares	batched, same configuration; LC/FC; no provider ratio
$2 \to 5$	$u\bar{u} \to e^-e^+d\bar{d}g$ 7 legs · QCD · 242 families (LC 114) OL·ML·Rc ≤ 4.5e−11 16 pts · median 5.3e−13	**12.1 s** ×1.8 0.63 GiB peak · plan 0.0960 MB **Rc 6.73 s** · ML 58.6 s	**13.9 ms** ×0.39 **Rc 35.3 ms** · OL 51.0 ms · ML 132 ms $B = 128$ pass ×0.30 numerator 48 %	**5.75 ms** LC/FC 0.41 $B = 128$ · FC 13.9 ms
	$u\bar{u} \to e^-e^+ggg$ 7 legs · QCD · 463 families (LC 164) ML·Rc·OL ≤ 4.3e−11 16 pts · median 1.4e−12 · †	**11.1 s** ×0.83 0.64 GiB peak · plan 0.179 MB **Rc 13.4 s** · ML 146 s	**47.9 ms** ×0.19 **Rc 251 ms** · OL 346 ms · ML 1.42 s $B = 128$ pass ×0.15 numerator 63 %	**9.42 ms** LC/FC 0.20 $B = 128$ · FC 47.9 ms
	$gg \to t\bar{t}e^-e^+g$ 7 legs · QCD · 579 families (LC 164) OL·Rc·ML ≤ 1.1e−8 16 pts · median 2.7e−11 · †	**11.8 s** ×0.43 0.64 GiB peak · plan 0.176 MB **Rc 27.7 s** · ML 165 s	**51.1 ms** ×0.13 **OL 404 ms** · Rc 620 ms · ML 3.04 s $B = 128$ pass ×0.097 numerator 57 %	**10.1 ms** LC/FC 0.20 $B = 128$ · FC 51.1 ms
	$gg \to t\bar{t}e^-e^+g$ 7 legs · EW · 1 599 families OL 2.7e−10 · ML 2.8e−9 16 pts · median 4.5e−10	**276 s** ×12 9.24 GiB peak · plan 0.445 MB **Rc 23.5 s** · ML 11.9 min	**131 ms** ×0.13 **Rc 987 ms** · OL 1.40 s · ML 7.07 s $B = 128$ pass ×0.10 numerator 46 % model: the workspace's py3 G_mu copy	***not supported*** refused by design: EW LC/NLC is never supported (user 2026-10-30, Q4; typed colour-declaration-coverage)
	$u\bar{u} \to ggggg$ 7 legs · QCD · 5 258 families (LC 2 712) ML 2.1e−12 · Rc 2.9e−10 3 pts · median 7.0e−11 · ¶	**88.2 s** ×0.0086 2.76 GiB peak · plan 14.7 MB **Rc 170.3 min** · ML 12.2 h	**1.31 s** ≥×0.0072 **Rc ≤182 s** · ML 569 s · OL no library $B = 128$ pass ×0.0023 numerator 62 %	**438 ms** LC/FC 0.34 $B = 128$ · FC 1.31 s
	$gg \to ggggg$ 7 legs · QCD · 4 707 families (LC 4 707) NJet DD 1.5e−7 10 pts · median 3.4e−9 · † · top-free (particle_veto = [t]), Nf = 5 ‡ pts 4, 8 above 2.0e−8: accepted (user 2026-10-25)	**133 s** 4.39 GiB peak · plan 20.2 MB ML diagrams 17.4 h, stopped in export · Rc unfinished	**4.99 s** $B = 128$ no gate (no matched provider) numerator 52 %	**2.68 s** LC/FC 0.54 $B = 128$ · FC 4.99 s
$2 \to 6$ (M-FINAL item 1, below Figure 1): pyAmpliCol only, no provider exists at $2 \to 6$, so no ratio, no gate and no colour · engine `2f3f76d9` (`work/v2-mf-8g-3`), native `loop-native-mf-2to6-2f3f76d9-lto` (2bae17f9…), an ancestor of this milestone's pin, not re-run on it · internal checks only				
$2 \to 6$	$gg \to gggggg$ 8 legs · QCD · 47 348 families · $N_f = 5$, no top **LC** $\varepsilon^{-2} \le$ 1.0e−11 · $E \le$ 1.4e−3 (2 pts above 10^{-5}) · binary64 = DD to 7.4e−5 · rescue 11/16 · NGluon DD to 6.3e−13 · **FC** $\varepsilon^{-2} \le$ 8.4e−11 · $E \le$ 5.1e−3 (2 pts above 10^{-5}) · Catani $\varepsilon^{-1} \le$ 7.1e−5 · binary64 = DD to 1.2e−4 · rescue 11/16 Ward FC not run, Ward LC not asserted	**55.3 min** FC 74.0 GiB peak · plan 486 MB · store 25.8 min artifact 0.53 GB · load 4.21 s (1.06 GiB) **LC** 46.1 min · 63.9 GiB · artifact 0.364 GB	**201 s** 1 block of 16 $B = 16$ no provider: no gate inputs 55 % peak 40.1 GiB	**98.0 s** 1 block of 128 $B = 128$ numerator 54 % peak 27.2 GiB
	$u\bar{u} \to gggggg$ 8 legs · QCD · 47 348 families · $N_f = 5$, no top **LC** $\varepsilon^{-2} \le$ 4.6e−12 · $E \le$ 6.2e−7 · binary64 = DD to 8.4e−11 · rescue 1/16 · Ward ≤ 6.5e−12 · **FC** $\varepsilon^{-2} \le$ 1.1e−11 · $E \le$ 5.9e−6 · Catani $\varepsilon^{-1} \le$ 3.3e−8 · binary64 = DD to 1.7e−10 · rescue 2/16 · Ward ≤ 2.1e−11	**73.6 min** FC 92.9 GiB peak · plan 814 MB · store 449 s artifact 0.143 GB · load 3.85 s (1.66 GiB) **LC** 20.5 min · 47.0 GiB · artifact 0.047 GB	**62.1 s** 1 block of 128 $B = 128$ no provider: no gate numerator 77 % peak 31.7 GiB	**10.9 s** 1 block of 128 $B = 128$ numerator 60 % peak 12.6 GiB
	$u\bar{u} \to u\bar{u}gggg$ 8 legs · QCD · 40 454 families · $N_f = 5$, no top **LC** $\varepsilon^{-2} \le$ 6.7e−14 · $E \le$ 7.5e−7 · binary64 = DD to 1.3e−11 · rescue 2/16 · Ward ≤ 4.3e−12 · **FC** $\varepsilon^{-2} \le$ 2.0e−13 · $E \le$ 6.5e−6 · Catani $\varepsilon^{-1} \le$ 6.6e−11 · binary64 = DD to 1.3e−11 · rescue 4/16 · Ward ≤ 8.5e−12 (DD 1.1e−28)	**29.8 min** FC 42.5 GiB peak · plan 315 MB · store 46.4 s artifact 0.038 GB · load 1.51 s (0.68 GiB) **LC** 506 s · 18.2 GiB · artifact 0.009 GB	**19.2 s** 1 block of 128 $B = 128$ no provider: no gate numerator 75 % peak 36.0 GiB	**2.67 s** 1 block of 128 $B = 128$ numerator 55 % peak 7.38 GiB

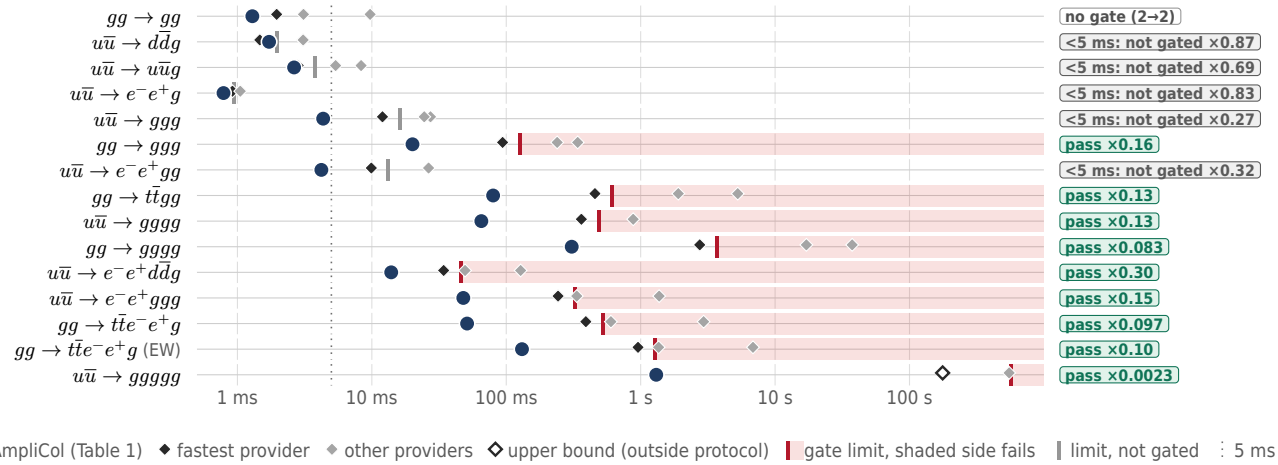


Figure 1: **Batched full colour against the gate** (finite part, $B = 128$; the rows of Table 1 with a matched provider).

Table 2: **Loop-induced processes** ($|A_{\text{loop}}|^2$; the equivalent of Table 1). How to read it: after Figure 1.

	Process · validation	Generation	$\lvert A_{\text{loop}}\rvert^2$ · per point, batched, default backend
no loop-induced package has landed at this head: every row **pending** (served by the package named in its cell) · default backend TIRust · nothing estimated			
	legs · corrections · PLAN §4.1 ids; largest deviation per reference (lit-val)	generate + plan export, peak memory; provider generation	median per point at $B = 128$; provider times (lit-ref); gate (lit-perf) or rule; backend
$2 \to 1 \cdot 1 \to 2$	$gg \to h$ 3 legs · QCD · LI-A1 validation pending (lit-val)	***pending*** served by lit-amp (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$h \to gg$ 3 legs · QCD · LI-A2 validation pending (lit-val)	***pending*** served by lit-amp (not landed at this head) providers pending (lit-ref): ML · Rc	
	$h \to aa$ 3 legs · QED · LI-A3, LI-E1 validation pending (lit-val)	***pending*** served by lit-ew (not landed at this head) providers pending (lit-ref): ML · Rc	
	$h \to za$ 3 legs · QED · LI-A4, LI-E1 validation pending (lit-val)	***pending*** served by lit-ew (not landed at this head) providers pending (lit-ref): ML · Rc	
	$aa \to h$ 3 legs · QED · LI-E2 validation pending (lit-val)	***pending*** served by lit-ew (not landed at this head) providers pending (lit-ref): ML · Rc	

	Process · validation	Generation	$\lvert A_{\text{loop}}\rvert^2$ · per point, batched, default backend
no loop-induced package has landed at this head: every row **pending** (served by the package named in its cell) · default backend TIRust · nothing estimated			
	legs · corrections · PLAN §4.1 ids; largest deviation per reference (lit-val)	generate + plan export, peak memory; provider generation	median per point at $B = 128$; provider times (lit-ref); gate (lit-perf) or rule; backend
$2 \to 2$	$gg \to aa$ 4 legs · QCD · LI-A5 validation pending (lit-val)	***pending*** served by lit-amp (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to zz$ 4 legs · QCD · LI-21 validation pending (lit-val)	***pending*** served by lit-amp (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to w^+w^-$ 4 legs · QCD · LI-22 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to za$ 4 legs · QCD · LI-23 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to hh$ 4 legs · QCD · LI-24, LI-A7 validation pending (lit-val)	***pending*** served by lit-amp (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to hz$ 4 legs · QCD · LI-25 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to hg$ 4 legs · QCD · LI-26, LI-A7 validation pending (lit-val)	***pending*** served by lit-amp (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$u\bar{u} \to hg$ 4 legs · QCD · LI-26, LI-A7 validation pending (lit-val)	***pending*** served by lit-amp (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$ug \to hu$ 4 legs · QCD · LI-26 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to zg$ 4 legs · QCD · LI-27 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to ag$ 4 legs · QCD · LI-27 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$e^+e^- \to ha$ 4 legs · QED · LI-A4, LI-E3 validation pending (lit-val)	***pending*** served by lit-ew (not landed at this head) providers pending (lit-ref): ML · Rc	
	$aa \to aa$ 4 legs · QED · LI-A6, LI-E2 validation pending (lit-val)	***pending*** served by lit-ew (not landed at this head) providers pending (lit-ref): ML · Rc	
	$aa \to zz$ 4 legs · QED · LI-E2 validation pending (lit-val)	***pending*** served by lit-ew (not landed at this head) providers pending (lit-ref): ML · Rc	
	$aa \to hh$ 4 legs · QED · LI-E2 validation pending (lit-val)	***pending*** served by lit-ew (not landed at this head) providers pending (lit-ref): ML · Rc	
	$e^+e^- \to hh$ 4 legs · QED · LI-E3 validation pending (lit-val)	***pending*** served by lit-ew (not landed at this head) providers pending (lit-ref): ML · Rc	
$1 \to 3$	$h \to ggg$ 4 legs · QCD · LI-37 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · Rc	
	$z \to ggg$ 4 legs · QCD · LI-37 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · Rc	

	Process · validation	Generation	$\lvert A_{\text{loop}}\rvert^2$ · per point, batched, default backend
no loop-induced package has landed at this head: every row **pending** (served by the package named in its cell) · default backend TIRust · nothing estimated			
	legs · corrections · PLAN §4.1 ids; largest deviation per reference (lit-val)	generate + plan export, peak memory; provider generation	median per point at $B = 128$; provider times (lit-ref); gate (lit-perf) or rule; backend
$2 \to 3$	$e^+e^- \to ggg$ 5 legs · QED · LI-E4 validation pending (lit-val)	***pending*** served by lit-ew (not landed at this head) providers pending (lit-ref): ML · Rc	
	$gg \to hhh$ 5 legs · QCD · LI-31 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to hgg$ 5 legs · QCD · LI-32, LI-A7 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to zzg$ 5 legs · QCD · LI-33 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to aag$ 5 legs · QCD · LI-33 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to w^+w^-g$ 5 legs · QCD · LI-33 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to zzz$ 5 legs · QCD · LI-34 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to hzg$ 5 legs · QCD · LI-35 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to agg$ 5 legs · QCD · LI-35 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$gg \to zgg$ 5 legs · QCD · LI-35 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$u\bar{u} \to hgg$ 5 legs · QCD · LI-36 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
	$u\bar{u} \to d\bar{d}h$ 5 legs · QCD · LI-36 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL · Rc	
$2 \to 4$	$gg \to hhgg$ 6 legs · QCD · LI-41 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL	
	$gg \to hggg$ 6 legs · QCD · LI-42 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL	
	$gg \to e^+e^-\mu^+\mu^-$ 6 legs · QCD · LI-43 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML · OL	
	$gg \to aagg$ 6 legs · QCD · LI-44 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): OL	
	$gg \to zzzz$ 6 legs · QCD · LI-44 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): ML	
$2 \to 5$	$gg \to e^+e^-\mu^+\mu^-g$ or $gg \to hgggg$ 7 legs · QCD · LI-51 validation pending (lit-val)	***pending*** served by lit-r2 (not landed at this head) providers pending (lit-ref): OL	
EFT	rank > N+2 EFT row (HC_NLO_X0 or SMEFT) EFT · LI-X1 · TIRust only validation pending (lit-val)	***pending*** served by lit-eft (not landed at this head) providers pending (lit-ref): ML	

Reading the summary

Table 1 has one row per test process, by multiplicity. A cell gives pyAmpliCol's value in bold and, at its right, the multiplier *pyAmpliCol ÷ the fastest provider* on this host (in bold among the provider values). A provider time marked ≤ ($u\bar{u} \to gggg$) comes from a run outside the protocol that does the production work and more (Recola's reference run evaluates the IR poles too), so it bounds that provider's production time from above, and the multiplier against it is a lower bound (≥); it is never a gate limit. Only run times are coloured: green faster, white at parity, red slower ($\log_2$ scale, clamped at ×1/16 and ×16). Generation cells carry their multiplier uncoloured (generation time is not the main metric, user 2026-10-24); leading colour has no provider ratio. Grey cells say why a value is missing. Run times are per phase-space point in batches, finite part, stability checks off; providers in production mode (Section 1). Every value is the **most recent measured configuration** of its cell (by git ancestry of the revisions; a newer check that only repeats part of an older comparison does not replace it): $2 \to 2$ to $2 \to 5$: the engine on `loop` with native per-point inputs, gate G1 at `ca602436` (2026-11-06). The engine before G-int (Python per-point inputs, G1 at `1c971289`) is history in Section 6. The column header of Table 1 names these configurations once.

Gate chips (Table 1 and Figure 1): the time over the binding limit (MadLoop's time, or 1.30 × the faster of Recola and OpenLoops), **pass** or **fail** when the whole uncertainty band is below or above 1, else **open**. The user's gate rule of 2026-10-24: a process whose time per point at $B = 128$ is below 5 ms has no blocking gate, its verdict is "below 5 ms: not gated (user 2026-10-24)" and its chip ("<5 ms: not gated") keeps the measured ratio ($u\bar{u} \to d\bar{d}g$ 1.72 ms, 0.87 of its limit, $u\bar{u} \to u\bar{u}g$ 2.64 ms, 0.69 of its limit, $u\bar{u} \to e^-e^+g$ 0.788 ms, 0.83 of its limit, $u\bar{u} \to ggg$ 4.34 ms, 0.27 of its limit and $u\bar{u} \to e^-e^+gg$ 4.21 ms, 0.32 of its limit). $gg \to gg$ ($2 \to 2$) has no gate; $gg \to t\bar{t}e^-e^+g$ (EW, `data/msm-g1.json`) is gated like the QCD rows since the final G1 retime (M-LIT-a0; at M-SM it was timed indicatively, [U5]); its leading colour is refused by design: EW LC/NLC is never supported (user 2026-10-30, Q4; typed colour-declaration-coverage). The gate is taken at $B = 128$ only (user 2026-10-28, `gate_decisions.gate_batch`): $B = 1$ and the $2 \to 5$ rows' $B = 16$ are measured and **reported** with their ratios (Table 12, Table 13), without a verdict or a verdict colour.

Stage shares (the bar of the full-colour cell; the largest is named; Table 14): inputs = trees, Born and counterterm inputs (on `loop` the executor's own timer of its native trees, since lc-perf); glue = point parameters, input conversion, assembly and the rest of the call. Times use the protocol's schedule B.

References and deviations. ML MadLoop 3.7.0, Rc Recola 2.2.5 (COLLIER), OL OpenLoops 2.1.5 (prebuilt libraries: no generation time), NJet DD NJet 3.1.2 in double-double. Deviation $|a - b| \,/\, \max(|a|, |b|, |B|\ \alpha_s/2\pi)$ of the binary64 finite part at the timed points (median: over the points, the largest over the references); † the largest lies within the binary64 stability estimate of the same configuration at that point; ¶ the largest is the reference's own: pyAmpliCol's double-double value at that point lies as far from it, and its binary64 value lies within E of the double-double one. ‡ $gg \to gggg$: the reference check's points 4 and 8 (1.5×10^{-7} and 2.6×10^{-8} from NJet DD) exceed its tolerance 2×10^{-8} in binary64: **accepted (user 2026-10-25): binary64 within the estimate, DD = NJet** (estimates E 1.5×10^{-6} and 3.4×10^{-7}; pyAmpliCol's double-double values equal the reference's); the check keeps its tolerance and, asserting double-double at these points since the user's E1 ruling (2026-10-29), passes (`data/milestone.json` `validation_decisions`).

Configurations (main-configuration rule of `scripts/build_summary.py`; `data/summary.json` names the configuration and source of every value). $2 \to 2$ to $2 \to 5$: the engine on `loop`, gate G1 at `ca602436`, build `loop-native-v2-slice-ca602436`, native per-point inputs (Section 6). Generation excludes model loading; unbatched development timings (e.g. a seven-gluon point with its rotation check) are not in this table (Section 8). G1 generates every row cold since milestone m1026a (`11542fd5`), in its own process with a fresh cache directory, so the counterterm caches are rebuilt in every run (Section 6, generation protocol).

Notes under a value (`data/milestone.json` `cell_notes`) qualify it and leave it as measured, in amber a known limitation of the value, in grey its run conditions: $gg \to t\bar{t}e^-e^+g$, full colour ("model: the workspace's py3 G_mu copy"): timed on reference/models/loop_qcd_qed_sm_Gmu_py3 (the harness's processes.json default, as at M-SM), not on the shipped patched UFO (UFOmodeldb 6ca4b120 + the user's issue-1 patch) that the M-SM campaign validated; the speed gate is taken on this timing.

The $2 \to 6$ rows (M-FINAL item 1; `data/2to6.json` through `summary.json` `preliminary`; their records and checks below) are outside the main-configuration rule: pyAmpliCol only (no provider exists at $2 \to 6$), so no multiplier, no gate and no colour in any cell; engine `2f3f76d9` (`work/v2-mf-8g-3`) for every record, an ancestor of this milestone's pin, not re-run on it (loop stores made before model compiler 17 do not load on the pin, so the rows keep their engine); generation = generate + plan export (store, artifact and load beside it; "FC" marks the full-colour generation, LC below it); the timed evaluation is binary64, finite part, checks off, one pinned thread, at the largest batch measured. The process cell lists the internal checks (ε^{-2}: double pole against the universal value; E: the recommended stability estimate, with the points above 10^{-5}; rescue: the stream points of 16 the default rescue policy's binary64 level sends up; Ward on the external gluons; binary64 against double-double; Catani at full colour; NGluon at leading colour on the all-gluon row) and, in grey, those not run or not asserted.

Figure 1 places each row's batched full-colour value of Table 1 (dots; the most recent measured configuration: $2 \to 2$ to $2 \to 5$: `ca602436`, native per-point inputs, G1 2026-11-06) against the providers (diamonds, the darkest the fastest) and the gate limit (red bar, the shaded side fails; grey bar: the limit of a process below 5 ms, not gated; dotted line: 5 ms), for the rows with a matched provider, at $B = 128$. The ratios and bands at every batch size: Table 13.

The loop-induced table

Table 2 follows Table 1's rules for loop-induced processes, $|A_{\text{loop}}|^2$ without a Born (`data/summary-li.json`, written by `scripts/build_summary.py` from the row registry `data/milestone.json` `loop_induced` and the loop-induced records). One row per

process of the M-LIT validation matrix (m-lit/PLAN.md §4.1; its ids in the process cell), grouped by multiplicity; the backend-unit, zero, identity, observable and complex-mass-scheme rows (LI-B1, LI-B2, LI-Z1, LI-Z1b, LI-Z2, LI-W1, LI-L1, LI-P1, LI-C1, LI-S1, LI-O1) stay in the validation tables. The cell is the batched $B = 128$ time per point (median, finite part, one pinned thread, checks off) on the default backend TIRust; NinjaRS per amplitude, labelled in the cell, where TIR does not serve the row yet; the other backend's time beside it when both ran. Its multiplier against the fastest matched provider (MadLoop's loop-induced production setting 7|6|1, OpenLoops ls, Recola; production mode, from lit-ref) colours it with Table 1's scale; generation is uncoloured. Rows of four legs or fewer ($2 \to 1$, $1 \to 2$, $2 \to 2$, $1 \to 3$) are reported without a gate; from $2 \to 3$ upward the gate is Table 1's (faster than MadLoop and at most 1.30× each matched OpenLoops and Recola time, on TIRust), "<5 ms: not gated" below 5 ms per point, no gate without a matched provider; verdict chips come with lit-perf's records, "no verdict yet" before them. Grey cells are explicit markers: **pending** (the package that serves the row, named in the cell, has not landed at this head, or the input file that fills the cell does not exist yet), **not run** (a landed package without a record of the row), and a provider's reason where it has no matched time (no library, beyond the 45-minute rule). Nothing is estimated or extrapolated; each cell shows the newest record of its row (git ancestry of the record revisions). At this head 0 of 42 rows are served; inputs: data/li-timing.json not written yet (the row's serving package), data/li-providers.json not written yet (lit-ref), data/lit.json not written yet (lit-val); the registry against lit-ref's manifest (tools/loop_reference/li_campaign.json): not written yet (lit-ref).

$2 \to 6$: records and checks (M-FINAL item 1)

The $2 \to 6$ rows of Table 1 come from data/2to6.json, which scripts/build_2to6.py derives byte for byte from the mf-2to6 run records (harness store records, stored-checks, Ward, NGluon and timing records, watchdog status files and the stores' generation.json; 106 sources). Engine (every record): work/v2-mf-8g-3 at 2f3f76d9 (d24a01ae (slice: mf-8g-2b integrated) + mf-8g-3), native loop-native-mf-2to6-2f3f76d9-lto (digest 2bae17f9…), an ancestor of this milestone's pin (integrated as mf-8g-3, df1d2a5d), not re-run on it. Protocol: timing: binary64, finite part only (poles off), no stability check, one thread on one pinned core, warm-up 16 points excluded, one block of B points, setup heap frozen out of the cyclic collector; indicative (shared host); particle veto PDG 6 (top-free, $N_f = 5$, user Q-A), 16 lanes. No external reference exists at $2 \to 6$: acceptance is internal (double pole against the universal value, Catani at full colour, NGluon at leading colour, Ward, double-double, the stability estimate) and the speed is reported without a gate (cells uncoloured).

Table 3: **The $2 \to 6$ records** (data/2to6.json rows): generation stages (wall time · peak), the stored artifact and its load, the timed evaluations (finite part, poles and checks off; warm-up excluded, including lazy compilation) and the checks (grey: not run or not asserted, with the reason in the data). Peaks are the harness's per-stage high-water marks and the timing process's maximum resident set (GiB = 2^{30} bytes).

Record	Generate	Export · plan	Store	Artifact · load	Evaluation	Checks
$gg \to gggggg$ LC · 47 348 families	23.8 min · 51.4 GiB	22.3 min · 63.9 GiB · 327 MB	21.5 min · 58.0 GiB	0.364 GB · 3.20 s, 0.74 GiB	**102 s/pt** at $B = 16$ (27.2 GiB); **98.0 s/pt** at $B = 128$ (27.2 GiB); warm-up 36.5 min	double pole, stability, double double, rescue, ngluon; ward not asserted
$gg \to gggggg$ FC · 47 348 families	14.3 min · 39.6 GiB	41.0 min · 74.0 GiB · 486 MB	25.8 min · 76.4 GiB	0.53 GB · 4.21 s, 1.06 GiB	**201 s/pt** at $B = 16$ (40.1 GiB); $B = 128$ not run; warm-up 64.7 min	double pole, stability, catani, double double, rescue; ward not run
$u\bar{u} \to gggggg$ LC · 23 437 families	495 s · 33.8 GiB	12.2 min · 47.0 GiB · 30.7 MB	198 s · 34.9 GiB	0.047 GB · 510 ms, 0.14 GiB	**9.54 s/pt** at $B = 16$ (11.1 GiB); **10.9 s/pt** at $B = 128$ (12.6 GiB); warm-up 200 s	double pole, stability, double double, rescue, ward
$u\bar{u} \to gggggg$ FC · 47 348 families	336 s · 16.0 GiB	68.0 min · 92.9 GiB · 814 MB	449 s · 89.3 GiB	0.143 GB · 3.85 s, 1.66 GiB	**65.0 s/pt** at $B = 16$ (30.7 GiB); **62.1 s/pt** at $B = 128$ (31.7 GiB); warm-up 18.3 min	double pole, stability, catani, double double, rescue, ward
$u\bar{u} \to u\bar{u}gggg$ LC · 13 318 families	248 s · 13.7 GiB	258 s · 18.2 GiB · 5.87 MB	14.6 s · 15.3 GiB	0.009 GB · 198 ms, 0.09 GiB	**2.58 s/pt** at $B = 16$ (4.85 GiB); **2.67 s/pt** at $B = 128$ (7.38 GiB); warm-up 46.0 s	double pole, stability, double double, rescue, ward
$u\bar{u} \to u\bar{u}gggg$ FC · 40 454 families	229 s · 8.27 GiB	26.0 min · 42.5 GiB · 315 MB	46.4 s · 41.1 GiB	0.038 GB · 1.51 s, 0.68 GiB	**18.2 s/pt** at $B = 16$ (34.6 GiB); **19.2 s/pt** at $B = 128$ (36.0 GiB); warm-up 359 s	double pole, stability, catani, double double, rescue, ward

Checks (data/2to6.json rows[].checks):

- **Poles.** The double pole equals the universal value at points 0–2 on all 6 records (≤ 8.4e−11). Catani's poles (native, from the stored dipole orbits) agree on the 3 full-colour records within 10 E (largest single-pole deviation 7.1e−5), in double-double at point 0 to binary64 resolution. At leading colour on $gg \to 6g$, NGluon (NJet) on exactly rotated points agrees in double-double to 6.3e−13 (finite part) and in binary64 within 10 E; the leading-colour single pole of the quark rows is not checked (the native Catani operator is full colour only, NGluon serves all-gluon rows only).
- **Double-double.** At point 0 the binary64 error against double-double is below its estimate E on all 6 records (largest 1.2e−4).
- **Ward** (external gluons, $\varepsilon_j \to p_j/p_j^0$, the real part of the interference). It passes on the 4 quark records in binary64 (virtual/scale ≤ 2.1e−11), on $u\bar{u} \to u\bar{u}gggg$ FC also in double-double. On $gg \to 6g$ LC every Ward point lies beyond the binary64 scope: recorded, not asserted

(inconclusive; each ratio below its E); the full-colour and double-double Wards of $gg \to 6g$ were not run (user, 2026-11-04 ~17:35: enough evidence).

- **Stability.** The recommended estimate exceeds 10^{-5} at points 0, 1 of $gg \to 6g$ (LC) and 0, 1 (FC); the default rescue policy's binary64 level sends 11 (LC) and 11 (FC) of 16 stream points up at 8 gluons, 1, 2, 2 and 4 on the quark records: rescue cases, not failures.
- **Caveats** (`global_caveats`): no external reference exists at $2 \longrightarrow 6$: acceptance is internal (Ward, rotations, Catani poles, double pole, double-double, NGluon at LC); timings are indicative: shared host (load 100-250 of 384 cores), one pinned thread, one block per B.
- **Not run or open** (`not_obtained`, as the data words it):
  - g g > g g g g g g (LC, N_f = 5, no top): ward not asserted (inconclusive: the harness verdict is FAILED (no point asserted) because every point has a leg with E beyond the binary64 scope (1e-9); the rows are recorded only, so this row carries no Ward evidence)
  - g g > g g g g g g (LC, N_f = 5, no top): binary64 E above 1e-5 at points [0, 1] (a rescue case, not a failure)
  - g g > g g g g g g (LC, N_f = 5, no top): ward-lc-dd not run (USER 2026-11-04 ~17:35: no further Ward / Catani / DD check job, enough evidence (stopped by the coordinator))
  - g g > g g g g g g (FC, N_f = 5, no top): timing B = 128 not run (USER 2026-11-04 ~17:48: 8-gluon B = 128 skipped, B = 16 measured (stopped by the coordinator in its timed block))
  - g g > g g g g g g (FC, N_f = 5, no top): ward not run (store-ward-fc, ward-fc, ward-fc-dd: USER 2026-11-04 ~17:35: no further Ward / Catani / DD check job, enough evidence (stopped by the coordinator))
  - g g > g g g g g g (FC, N_f = 5, no top): binary64 E above 1e-5 at points [0, 1] (a rescue case, not a failure)
  - u u~ > g g g g g g (LC, N_f = 5, no top): LC single pole unchecked (the native Catani operator is full colour only and NGluon serves only all-gluon rows); the LC double pole is checked
  - u u~ > g g g g g g (FC, N_f = 5, no top): ward-fc-dd not run (USER 2026-11-04 ~17:35: no further Ward / Catani / DD check job, enough evidence (stopped by the coordinator))
  - u u~ > u u~ g g g g (LC, N_f = 5, no top): LC single pole unchecked (the native Catani operator is full colour only and NGluon serves only all-gluon rows); the LC double pole is checked

Optimisation rounds at seven legs (M-FINAL)

Table 4: **The optimisation rounds at seven legs** (`data/rounds-7leg.json`, from `m-final/round-table-7leg.md`): $gg \to ggggg$ (top-free, $N_f = 5$) and, below, $u\overline{u} \to 5g$ at full colour, the baseline (for $gg \to ggggg$ the 8g0 records, 8g0-report records, before round 1; for $u\overline{u} \to 5g$ mf-8g-2's A/B base side) against the slice after round 2 (`d24a01ae`), one pinned thread, s and peak GiB per stage, factors now/baseline; values bitwise through every round, the stored artifacts byte-identical between the rounds' records; indicative (shared host). The $2 \to 6$ records above ran on the engine after these rounds plus mf-8g-3.

Stage	LC: baseline → now	Full colour: baseline → now
generate	281 s · 4.44 GiB → **62.0 s · 3.30 GiB** ×0.22 · ×0.74	279 s · 4.05 GiB → **46.0 s · 2.52 GiB** ×0.16 · ×0.62
export	349 s · 10.1 GiB → **93.0 s · 4.05 GiB** ×0.27 · ×0.40	353 s · 9.41 GiB → **93.0 s · 4.54 GiB** ×0.26 · ×0.48
store	161 s · 10.8 GiB → **50.0 s · 3.82 GiB** ×0.31 · ×0.35	169 s · 11.0 GiB → **36.0 s · 4.66 GiB** ×0.21 · ×0.42
generation total (wall, max peak)	13.2 min · 10.8 GiB → **205 s · 4.05 GiB** ×0.26 · ×0.38	13.4 min · 11.0 GiB → **175 s · 4.66 GiB** ×0.22 · ×0.42
artifact (pacbin)	70.2 MB → **16.9 MB** ×0.24	83.1 MB → **21.2 MB** ×0.26
load (time, resident)	93.0 s · 8.43 GiB → **18.6 s · 0.73 GiB** ×0.20 · ×0.09	106 s · 9.69 GiB → **15.3 s · 1.03 GiB** ×0.14 · ×0.11
evaluation B = 16 (per point, peak)	6.20 s · 13.1 GiB → **3.95 s · 3.73 GiB** ×0.64 · ×0.28	11.8 s · 15.1 GiB → **5.10 s · 4.41 GiB** ×0.43 · ×0.29
evaluation B = 128 (per point, peak)	5.90 s · 24.4 GiB → **4.02 s · 5.73 GiB** ×0.68 · ×0.23	11.1 s · 21.0 GiB → **5.22 s · 6.15 GiB** ×0.47 · ×0.29
load (time, resident) ($u\overline{u} \to 5g$)	—	10.3 s · 1.90 GiB → **3.50 s · 0.34 GiB** ×0.34 · ×0.18
evaluation B = 16 (per point, peak) ($u\overline{u} \to 5g$)	—	1.62 s · 4.57 GiB → **1.40 s · 1.18 GiB** ×0.86 · ×0.26
evaluation B = 128 (per point, peak) ($u\overline{u} \to 5g$)	—	1.65 s · 6.04 GiB → **1.32 s · 3.65 GiB** ×0.80 · ×0.60

The spot check re-measured the LC cells within 5 % (evaluation) and 0.5 % (peaks); a first generate run on a core with a busy SMT sibling was 1.5× slower, the repeat on an idle core reproduced 62 s within 2 %. After round 2 the store's companion re-export is the largest remaining generation piece (43.7 of 50 s LC, 29 of 36 s FC at 7 legs; ≈950-1,100 s at 8 legs).

1 Status at this milestone: M-LIT-a0

On loop (engine ca602436; merges since 7266c160, engine c8885cc5)

936e81b3	CI	Lightweight CI fixes
ec99ad04	1.0.2	Version 1.0.2
5fee3347	#40	HEFT H g g U(1) subtraction
b81b77b3	CI	Eager-portability timeout
053a0778	SDK	C-ABI example client
36e60f6c	CI	Restricted-Symbolica test isolation
ca602436	CI	Rust workspace tests

Before (engine c8885cc5): v2 waves 0-1 (gates S1, G0), native executor, colour, stability API, schedule M, observables, NJet reference, public v2 engine with plan v4 artifacts, counterterms as plain tree insertions (addendum J), runtime pole selection, Rust A0-A3, the G1 harness, G-int (native per-point inputs and generation stages), the summary-first report; at milestone m1025a ct-perf, ew-poles (EW guards, QED Catani operator), lc-perf, artifact-portability and ew-lines; at m1025c native-dag A4, ew-gauge EWI-4a, precision-ap, ew-ct EWI-0/EWI-3, measure-g1-7leg (G1 at 2->5), reference-relocate, model-determinism, native-lowering D phase 1 part 1 and onshell-consistency; at m1026a ew-gauge EWI-4b, native-seeds B phase 1, native-cells C phase 1, genericity-fixes, ew-poles EWI-6b, fermion-lines and generation-scale-2; at m1028a ew-ct EWI-3R, ew-epsilon EWI-7b, fail-closed, sdk-parity, ew-leptons EWI-5, identical-flavour-lc and gate-b128; at M-QCD onshell-threshold, fail-closed-2, glibc-floor, MQ-WARD, MQ-FFT, MQ-REF, MQ-MODEL, MQ-STAB, MQ-VALIDATE, MQ-NC, MQ-5F, MQ-EWJOIN, MQ-SELECT, MQ-LEGACY-REMOVE and g1-accuracy; at M-SM ew-native, ew-validate, MS-REF, MS-PHOTON, MS-CMS-N, MS-MODEL, MS-CMS, MS-GEN, MS-POLES, MS-UFODB, MS-UFO-PATCH, MS-VALIDATE, ms-policy-fix and MS-MASTERS; at M-FINAL-b MF-CORR, MF-RESCUE, MF-SMALL, MF-GM-HC, MF-UP (SymJIT 2.28.3), MF-CLEAN, MF-MAP, MF-DEC-1, MF-8G-0, MF-RUST, MF-DEC-2, MF-SMALL2 and the test-only mf-map-fix (f3d9e361); at M-FINAL MF-SDK, mf-map-fix, MF-8G-2, MF-8G-1b, MF-SMALL3, MF-MASKS, MF-8G-1c, MF-PERF-OPT, MF-8G-1, MF-HC05, MF-8G-2b, lp-models, lp-colour, MF-8G-3, MF-MASKS-ROWS, lp-precision, MF-2TO6, MF-LIMITS-PY, MF-LIMITS-RUST, MF-NORM and MF-SYMBOLICA (c8885cc5). **Native build** `loop-native-v2-slice-ca602436` (digest `4cbb549f`…). **Integrated after it, not re-timed:** `a2ade899` lit-ci (CI only: licensed loop jobs in tests.yml (no native input changed, native digest 4cbb549f)); `35a5a9b2` lit-col (loop-induced amplitude colour basis (new native digest b044768c); the Born-interfered path not re-timed: the package's Born A/B is byte-identical (slice-status 2026-11-06 ~10:25)).

M-LIT-a0 (the first milestone push of M-LIT, loop-induced processes and the TIRust backend; user green light 2026-11-06 ~07:05; `m-lit/PLAN.md` §3.2 LIT-RETIME)

- **The final G1 retime** (deferred from M-FINAL, `m-final/PLAN.md` §2.4): every G1 loop artifact regenerated (model compiler 17) and every G1 row re-timed on `ca602436` (`loop-native-v2-slice-ca602436`) with the G1 protocol (finite part, $B = 128$, one pinned thread, checks off, providers' production records; no quiet window); the EW row is gated from now on; the $2 \to 6$ rows stay `data/2to6.json`'s. per point at B = 128 against M-SM's G1 (71e2041e): full colour x0.69-1.12 (median x0.84, 15 rows), leading colour x0.67-1.00 (median x0.88, 15 rows); g g > g g g g g g full colour 7.26 -> 4.99 s. This run is the baseline of every M-LIT no-regression A/B of the Born-interfered default (NinjaRS).
- **The loop-induced summary table** (Table 2, `data/summary-li.json`): the 42 process rows of the M-LIT validation matrix by multiplicity, every row **pending**: no loop-induced package has landed at this push. Loop-induced requests are still refused typed (`loop-induced`) at this head.

M-FINAL, closed 2026-11-05 (engine `c8885cc5`; `m-final/PLAN.md`, the user's rulings of 2026-10-31 ~11:00, the lean close of 2026-11-05 ~07:30)

- **Item 1, 2 → 6** (Table 1, records): $gg \to 6g$ ($N_f = 5$, no top), $u\bar{u} \to 6g$ and $u\bar{u} \to u\bar{u}gggg$ at LC and full colour generated, stored (every step below 150 GB), loaded and evaluated on one engine (`2f3f76d9`, an ancestor of the pin); every asserted internal check passes, no provider exists.
- **Item 2, decays 1 → *n*** (M-FINAL-b, Section 2): XX01–XX08 = MadLoop's quadruple precision, QCD and EW.
- **Item 4:** the rescue stack (Python native evaluate API only; M-FINAL-b); HC_NLO_X0 HC01, HC02, HC03, HC04, HC05, HC06 = MadLoop 3.7.0 (Table 6; HC05 since mf-hc05); the optimisation rounds (seven-leg table: generation ÷3.9 LC / ÷4.6 FC, load resident ÷12 / ÷9, $B = 128$ peak ÷4.3 / ÷3.4, values bitwise); the Recola correlation oracles (M-FINAL-b). The final G1 retime followed at M-LIT-a0 (box above).
- **Also since M-FINAL-b:** the SDK v4 extras (mf-sdk); lsimon125's PRs #26 #29 #30 #32 #34 #36 #37 #38 #39 (#24 adapted with credit): #38 fixes a real 4× over-count of compiled-mode HEFT quark tree amplitudes (built-in SM-HEFT and HC_NLO_X0, now = MG5 and the MadLoop Borns), #34 moves binary64 quark-line values by rounding only ($\lesssim 10^{-12}$, accepted); the built-in SM fixes; the mechanical hard-coded limits lifted; Symbolica off the fork (boxes below).
- **Validation scope:** the campaigns of M-FINAL-b (`617b59cd`: M-SM, M-QCD, decays, HC, correlations) stand, plus package evidence (A24 passes rule 4 with native Catani: mf-8g-3; HC05 = MadLoop: mf-hc05). The pin itself carries its integrators' targeted tests; the global tests run in CI on `loop` (no local whole-suite sweep, wheel self-test or retime: user 2026-11-05 ~07:30).

Gate G1 (Section 6; finite part, $2 \to 3$ and beyond)

- **Re-timed at M-LIT-a0** (the final retime deferred from M-FINAL): every G1 row on `ca602436` (`loop-native-v2-slice-ca602436`), the EW row gated; per point at B = 128 against M-SM's G1 (71e2041e): full colour x0.69-1.12 (median x0.84, 15 rows), leading colour x0.67-1.00 (median x0.88, 15 rows); g g > g g g g g g full colour 7.26 -> 4.99 s. The evaluation and memory rounds of M-FINAL (seven-leg table) are in these cells now.
- **Most recent configuration** (Table 1 and Figure 1): native per-point inputs on `loop` (G1 at `ca602436`) at $B = 128$: 9 of 9 gated processes pass; $u\bar{u} \to d\bar{d}g$, $u\bar{u} \to u\bar{u}g$, $u\bar{u} \to e^-e^+g$, $u\bar{u} \to ggg$ and $u\bar{u} \to e^-e^+gg$ run below 5 ms per point: **not gated** (user 2026-10-24; 0.87, 0.69, 0.83, 0.27, 0.32 of the limit). $B = 1$ (16 lanes, the runs of `28a16255`, not re-timed) is **reported**, not gated (user 2026-10-28): ×1.0–5.3 of the limit (Table 13).
- **2 → 5** (full and leading colour at $B = 128$; $B = 16$ reported): 5 of 5 gated rows pass ($u\bar{u} \to e^-e^+d\bar{d}g$ 13.9 ms, 0.30 of its limit; $u\bar{u} \to e^-e^+ggg$ 47.9 ms, 0.15 of its limit; $gg \to t\bar{t}e^-e^+g$ 51.1 ms, 0.10 of its limit; $gg \to t\bar{t}e^-e^+g$ (EW) 131 ms, 0.10 of its limit; $u\bar{u} \to ggggg$ 1.31 s, 0.0023 of its limit, ≥×0.0072 against Recola's upper bound); $gg \to ggggg$ 4.99 s at full and 2.68 s at leading colour per point, no matched provider and no gate.
- **Before G-int** (Python per-point inputs, G1 at `1c971289`; history, Section 6): failed every gated process it generated, at ×11–20 of the binding limit at $B = 128$.
- **EW row** (`data/msm-g1.json`; gated since this retime, indicative at M-SM: [U5]): $gg \to t\bar{t}e^-e^+g$ (EW) 131 ms per point at $B = 128$, ×0.13 of Recola's time, pass ×0.10 (Section 4).

Built-in SM fixes (tree level; mf-norm 049a9389, with mf-limits-py d958f2ea for the initial-state averages): the built-in SM now equals the UFO sm in every lane. What was wrong, since when (2f876d1b, 2026-07-16, where the built-in model files enter this history, also on origin/main; 43e6fa5e, 2026-07-22) and by how much:

Process class	Lanes	Since	Effect (old value against the UFO sm)
every process without a coloured leg, decays included	all	2f876d1b	too small by $1/(8\pi\alpha)^2 = 27.797$ (old = $4e^4 \approx 0.036$ × correct); a run-time change of α rescaled them with the wrong power
initial Z or W; initial Higgs	all	2f876d1b	×3/2 per Z/W, ×1/2 per Higgs (spin averages; fixed by mf-limits-py)
$WW\gamma$ / WWZ vertices and four-boson contacts	recurrence	43e6fa5e	counted twice (old recurrence value against the compiled lane): $\gamma\gamma \to W^+W^-$ ×16, $e^+e^- \to W^+W^-$ ×4 (the compiled lane itself lacked the neutrino exchange, next row), $u\bar{u} \to W^+W^-$ point-dependent
$e^+e^- \to W^+W^-$; $e^+e^- \to \nu_e\bar{\nu}_e$	all	2f876d1b	the t-channel neutrino exchange missing (e.g. 245× at one 500 GeV point); $\nu_e\bar{\nu}_e$ refused (missing lepton/neutrino currents)
electroweak exchange between quark lines	all	2f876d1b	quark EW currents missing: QED-allowed requests silently gave the QCD-only value ($u\bar{u} \to d\bar{d}$: 1.445, 0.105, 1.079 × at three 700 GeV points); purely EW four-quark processes refused
the $HHHH$ contact ($hh \to hh$, ...)	compiled, eager, on-the-fly	2f876d1b	21.8, 97.8, 48.7 × at 700 GeV (and a further 1/27.797 before the colourless fix)
the $HHVV$ contact ($hh \to ZZ$, $hh \to W^+W^-$)	all	2f876d1b	204, 11.4, 72 × and 381, 15.6, 121 × at 700 GeV

Regression test: tests/integration/test_builtin_sm_ufo_sweep.py (default tier): every SM $2 \to 2$ process of the particle table (1993, reduced to 154 classes by crossing, charge-parity conjugation and the vertex-table relabellings), every $1 \to 2$ decay, 12 sample processes (seven $2 \to 3$, four $1 \to 3$, one $2 \to 2$) and three four-quark requests at explicit orders, in the compiled, eager, on-the-fly and recurrence lanes, against committed UFO sm values (MG5 3.7.0, no_masses) to 10^{-10} relative; where the UFO has no tree amplitude the built-in must refuse; a nightly variant runs 901 $2 \to 2$ processes against a freshly compiled UFO. Coloured processes: the compiled lane bitwise unchanged, the prepared lanes at ulp level ($\leq 1.3 \times 10^{-15}$, 9.6×10^{-14} at one ill-conditioned point; accepted). Typed refusals as the UFO sm: four-quark processes without gluon exchange in the recurrence lane, a massive top on a massless-b line (dirac-weyl-closure). The one-loop engine is unaffected (it refuses the tree-only built-in SM).

M-SM, kept (QCD + QED/EW on the shipped loop_qcd_qed_sm(_Gmu); m-sm/PLAN.md §5; validated at M-FINAL-b, not re-run here: user 2026-11-04 ~17:35)

- **Validation campaign** at 617b59cd (Section 4, Table 30): 71 rows in 10 classes, 69 pass, 2 partial, 0 fail (LL02, SC02: F7, accepted); no deviation above 10^{-5}, no wrong pole or sector; against M-SM's campaign no status change and no deviation change above 10^{-5} (40 values in 8 rows moved by $\leq 7.4 \times 10^{-11}$). Largest: DD from the reference 4.9×10^{-9} (SC02), binary64 from DD 1.9×10^{-10} (TV06), Catani DD 7.2×10^{-26} (CM07) at rule 4′s standard tolerance (relaxed by the user's ruling: TT05 $1.1\times10^{-29} \leq 1\times10^{-14}$, TT06 $4.2\times10^{-30} \leq 1\times10^{-13}$, SC03 $1.1\times10^{-6} \leq 1\times10^{-5}$, SC05 $4.1\times10^{-14} \leq 1\times10^{-12}$).
- **CMS** (Table 9): 9 rows = MadLoop CMS; the acceptance row CM07 ($gg \to t\bar{t}e^-e^+g$) DD 6.9×10^{-13}, binary64 from DD 5.6×10^{-11}, native inputs.
- **Public route** on both shipped EW UFOs: test_loop_v2_ew_public_models.py 25/25 at 617b59cd; wheel self-test 7/7 in two fresh venvs, its EW virtual from the shipped prepared artifact = MadLoop to 5.4e-15.
- **EW products:** per-leg photon tagging (loop.external_photons, default "final"), poles uv | ir | split by the model's regulators, native complex pole masses, EW generation stages in Rust, binary64 of a stored artifact without Symbolica (Table 25).

M-QCD, kept (QCD on the stock loop_sm; m-qcd/PLAN.md §5; validated at M-FINAL-b, A24 from mf-8g-3)

- **Campaign** at 617b59cd (2026-11-02, Section 5, Table 31): 78 rows, 77 pass, 0 partial, 1 fail (F7, accepted). **A24** ($gg \to ggggg$) passes since mf-8g-3: rule 4 with the native Catani operator from a stored QCD record (DD 3.5×10^{-26}, binary64 3.2×10^{-12} rule 1 = NJet's DD to 8.0×10^{-17} F11 resolved; the package's nightly record, re-run by its evidence audit on 2f3f76d9). At M-FINAL-b, against M-SM's re-run, 42 of 1212 values moved, all binary64 (the largest, information only: A14′s finite part against the providers $1.5 \times 10^{-8} \to 5.6 \times 10^{-8}$), every double-double value bitwise. Largest: DD from the reference 1.2×10^{-11} (A18), binary64 from DD 2.3×10^{-10} (C05), Catani DD 1.2×10^{-22} (A09).
- **At e18a024b** (not repeated: no status changed): public route on the shipped loop_sm, one row per class, DD from MadLoop $\leq 3.6\times10^{-13}$; FFT seed contraction = direct to rounding, ×0.46–5.9 of direct (Table 11). **Colour:** full colour by default; LC and NLC pass rule 7 on 13 rows. **Ward** up to seven gluons (passed (MQ-WARD)).

Dependencies (dependencies/manifest.json; no fork: user 2026-11-02 ~19:00 / ~21:30)

- **Symbolica** = symbolica-dev community 36cbac3a plus the single consolidated PR symbolica-dev/symbolica#163 (branch loop_me_pyamplicol, 20a9da94, 6 commits; moves upstream once merged); the alphal00p fork is gone from the graph (mf-symbolica).
- **NinjaRS** main b3a29c85 (PRs #7–#12 and #14 merged: UV residues from ninjars::fast::TensorUvResidues, the light-like spinors of #10) and **oneloopmaster** main 128ea664.
- **Python:** the Community build of 20a9da94 or PyPI symbolica 3.0.3; every stored saved evaluator is canonical (referenced

Hard-coded limits (mechanical ones: user 2026-11-04)

- **Lifted:** Python (mf-limits-py d958f2ea): compile timeout and slots, exact rational charges and grouping factors, built-in SM $1 \to n$ decays, uncapped mappings and fixpoints, N_c powers at any multiplicity; three tuning bounds became settings (color.*, generation.validation.*). Rust (mf-limits-rust 72fe20dc): one rank admission from the reducer, one leg boundary 1–64, checked narrowing, the S_m colour-FFT degree while $m!$ fits usize; values bitwise.
- **Not supported** (generalisations; typed refusals, no milestone planned): rank $> N + 2$ (NinjaRS, oneloopmaster), multi-fermion loop vertices, coloured n-point EFT contacts, nested singlet

symbols only, numbered by name), so loop artifacts are byte-identical from either.

contacts, octet/sextet fermions, an LC/NLC menu, binary64-only paths, C-API precision.

2 → 6 (M-FINAL item 1, **measured** Table 1, records and checks)

- **On one engine** (`data/2to6.json`; `2f3f76d9`, an ancestor of the pin), per point LC / FC: $gg \to 6g$ 98.0 s / 201 s (FC at $B = 16$), $u\bar{u} \to 6g$ 10.9 s / 62.1 s, $u\bar{u} \to u\bar{u}ggg$ 2.67 s / 19.2 s generation up to 73.6 min and 92.9 GiB (generate + plan export); every store step below 99 GB (budget 150 GB).
- **Internal checks** (no reference exists): double pole, Catani (FC), NGluon (LC, 8 gluons), double-double and the quark rows' Ward pass; the 8-gluon LC Ward is recorded, not asserted (beyond the binary64 scope); at 8 gluons most points are rescue cases (binary64 $E > 10^{-5}$).
- **Not run** (user, 2026-11-04): the 8-gluon full-colour and double-double Wards and the $u\bar{u} \to 6g$ FC double-double Ward (~17:35: enough evidence), the 8-gluon full-colour $B = 128$ timing (~17:48).

This milestone, M-LIT (in progress; user 2026-11-03 ~21:15, green light 2026-11-06 ~07:05; `m-lit/PLAN.md`)

- **Loop-induced processes** $|A_{\text{loop}}|^2$ (refused typed at this push, `loop-induced`): every row of Table 2 pending.
- **TIRust** (`huashengshao/TIRust`) as a tensor-integral reduction backend next to NinjaRS: the loop-induced default, on which the loop-induced speed gates are taken; M-LIT closes only when it lands.
- **Done at M-LIT-a0:** the final G1 retime, with the EW speed gate binding.

Missing or incomplete (typed refusal keys, listed and shown by the dry run)

- **This milestone (M-LIT, in progress):** loop-induced processes, including loop-induced decays ($h \to gg$) and the loop-squared members of multiparticle requests (refused typed at this push, `loop-induced`; a multiparticle request generates its members with a Born and records the others as skipped; every row of Table 2 pending), with TIRust as a tensor-integral reduction backend (the loop-induced default).
- **Refused (typed):** decay chains and `$` (`madloop-syntax`); Majorana fermions (`majorana-fermion`) and multi-fermion contact vertices (`multi-fermion-contact`); numerical process-group analysis (`process-group-analysis`; exact aliases and members are served); photon-only and partial EW vetoes on the NLOCT UFOs (`veto-unlabelled-counterterms`: their counterterm couplings are not split by loop content); the complex-mass scheme for a width whose mass is an internal parameter (`cms-internal-mass`: `loop_sm`'s W), an external unstable species under CMS (`cms-external-width`), a Yukawa input under CMS (`cms-yukawa-input`); a width outside CMS (`electroweak-width`, `coloured-width`; a decaying particle has zero own width); tagged photons without an $\alpha(0)$ input (`external-photon-scheme`); a spin-correlated leg without a vector source (refused at generation). HC_NLO_X0: $u\bar{u} \to x_0ggg$ (the U(1)-leak check), single-trace and C-odd five-point R_2 colours. **Tree level:** four-quark processes without gluon exchange in the recurrence lane (the other lanes serve them) and Dirac–Weyl currents the tree lanes cannot close (`dirac-weyl-closure`: a massive top on a massless-b line), both as for the UFO `sm`.
- **Not supported** (no milestone planned): reduction rank $> N + 2$ and the other capability generalisations of the limits box; helicity and colour-flow selection at NLC and full colour (LC has it; refused for mapped processes); the publication-grade timing campaign and the provider protocol-mode columns (user Q11).
- **Not supported, by design:** EW LC/NLC (`colour-declaration-coverage`, permanent: user 2026-10-30; no plan to lift it).
- **Accepted (validated at M-FINAL-b):** F7 — LL02 and SC02 partial (rule 4 not formed: the massive-muon G_μ copy writes the field renormalisation of massless leptons; rule 1 = MadLoop); SC03′s τ-mass relaxation of rule 4 to 10^{-5} (the EW counterterms omit the τ mass, the user's retained model limitation; MadLoop has the same offset); M-QCD's A09@collinear (F7: typed `Unstable` / `CutFailed{DegenerateGram}`, never a value).
- **Recorded test outcomes** (not gate failures; the integrators' records): `rusticol-capi` `static_c_api_advertises_all_symbolica_free_f64_capabilities` fails as before mf-limits-rust (a legacy-JIT capability advertised); under PyPI symbolica 3.0.3 six double-double rows of `test_loop_v2_sdk_extras.py` differ at ~10^{-30} relative (the Community build passes); the tree prepared-model kernels are not byte-reproducible run to run (operand order; their state sections are canonical).
- **Compatibility notes:** the shipped EW UFOs are UFOmodeldb `6ca4b120` plus the user's issue-1 patch (PROVENANCE records the local delta); MadLoop never inserts R2_G0tt_b on the unpatched models (mg5amcnlo#475): pyAmpliCol keeps the UFO-complete R_2 insertion (user Q7).
- **No longer missing** (since M-FINAL-b): the $2 \to 6$ processes; HC05 ($gg \to x_0gg$, its five-point $fffff$ R_2 contact); A24′s rule 4 (native Catani); the built-in SM tree defects; NinjaRS's light-like spinors along $-z$ (#10: the five xfails are plain cases, large unit factors keep the virtual); the alphal00p Symbolica fork; the SDK v4 extras. Earlier (M-FINAL-b): $1 \to n$ decays, permutation aliases and process-group members, Recola oracles of correlated Borns and virtuals, the stability rescue stack (Python API).

Performance gates and timing protocol: definitions

- **Gate.** Faster than MadLoop and at most 1.30× each matched Recola and OpenLoops time, on the **finite part**, from $2 \to 3$ upward, **taken at $B = 128$ only** (user 2026-10-28, MASTER_PLAN §5.8: $B = 1$ and the $2 \to 5$ rows' $B = 16$ are measured and reported with their ratios, without a verdict; a configuration without a $B = 128$ timing has none); $2 \to 2$ is reported without a gate. A process whose time per point at $B = 128$ is below 5 ms has no blocking gate: its ratio is shown, its verdict is "below 5 ms: not gated (user 2026-10-24)" (user decision; `data/milestone.json` `gate_decisions.threshold_rule`); optimisation effort goes to higher multiplicities. The electroweak row is timed and gated like the QCD rows since the final G1 retime (M-LIT-a0; `data/msm-g1.json`), after its indicative timing at M-SM ([U5]). Aspiration: fastest of all tools from $2 \to 3$.
- **Protocol.** Providers in production mode with poles off where they allow it (OpenLoops `truepoles 0`, Recola without IR poles, MadLoop production build); pyAmpliCol finite-only (`poles = "none"`) with stability checks off, one pinned thread, binary64, 16 lanes, the same frozen points (pyAmpliCol receives their exact rationals, the providers their binary64 decimal text); since milestone m1025a pyAmpliCol's setup heap is frozen out of Python's cyclic collector before the warm-up (a loaded runtime holds no generation objects; in the committed harness since `de8d3eb8`, which times G1 at `ca602436`; Section 6). Statistic: median per point; the uncertainty band combines pyAmpliCol's block spread and the provider's repeat spread; a gate whose band contains the limit is **open**.
- **Timing boundary.** pyAmpliCol: the public v2 runtime call (`NativeExecutor.evaluate`) for a batch of B distinct points, everything included (per-point inputs, loop stage, assembly); the public `Runtime.evaluate` of a loaded artifact adds its input parsing (Section 6). Providers have no batch API: one time per point serves every B. Generation: pyAmpliCol `LoopSlice.generate` plus the plan export, in G1 on one pinned core, cold (fresh cache directory) since `11542fd5` (m1026a); a development record's generation in Table 1 ran on the shared host under load (threads per increment: Table 22), without a multiplier in a G1 row; MadLoop graph generation, export and compilation; Recola process generation; OpenLoops libraries are prebuilt (no generation time).
- **Labels.** Table 1 and Figure 1 show the most recent measured configuration of each cell (main-configuration rule of `scripts/build_summary.py`; `data/summary.json` names the configuration and source of every value); the column header of Table 1 states it once, with branch and commit, and "Reading the summary" says when it is a development measurement not yet on `loop`. The detailed G1 tables (Section 6) show the configuration on `loop` and, labelled, its history: the runs of milestones M-SM (`71e2041e`), M-QCD (`e18a024b`), m1028a (`28a16255`), m1026a (`11542fd5`) and m1025a (`1d27426b`) and of G-int (`4e9b5a9e`), the Python per-point inputs before G-int and the development runs of `39824f9b`. Missing values are marked, never estimated or extrapolated; test durations are never used as event timings. The legacy per-diagram engine's timings of 18 October are history (data kept in `data/measurements.json`, not shown).

2 Decays $1 \to n$ (M-FINAL item 2)

Table 5: **Decay capability matrix** at 617b59cd (native loop-native-v2-slice-b36af289, campaign of 2026-11-01 data/dec-matrix.json from data/dec.json; per-row data: Table 29): every row against MadLoop's quadruple-precision reference of the same decay (MF-DEC-1 fixtures; the decaying particle at its real mass with zero own width, its absorptive self-energy dropped, as MadLoop does).

Row	Process	Corr.	Model · scheme	Sectors	DD vs ML	b64/DD	Poles (DD)	LC/NLC	Public · inputs	Status
XX01 1→2	$t \to bw^+$	QCD	loop_sm no_b_mass · QCD, 5F	QCD2 QED2	3.9e−16	1.4e−14	1.2e−31 Catani	N_c scan ✓	public ✓	pass
XX02 1→2	$t \to bw^+$	QCD+QED	$G_\mu \cdot G_\mu$	QCD2 QED2, QED4	2.9e−15	2.0e−13	2.3e−30 Catani + QED	—	native inputs ✓	pass
XX03 1→2	$z \to e^-e^+$	QED	qcd_qed_sm · $\alpha(m_Z)$	QED4	2.9e−15	2.0e−15	9.3e−31 Catani + QED	—	—	pass
XX04 1→2	$z \to u\bar{u}$	QCD+QED	qcd_qed_sm · $\alpha(m_Z)$	QCD2 QED2, QED4	6.1e−15	3.8e−13	1.9e−30 Catani + QED	—	—	pass form factor
XX05 1→2	$w^+ \to e^+$ ve	QED	qcd_qed_sm · $\alpha(m_Z)$	QED4	1.8e−15	5.4e−14	7.5e−31 Catani + QED	—	—	pass
XX06 1→2	$h \to b\bar{b}$	QCD	loop_sm · QCD, 4F	QCD2 QED2	4.1e−16	3.0e−15	1.9e−32 Catani	—	—	pass
XX07 1→3	$t \to bw^+g$	QCD	loop_sm no_b_mass · QCD, 5F	QCD4 QED2	5.6e−15	1.4e−13	4.5e−31 Catani	N_c scan ✓	public ✓	pass Ward
XX08 1→3	$z \to u\bar{u}g$	QCD+QED	qcd_qed_sm · $\alpha(m_Z)$	QCD4 QED2, QCD2 QED4	9.6e−14	5.2e−11	5.7e−27 Catani + QED	—	—	pass

Reading the matrix. The decay campaign re-ran every row at the pin's Python tree (617b59cd, the native loop-native-v2-slice-b36af289): 8 pass, 0 partial, 0 fail, as at its integration (no status change; the largest metric change 7.9×10^{-31}). Rows XX01–XX08 cover $1 \to 2$ and $1 \to 3$, QCD (loop_sm, five and four flavours) and EW (QED and joint QCD+QED corrections on both shipped EW UFOs, G_μ and $\alpha(m_Z)$); the deviations are the largest over every Laurent coefficient, sector, point and gating variant: DD from MadLoop's quadruple precision (rule 1), binary64 from DD where binary64 is asserted (rule 2), the poles against Catani, plus the QED operator for EW corrections, in DD (rule 4). XX02 runs on the native per-point inputs; XX01@public and XX07@public generate through Generator on the shipped loop_sm and evaluate the stored artifact with the recommended checks. **Further checks:** XX04: $Z \to u\bar{u}$ (QCD) equals the analytic massless vector form factor (1 test, pass) XX07: the Ward identity on the external gluon of $t \to bw^+g$ (8 checks; virtual 5.0×10^{-31} and Born 1.3×10^{-31} of the physical scale in DD, binary64 within its estimate) XX01 and XX07: LC and NLC are the kept powers of the N_c scan (3 tests, pass). Loop-induced decays are refused typed (loop-induced, e.g. $h \to gg$), chains and $ too (madloop-syntax).

3 Model genericity and correlation oracles (M-FINAL item 4)

Table 6: **HC_NLO_X0 rows against MadLoop 3.7.0** (data/hc.json: the campaign rows at 617b59cd, loop-native-v2-slice-b36af289, tests/unit/test_loop_v2_hc.py 15/15 passed; HC05 from the mf-hc05 package record (integrated e3d4ef4d)): the Higgs characterisation UFO at NLO (v1.3) with its restriction and parameter cards (κ settings), the largest relative deviation deviation |ours - MadLoop| / max(|ours|, |MadLoop|, |Born| alpha_s/(2 pi)) per Laurent coefficient (acceptance 10^{-5}), the poles against Catani, the rotation estimate and double-double where run.

Row	Process	Restriction · veto	Pts · values	Born	Finite	ε^{-1}	ε^{-2}	Catani	Rotation	Status · DD
HC01	$gg \to$ x0	heft · / t	1 · 6	3.5e−16	1.3e−15	2.0e−15	8.4e−16	1.9e−15	1.9e−14	pass
HC02	$gg \to$ x0 g	heft · / t	2 · 12	1.7e−15	1.8e−14	3.1e−15	2.0e−15	5.2e−15	1.4e−13	pass DD 1.5e−14
HC03	$u\bar{u} \to$ x0 g	heft · / t	2 · 8	1.4e−15	2.4e−13	1.4e−14	2.6e−15	3.4e−14	6.5e−13	pass
HC04	$u\bar{u} \to$ x0 z	no_b_mass	2 · 8	1.5e−14	8.4e−14	8.2e−15	8.1e−15	2.4e−14	1.2e−13	pass
HC05	$gg \to$ x0 gg	heft · / t	2 · 4	≤ 2.6e−12 (every coefficient)				DD and rotation: the nightly row passed		pass mf-hc05 record
HC06	$gg \to$ x0 $t\bar{t}$	no_b_mass	2 · 8	2.2e−15	2.9e−14	2.2e−14	1.6e−14	9.6e−15 (ε^{-2} only)	6.2e−13	pass

The HC rows ran through the generic fixes of MF-GM-HC (e2f0c78e): the UFO namespace head, the Feynman-gauge refusal only where massive vectors run in the loops (ew_policy.require_loop_gauge), positional UFO counterterm keys with zero extra positions (other key lengths are the listed source refusal ufo-counterterm-keys) and a four-point contact with a colour-singlet leg split through an adjoint auxiliary. HC05 ($gg \to$ x0 gg, refused typed at M-FINAL-b) passes since mf-hc05 (e3d4ef4d): the five-point R2 contact [G G G G X0] (colour f*f*f*f) is lowered through the generic exact N_c = 3 projection of vocabulary-refused counterterm colours onto products of at most two primitive adjoint tensors (loops/v2/ct_colour_projection.vocabulary_colours; zero residual required, else unlowered); binary64 = MadLoop 3.7.0 within 2.6×10^{-12} over every Laurent coefficient at 2 points × 2 scales, and the nightly row with double-double and a rotation passed (the package's record, not re-run here; HC01–HC04 and HC06 bitwise unchanged). Still refused typed: $u\bar{u} \to x_0ggg$ (the U(1)-leak check) and single-trace or C-odd five-point R_2 colours.

Table 7: **Colour, spin and spin-colour correlations against Recola** (data/corr.json at 617b59cd, tests/unit/test_loop_v2_corr_recola.py 12/12 passed; references tests/fixtures/loop/required/corr/<process>.json (Recola2 2.2.5, binary64)): per process and kind the entries and points, the correlated Born of the tree engine and of the one-loop runtime relative to Recola's (acceptance 10^{-11}), and the one-loop Laurent coefficients relative to the row's largest Recola coefficient (acceptance 10^{-8}).

Process	Kind	Entries · pts	Tree Born	Loop Born	Finite	ε^{-1}	ε^{-2}	Status
$u\bar{u} \to gg$	plain Born	2 · 2	1.7e−15	9.2e−16	9.4e−14	2.7e−14	1.3e−15	pass
$u\bar{u} \to gg$	colour	24 · 2	1.4e−15	7.7e−16	1.7e−13	8.1e−14	2.6e−14	pass
$u\bar{u} \to gg$	spin	4 · 2	1.8e−15	7.2e−16	1.0e−13	2.5e−14	2.1e−15	pass
$u\bar{u} \to gg$	spin-colour	12 · 2	1.5e−15	6.9e−16	2.5e−13	6.8e−14	4.1e−15	pass
$gg \to t\bar{t}$	plain Born	2 · 2	1.1e−15	6.7e−16	1.2e−15	1.5e−16	2.7e−17	pass
$gg \to t\bar{t}$	colour	24 · 2	2.0e−15	9.6e−16	2.8e−14	2.6e−15	6.5e−15	pass
$gg \to t\bar{t}$	spin	4 · 2	2.0e−15	1.2e−15	2.1e−15	1.1e−15	1.7e−15	pass
$gg \to t\bar{t}$	spin-colour	12 · 2	1.7e−15	6.3e−16	1.3e−14	2.4e−15	1.4e−15	pass
$u\bar{u} \to d\bar{d}g$	plain Born	2 · 2	4.8e−15	2.0e−14	1.2e−10	4.7e−11	7.4e−12	pass
$u\bar{u} \to d\bar{d}g$	colour	40 · 2	1.3e−14	6.5e−14	4.1e−10	1.8e−10	3.4e−11	pass
$u\bar{u} \to d\bar{d}g$	spin	2 · 2	5.1e−15	1.9e−14	9.7e−11	3.9e−11	6.5e−12	pass
$u\bar{u} \to d\bar{d}g$	spin-colour	8 · 2	1.2e−14	2.8e−14	3.5e−10	1.7e−10	3.4e−11	pass
$gg \to ggg$	plain Born	2 · 2	5.0e−15	1.7e−14	1.9e−12	7.7e−13	4.0e−13	pass
$gg \to ggg$	colour	40 · 2	9.5e−15	2.7e−14	1.0e−10	4.0e−11	6.1e−12	pass
$gg \to ggg$	spin	10 · 2	7.4e−15	2.3e−14	3.1e−12	1.5e−12	5.2e−13	pass
$gg \to ggg$	spin-colour	40 · 2	1.3e−14	2.6e−14	1.1e−10	3.9e−11	7.0e−12	pass

4 Electroweak capability matrix (M-SM)

Table 8: **Electroweak capability matrix** at 617b59cd, per class of m-sm/PLAN.md §1.1 (campaign of 2026-11-01/02 at the M-FINAL-b pin, M-SM's of 2026-10-31 re-run, on the shipped patched EW UFOs, data/msm.json; legend below the table, rows: Table 30).

Class	Rows	Public route	DD vs ref.	binary64	Catani (DD)	Native inputs	CMS	Status
W-lep lepton-only [virt=QED], lines through the loop, Bhabha, W pairs, tops · up to $2 \to 3$ · qcd_qed_sm, G_μ, G_μ, m_μ	**7** pass 6 partial 1	LL01, self-test (LL03) ✓ bitwise the tools view; 1 against MadLoop QP	**4.6e−13** LL08	**2.3e−12** LL08 · ref. 2.3e−12†	**3.5e−28** LL08 · LL02 not formed	— public route: native	CM01 ✓ same processes	accepted 6 of 7 pass; F7 LL02, SC02 (accepted)
W-DY neutral and charged Drell-Yan, lepton-hadron, + jets, alpha(MZ) twin · up to $2 \to 4$ · qcd_qed_sm, G_μ	**7** pass 7	DY01 ✓ bitwise the tools view; 1 against MadLoop QP	**9.1e−13** DY05	**4.7e−12** DY02 · ref. 4.7e−12	**1.1e−28** DY07	— public route: native	CM02, CM03, CM05 ✓ same processes	pass 7 of 7 pass; F4 DY05, VJ03, VV06 (open)
W-V V/photon + jet, VV, VVV, VH, VBF H · up to $2 \to 3$ · G_μ	**13** pass 13	VJ02, VV02, VJ03 ✓ bitwise the tools view; 2 against MadLoop QP	**5.3e−11** VV07	**9.2e−11** VV07 · ref. 4.2e−9†	**1.1e−26** VV07	— public route: native	—	pass 13 of 13 pass; F4 DY05, VJ03, VV06 (open)
W-top t t~ (joint QCD+QED), + jet, four tops, massive b, t t~ V/H/ll(+g) · up to $2 \to 5$ · G_μ	**11** pass 11	TT01 ✓ bitwise the tools view; 1 against MadLoop QP	**1.4e−11** TT03	**1.9e−10** TV06 · ref. 9.6e−5†	**6.4e−26** TV07 · TT05 ≤ $1{\times}10^{-14}$ · TT06 ≤ $1{\times}10^{-13}$	TT04, TV07 ✓ public route: native	CM04, CM07 ✓ same processes	pass
W-singletop s- and t-channel, tW; massless and massive b · up to $2 \to 2$ · G_μ	**4** pass 4	— no public case	**3.0e−12** ST04 · MadLoop + Recola	**2.3e−11** ST04 · ref. 2.3e−11†	**4.8e−27** ST04	— campaign rows	—	pass
W-jets QCD processes with EW orders, every split order · up to $2 \to 3$ · G_μ	**4** pass 4	JJ01 ✓ bitwise the tools view	**6.9e−13** JJ04	**7.6e−12** JJ04 · ref. 8.3e−12	**2.0e−28** JJ04	— public route: native	—	pass
W-photon final and initial photons, tagged per leg · up to $2 \to 2$ · G_μ	**4** pass 4	VJ03 ✓ bitwise the tools view	**9.7e−15** PI04	**2.9e−13** PI03 · ref. 2.9e−13	**1.0e−29** PI04	— public route: native	CM09 ✓ same processes	pass
W-CMS internal widths > 0, coloured too; tagged photons; loop_sm · up to $2 \to 5$ · G_μ, loop_sm	**9** pass 9	— no public case	**6.9e−13** CM07	**1.8e−10** CM10 · ref. 4.8e−9†	**7.2e−26** CM07 · CM08 no sector (uncorrected-exchange) · CM10 no sector (uncorrected-exchange)	CM07 ✓ campaign rows	9 rows ✓ against MadLoop CMS (Table 9)	pass
W-multi multiparticle labels, no-Born members skipped (Q2) · up to $2 \to 2$ · G_μ	**2** pass 2	— no public case	**4.2e−14** MP01	**1.5e−12** MP01 · ref. 1.5e−12†	**1.0e−29** MP02	— campaign rows	—	pass
W-scheme alpha(MZ), massive mu, CKM, massive b, massive initial states, eps sign · up to $2 \to 4$ · G_μ, G_μ, m_μ, loop_sm	**10** pass 9 partial 1	— no public case	**4.9e−9** SC02	**1.1e−10** SC06 · ref. 3.4e−4†	**1.1e−6** SC03 · SC02 not formed · SC03 ≤ $1{\times}10^{-5}$ · SC05 ≤ $1{\times}10^{-12}$	SC06, SC07, SC08, SC09 ✓ campaign rows	CM03 ✓ same processes	accepted 9 of 10 pass; F7 LL02, SC02 (accepted)

Rows: campaign rows and their status (passing every rule of their tier; partial: a rule not formed); class status **accepted**: every other row is a finding the coordinator accepted. Models: G_μ = loop_qcd_qed_sm_Gmu, qcd_qed_sm = loop_qcd_qed_sm (both shipped: UFOmodeldb 6ca4b120 + the user's issue-1 patch), G_μ, m_μ its massive-muon copy. **Public route**: the campaign's cases of test_loop_v2_ew_public_models.py (Generator on the shipped UFO with its shipped profile, Runtime.load, evaluate; Born and virtual bitwise the tools view at two points, and = MadLoop's quadruple precision where the case names a fixture) at the row of their process, and the wheel self-test's EW virtual from the shipped prepared artifact. **DD vs ref.**: largest relative deviation of the double-double Laurent coefficients (Born, finite, poles, every sector) from the row's reference (rule 1; MadLoop in quadruple precision, "double": its double-precision control exports). **binary64**: from DD where binary64 is asserted (rule 2, estimate $E \leq 10^{-9}$); "ref." from the reference at every point (information; † the class has points with $E > 10^{-9}$, where DD is asserted). **Catani**: DD poles against the Catani operator (rule 4, every point and predicted sector; "≤" a tolerance relaxed by the user's ruling of 2026-10-29 ~10:20; "no sector": no predicted sector, typed reason in Table 9). **Native inputs**: the rows the campaign ran on the native per-point inputs (production route); the public route is native. **CMS**: the complex-mass rows of the class's processes.

Reading the matrix. The campaign ran every row of its tier at the pin 617b59cd (native loop-native-v2-slice-b36af289) as whole test files, one process each: 23 nightly file jobs, four restriction-smoke files, the default tier and the campaign tier (TV07, its QCD twin and CM07 through the harness CLI), the manifest's points in binary64 and double-double (rules 1–4; rule 5 and the extras on the rows that name them; m-sm/PLAN.md §1.4 under the user's ruling of 2026-10-29 ~10:20: a deviation below ~10^{-5} is not a finding). Its counts: 69 pass, 2 partial, 0 fail, 0 pending, 0 not run; the 13 smokes and multiparticle expansions all pass. The two partial rows are the accepted F7: rule 4 is not formed on the massive-muon G_μ copy, which writes the field renormalisation of massless leptons (rule 1 = MadLoop). Notes of the campaign (data/msm.json coverage_notes): rule 4 is not formed on LL02 and SC02 (F7): a model property, the rows report partial; the Ward extra is a diagnostic under the complex-mass scheme (CM05, CM07, CM09; PLAN 1.4/1.5) and not run on VV05 and ST03 in the real-mass scheme (F3); rule 7 (nc) is not run: EW LC/NLC is never supported (USER Q4, 2026-10-30 ~08:10; typed colour-declaration-coverage); a row whose evaluation raises records the error as its one failed check (the F1 and F2 rows did until the fixes: 52eaff60, and 19619e62 / ms-ufo-patch's decided_number); on the shipped patched models (MS-UFO-PATCH) MadLoop's exports call every R2 coupling the engine inserts (the R2 rule finds no omission; on the unpatched models R2_G0tt_b on LL10, TT04, SC09, TT05).

Table 9: **Complex-mass scheme against MadLoop CMS** (data/msm-matrix.json cms, from data/msm.json and the manifest): CM01–CM05 and CM07–CM10 at 617b59cd (CM06 was cut: its policy is an MS-CMS refusal test). Widths fixed at generation (GeV; "card widths": the restriction's), the per-point inputs, the points and checks (variants included), the largest deviation of double-double from MadLoop CMS (rule 1; quadruple precision, CM01 against its double-precision exports), binary64 from DD (rule 2), the poles against Catani in DD (rule 4) and the Ward identity in DD, a diagnostic only under CMS (m-sm/PLAN.md §1.4, §1.5). CM07 is the acceptance row ([U4]).

Row	Process	Model	Widths	Inputs	Pts · checks	DD	b64/DD	Catani	Ward (DD)	Variants · notes · status
CM01	$e^-e^+ \to m^-m^+$	G_μ	card widths	python	4 · 278	1.2e−13	1.0e−13	2.0e−30	—	zero, small, source pass
CM02	$u\bar{u} \to e^-e^+$	G_μ	WT 1.37758 WZ 2.49877 WW 2.09291 WH 0.00407468	python	4 · 134	8.8e−15	1.5e−13	2.2e−29	—	pass
CM03	$u\bar{d} \to e^+\,\mathrm{ve}$	G_μ	WT 1.37758 WZ 2.49877 WW 2.09291 WH 0.00407468	python	4 · 134	2.2e−14	3.9e−12	3.1e−30	—	pass
CM04	$gg \to t\bar{t}$	G_μ	WT 0 WZ 2.49877 WW 2.09291 WH 0.00407468	python	4 · 230	1.3e−15	1.5e−14	9.1e−31	—	pass
CM05	$u\bar{u} \to e^-e^+g$	G_μ	WT 1.37758 WZ 2.49877 WW 2.09291 WH 0.00407468	python	4 · 142	1.1e−13	5.1e−13	4.8e−30	3.7e−31	pass
CM07 U4	$gg \to t\bar{t}e^-e^+g$	G_μ	WT 0 WZ 2.49877 WW 2.09291 WH 0.00407468	native	3 · 86	6.9e−13	5.6e−11	7.2e−26	1.2e−3	pass
CM08	$u\bar{d} \to e^+\,\mathrm{ve}\,b\bar{b}$	G_μ	WT 1.37758 WZ 2.49877 WW 2.09291 WH 0.00407468	python	4 · 66	1.3e−14	1.8e−13	—	—	QCD2 QED8 unpredicted (uncorrected-exchange) pass
CM09	$u\bar{d} \to e^+\,\mathrm{ve}\,a$	G_μ a0	WT 1.37758 WZ 2.49877 WW 2.09291 WH 0.00407468	python	4 · 142	2.7e−14	1.6e−12	2.4e−29	2.9e−30	pass
CM10	$u\bar{d} \to w^+b\bar{b}$	loop_sm no_b_mass	WT 1.4915 WZ 0 WW 0 WH 0	python	4 · 66	1.6e−13	1.8e−10	—	—	QCD2 QED6 unpredicted (uncorrected-exchange) pass

G1, electroweak row (data/msm-g1.json, the gate tool on the EW records; gated since the final G1 retime, M-LIT-a0, after its indicative timing at M-SM, [U5]). $gg \to t\bar{t}e^-e^+g$ (EW) on loop_qcd_qed_sm_Gmu_py3 (the workspace's Python-3 reference copy of loop_qcd_qed_sm_Gmu, the harness default; not the shipped patched UFO, which the campaign used), native per-point inputs, full colour, binary64 finite part, checks off, one pinned core (184), at $B = 128$: **131 ms per point** (±0.7 %) against Recola 987 ms, OpenLoops 1.40 s and MadLoop 7.07 s (provider records reused, not re-run): ×0.133 of Recola's time, 0.102 of 1.30 × Recola, gate **pass** (0.102 of the binding limit) (Table 1 row, Figure 1). Generation 212 s + 63.2 s, peak 9.24 GiB. Accuracy: against its fixture (3 points) 2.0×10^{-10} (DD asserted at point 2, where $E > 10^{-9}$; binary64 4.7×10^{-10}); against the providers at 16 stream points in binary64: ML 2.8×10^{-9}, OL 2.7×10^{-10}. Earlier run (MS-VALIDATE, 71e2041e): $B = 128$ 153 ms.

5 QCD capability matrix (M-QCD)

Table 10: **QCD capability matrix** at 617b59cd, per class of m-qcd/PLAN.md §1.1 (campaign of 2026-11-02, data/mqcd.json; legend below the table, rows: Table 31).

Class	Rows	Public route	DD vs ref.	binary64	Catani (DD)	LC/NLC	FFT	Status
Q-massless gluons and massless quark lines, every initial state · up to $2 \to 5$	**22** pass 21 fail 1	A12 ✓ DD 3.8e−14 · 3 pts	**1.2e−11** A18	**1.1e−10** A18 · ref. 7.3e−7†	**1.2e−22** A09	rule 7 ✓ 8 rows · smoke pass	**×0.46-5.92** seed · 9 rows; total ×0.93-1.29	accepted 21 of 22 pass; F7 A09@collinear (accepted)
Q-top t t~ + jets, t t~ with quark lines, four tops · up to $2 \to 4$	**6** pass 6	M01@P0, M01@P2 ✓ DD 2.0e−15 · 8 pts	**3.4e−13** B03	**1.2e−10** B01 · ref. 1.7e−7†	**3.0e−26** B01	rule 7 ✓ 1 row · smoke pass	**×1.13** seed · 1 row; total ×0.96	pass
Q-heavy massive b (4FS) and c, b-Yukawa Higgs · up to $2 \to 4$	**10** pass 10	C01 ✓ DD 9.3e−16 · 4 pts	**5.0e−14** C02	**2.3e−10** C05 · ref. 2.3e−10	**1.1e−28** C09	rule 7 ✓ 1 row · smoke pass	auto no cost row	pass
Q-EWBorn QCD corrections to Borns with QED >= 1 · up to $2 \to 5$	**25** pass 25	D05 ✓ DD 2.3e−14 · 4 pts	**1.8e−12** D28	**1.4e−11** D28 · ref. 3.1e−11†	**7.3e−28** D29 · D19 no sector	rule 7 ✓ 1 row · smoke pass	auto no cost row	pass
Q-ttX t t~ + H/Z, t t~ + leptons, e+ e- > t t~ (h) · up to $2 \to 5$	**4** pass 4	E01 ✓ DD 3.6e−13 · 4 pts	**3.6e−13** E01	**7.4e−12** E01 · ref. 5.7e−9†	**5.7e−27** E06	rule 7 ✓ 1 row · smoke pass	auto no cost row	pass
Q-singletop s- and t-channel, tW, 4FS t-channel · up to $2 \to 3$	**5** pass 5	F03 ✓ DD 1.7e−15 · 4 pts	**8.0e−15** C07	**1.6e−14** C07 · ref. 1.6e−14	**3.6e−30** F04 · C07 no sector	rule 7 ✓ 1 row · smoke pass	auto no cost row	pass
Q-mixed several Born orders, every squared split order · up to $2 \to 2$	**2** pass 2	G02 ✓ DD 2.0e−15 · 4 pts	**4.6e−15** G01	**3.1e−14** G02 · ref. 2.9e−14	**3.7e−30** G02	smoke pass no rule-7 row	auto no cost row	pass
Q-CKM non-diagonal CKM on W lines · up to $2 \to 2$	**3** pass 3	H01 ✓ DD 1.4e−15 · 4 pts	**1.4e−15** H01	**7.2e−14** H01 · ref. 7.3e−14	**4.7e−29** H01	smoke pass no rule-7 row	auto no cost row	pass
Q-multi multiparticle labels (p p > t t~) · up to $2 \to 2$	**1** pass 1	M01@P0, M01@P2 ✓ DD 2.0e−15 · 8 pts	**2.0e−15** M01	**4.1e−14** M01 · ref. 4.3e−14	**4.0e−30** M01	smoke pass no rule-7 row	auto no cost row	pass

Rows: campaign rows and their status (passing every rule of their tier; partial: a rule not formed); class status **accepted**: every other row is an open finding the coordinator accepted for the gate. **Public route** (at e18a024b): the class representative generated by Generator on the shipped loop_sm, Runtime.load and evaluate, against MadLoop (rules 1–3; Q-top and Q-multi by the campaign's public variants of $pp \to t\bar{t}$). **DD vs ref.**: largest relative deviation of the double-double Laurent coefficients (Born, finite, poles, every sector) from the binding reference (rule 1; MadLoop in quadruple precision unless named). **binary64**: from DD where binary64 is asserted (rule 2, estimate $E \le 10^{-9}$); "ref." from the reference at every point (information; † the class has points with $E > 10^{-9}$, where DD is asserted). **Catani**: DD poles against the Catani operator (rule 4, every point and predicted sector). **LC/NLC**: rule 7 (formal-N_c projection = N_c scan, and MQ-NC's independent oracles) and the smoke on every shipped restriction. **FFT**: seed contraction FFT ÷ direct of the class's cost rows (Table 11); auto runs the certified FFT wherever the plan has a certified pair.

Reading the matrix. The campaign ran every row of its tier at the pin as whole test files, one process each, the manifest's points in binary64 and double-double (rules 1–4; rules 5–7 and the extras on the rows that name them; m-qcd/PLAN.md §1.4 under the user's ruling of 2026-10-29 ~10:20: a deviation below ~10^{-5} is not a finding). Its counts: 77 pass, 0 partial, 1 fail, 0 pending, 0 not run. The two non-passing rows are accepted limitations, not wrong values: A24 ($gg \to ggggg$, F11) forms no rule 4 because the seven-gluon colour-correlated trees exceed addendum H's memory budget, and its poles equal NJet's double-double values (rule 1); A09 ($u\bar{u} \to ggg$, F7) fails only through its collinear variant, where the executor returns typed failures and no value (binary64 Unstable at $\theta = 10^{-3}$–10^{-5}, DD CutFailed{DegenerateGram} at 10^{-4}, 10^{-5}); A09 itself and its soft variant pass. D19 and C07 have no predicted pole sector (MQ-EWJOIN's separated-exchange rule: their QCD loop selection omits IR-divergent loops, as MadLoop's does) and are checked against MadLoop per split order. **Re-run at 617b59cd** (2026-11-02): no status or check count changed; against M-SM's re-run 42 of 1212 values moved, all binary64 (QCD-only rows too; likely the summation order of SymJIT 2.28.3, not proven; every double-double value bitwise; Section 1); the public-route and FFT columns are M-QCD's runs at e18a024b (public route on a scratch manifest copy).

Table 11: **FFT seed contraction against the direct one** (MQ-FFT, color.contraction): time for the batch of 128 G1-stream points (binary64, poles none, checks off, one pinned thread; g g > u u~ d d~ on the six-gluon stream; one paired round, best of two below $2 \to 5$; indicative), FFT / direct and their ratio (green faster, red slower; Table 1's scale), for the whole evaluation and for the executor's seed-contraction stage. The values agree to rounding at stable points; at the cost runs' worst binary64 points (21 point/carrier checks, binary64-unstable: E up to $O(10^2)$) FFT − direct stays within 2.8 % of the acceptance envelope $10(E_{\mathrm{FFT}} + E_{\mathrm{direct}}) + 16uS$ in binary64 and within 1.7 % in DD (relative $\lesssim 1.1\times10^{-19}$); test_loop_v2_seed_fft.py (default, nightly, seven legs) passed at e18a024b.

Process	Class	group	pairs	evaluation, s	ratio	seed contraction, ms	ratio
$gg \to gg$	Q-massless	S_3	2	0.266 / 0.205	**×1.29**	2.00 / 0.339	**×5.92**
$gg \to ggg$	Q-massless	S_4	2	2.58 / 2.76	**×0.93**	10.5 / 7.03	**×1.50**
$gg \to gggg$	Q-massless	S_5	2	49.3 / 52.1	**×0.95**	180 / 255	**×0.70**
$gg \to ggggg$	Q-massless	S_6	2	1275 / 1246	**×1.02**	8 160 / 17 700	**×0.46**
$u\bar{u} \to ggg$	Q-massless	S_3	19	0.800 / 0.825	**×0.97**	11.4 / 6.73	**×1.69**
$u\bar{u} \to gggg$	Q-massless	S_4	40	10.6 / 10.2	**×1.04**	357 / 312	**×1.15**
$u\bar{u} \to ggggg$	Q-massless	S_5	74	240 / 247	**×0.97**	24 300 / 34 700	**×0.70**
$gg \to t\bar{t}gg$	Q-top	S_4	40	13.2 / 13.7	**×0.96**	224 / 198	**×1.13**
$gg \to u\bar{u}d\bar{d}$	Q-massless	S_2	2340	4.95 / 4.55	**×1.09**	260 / 89.8	**×2.89**
$u\bar{u} \to d\bar{d}g$	Q-massless · no certified pair: direct	—	—	0.519 / 0.420	**×1.24**	2.88 / 2.88	**×1.00**

6 Gate G1: end-to-end timings of the v2 engine

G1 times the complete Born–virtual evaluation, **finite part only**, through the call of the public v2 runtime (`NativeExecutor.evaluate`, which `V2VirtualRuntime` calls) with stability checks off, on one pinned core and one thread, binary64, 16 SIMD lanes, in batches of B distinct points of a frozen stream (REQUIREMENTS_ADDENDUM D, plan §5.8); pyAmpliCol receives the stream's exact rational points, the providers their binary64 decimal text. The configurations:

- **native inputs, on `loop`** — the engine at `ca602436` (milestone M-LIT-a0, the final G1 retime: M-FINAL's engine `c8885cc5` with the 1.0.2 release and its CI fixes, Section 1; build `loop-native-v2-slice-ca602436`, digest `4cbb549f`): the executor forms the trees, the Born and the counterterm inputs per point (binary64 SymJIT plane applications, tree closures in lanes of points), skips the counterterm interactions that only feed poles in finite-only mode, then runs the loop stage. $gg \to ggggg$ is timed from its stored process. The public path ($u\bar{u} \to d\bar{d}g$, `public_path_check` in `data/g1.json`): u u~ > d d~ g, public Generator artifact at ca602436: point_inputs native-v2, plan sha256 9aa235a3 equal to the harness's (and NOT equal to 71e2041e's 64c32c44: the plan format changed since); Runtime.load in a fresh interpreter under the GenerationSentinel: V2VirtualRuntime, native process loaded, executor inputs native, generation calls none, load 0.84 s; the loaded executor's binary64 values equal the timed run's bitwise at 128 of 128 points; public Runtime.evaluate (poles none, core 218) 9.16 ms/pt at B = 1 and 1.73 ms/pt at B = 128 (the campaign's harness run: 1.72 ms/pt on core 136). **Values against the previous campaign** (`values_note`): binary64 values of ca602436 against 71e2041e (m1031-msm's deciding records) at every timed point: bitwise in 0 of 30 configurations; changed: g g > g g g g g FC 128/128 points, median 1.0e-09, max 6.6e+01; g g > g g g g g LC 128/128 points, median 5.1e-10, max 1.9e+01; u u~ > g g g g g LC 128/128 points, median 5.1e-13, max 7.5e-03; u u~ > g g g g g FC 128/128 points, median 1.3e-12, max 3.4e-03; g g > g g g g LC 128/128 points, median 3.0e-11, max 1.2e-03; g g > g g g g FC 128/128 points, median 4.3e-11, max 5.6e-04; g g > g g g LC 128/128 points, median 2.0e-13, max 4.3e-04; g g > g g g FC 128/128 points, median 3.1e-13, max 1.6e-04; u u~ > e- e+ g g g LC 128/128 points, median 6.7e-14, max 8.7e-06; u u~ > e- e+ g g g FC 128/128 points, median 2.6e-13, max 5.9e-06; u u~ > u u~ g FC 128/128 points, median 2.9e-14, max 2.6e-06; u u~ > u u~ g LC 126/128 points, median 2.1e-14, max 2.2e-06; u u~ > e- e+ g g LC 128/128 points, median 1.5e-14, max 1.2e-08; u u~ > e- e+ g g FC 128/128 points, median 3.9e-14, max 7.5e-09; g g > t t~ g g FC 128/128 points, median 4.2e-13, max 6.6e-09; g g > t t~ g g LC 128/128 points, median 2.0e-13, max 5.1e-09; u u~ > g g g g LC 128/128 points, median 1.9e-13, max 2.9e-09; g g > t t~ e- e+ g FC 128/128 points, median 2.1e-12, max 2.1e-09; u u~ > d d~ g LC 128/128 points, median 1.9e-14, max 1.2e-09; g g > t t~ e- e+ g LC 128/128 points, median 4.8e-13, max 9.9e-10; u u~ > g g g g FC 128/128 points, median 3.6e-13, max 8.4e-10; u u~ > e- e+ d d~ g LC 128/128 points, median 7.6e-14, max 7.7e-10; u u~ > g g g FC 128/128 points, median 2.4e-14, max 7.2e-10; u u~ > e- e+ d d~ g FC 128/128 points, median 1.2e-13, max 4.1e-10; u u~ > g g g LC 128/128 points, median 2.8e-14, max 2.5e-10; u u~ > d d~ g FC 128/128 points, median 2.2e-14, max 2.4e-10; u u~ > e- e+ g FC 128/128 points, median 5.0e-15, max 1.2e-11; g g > g g FC 124/128 points, median 6.8e-15, max 1.2e-11; u u~ > e- e+ g LC 128/128 points, median 5.3e-15, max 1.6e-12; g g > g g LC 126/128 points, median 5.7e-15, max 3.3e-13. Changes above 1e-5: g g > g g g, g g > g g g, g g > g g g g, g g > g g g g, g g > g g g g g, g g > g g g g g g, u u~ > g g g g g, u u~ > g g g g g (binary64-unstable points, see the stability rows). Binary64-unstable configurations moved (lita0-gg_ggg-full.json, lita0-gg_gggg-full.json, lita0-gg_ggggg-full.json, lita0-uux_ggggg-full.json): the stability rows of e18a024b at their unstable points (m1029-x*, m1029-p66, estimate E >= 1e-2) stay inputs. Stability rows re-run here against 71e2041e's at the same point and arithmetic: finite parts bitwise at 24 of 65 (binary64 and double-double); the others (g g > g g double, g g > g g g double, g g > g g g g double, g g > g g g g g double, g g > t t~ e- e+ g double, g g > t t~ g g double, u u~ > d d~ g double, u u~ > e- e+ d d~ g double, u u~ > e- e+ g double, u u~ > e- e+ g g double, u u~ > e- e+ g g g double, u u~ > g g g double, u u~ > g g g g double, u u~ > g g g g g double, u u~ > u u~ g double) within 1.4e-06. Points moved by more than 1e-5: 43 (in 8 configurations), 43 within the binary64 stability estimate E recorded at that point for the process (this campaign's or an earlier campaign's stability rows, lita0-x*-stability-* added for the points without one). The timing runs evaluate every point with checks off. Earlier campaigns (M-SM `71e2041e`, M-QCD `e18a024b`, m1028a `28a16255`, m1026a `11542fd5`, m1025a `1d27426b`, G-int `4e9b5a9e`, the development campaign of `39824f9b`) are `earlier_runs`, history;
- **Python inputs** history — the engine on `loop` before G-int, G1 at `1c971289` (LC at `d591a2f7`): trees, Born amplitudes and counterterm scalars formed per point in Python, then the loop stage natively.

Timing protocol at `ca602436` (`timing_protocol` in `data/g1.json`): binary64, finite part only (poles none), no stability check, one thread on one pinned core (taskset, RAYON_NUM_THREADS=1), setup heap frozen (Run.freeze_heap), warm-up excluded, the frozen streams; B = 128 only (user 2026-10-28): 3 blocks of 128 points (2->5: after a 16-point warm-up), full and leading colour. Every configuration re-timed (the native digest changed: f8d77c97 -> 4cbb549f),

every G1 loop artifact regenerated (MODEL_COMPILER_VERSION 17); B = 1 and the 2->5 rows' B = 16 not re-timed: their 28a16255 (m1028a) timings stay in each configuration's earlier_runs (revision 28a16255). The baseline of M-LIT's no-regression A/Bs (lit-tirust step 3, lean-gate item 4b). Host: AMD EPYC 9754, shared host, cores 128-218 without SMT siblings: the final G1 retime (lit-retime, M-LIT-a0) on 2026-11-06 (09:03-10:24 UTC), at most one job per CCX of 8 cores (CCXs 128-215), each on the least busy core of a CCX whose least busy core was below 50 % busy (cores per record: host_cpus); host load average 104-153 (other users' and other packages' processes); no quiet window (m-final/PLAN.md §2.4), one round, no repeats for contention. **Against the previous campaign:** per point at B = 128 against M-SM's G1 (71e2041e): full colour x0.69-1.12 (median x0.84, 15 rows), leading colour x0.67-1.00 (median x0.88, 15 rows); g g > g g g g g g full colour 7.26 -> 4.99 s.

Generation (`generation_protocol`): every row generated cold in its own process (generation in the timing process; g g > g g g g g g in its storing process), fresh PYAMPLICOL_CACHE_DIR and compiled-model directory per job; stages model / generate / export with their peak RSS and plan_mb, as in m1031-msm. Against 71e2041e model + generate + export changed by -54 % to -2 % (median -19 %) on the 28 generated configurations; plan structure (families, lanes, kernels: compare.py) equal everywhere, plan sizes x0.06-0.21 (the binary plan member is the stored form since mf-8g-1); g g > g g g g g: FC generate 44 s + export 89 s, peak 4.4 GiB, plan c1d2ce4a (differs from 71e2041e's; there 171 + 213 s); LC generate 60 s + export 64 s, peak 4.1 GiB, plan 286a85b3 (differs from 71e2041e's; there 178 + 213 s). The provider records are reused unchanged (`provider_records`). The electroweak row is timed by the same protocol in `data/msm-g1.json` and gated (Section 4); the tables of this section are `data/g1.json`'s.

Table 12: **Gate G1, full colour**, ms per point (gate statistic: median), finite part, one thread. Rows $B = 1$ and $B = 128$ (the Python-input $2 \to 4$ rows only at $B = 1$: their Python inputs took seconds per point). Multiplier and colour: time ÷ the fastest provider of the row (legend of Table 1). Grey rows: generation — pyAmpliCol generation plus plan export with its peak memory; MadLoop graph generation, export and compilation · the share of its loop-numerator coefficients in its time (Table 14); Recola process generation; OpenLoops: the prebuilt library and its load plus first call. MadLoop $u\bar{u} \to gggg$ was not exported: the $gg \to gggg$ export took 2.4 h here (HANDOFF §0.1), and OpenLoops binds the gate. Gate: upper line Python inputs (history), lower line native inputs (`loop`); the number is the time over the binding limit (MadLoop's time or 1.30× the faster of Recola and OpenLoops); **open**: the band contains the limit; **not gated**: the configuration runs below 5 ms per point at $B = 128$ (user rule of 2026-10-24; ratio as measured); **reported**: the verdict is taken at $B = 128$ only, so the $B = 1$ ratios are measured and shown without a verdict (user 2026-10-28; bands: Table 13); ° native $B = 1$: the run of `28a16255`, not re-timed at `ca602436` (uncoloured, history; its ratio over this campaign's limit).

Process	B	Python inputs^{hist}	Native inputs, loop	MadLoop	Recola	OpenLoops	Gate
$gg \to gg$ $2 \to 2$, no gate	1	**140** ×69	15.0° ×7.4	**2.02** fastest	**10.0** ×5.0	**3.20** ×1.6	no gate
	128	**132** ×66	**1.29** ×0.64	**2.02** fastest	**10.0** ×5.0	**3.20** ×1.6	no gate
	gen.	10.4 s · 0.3 GiB	5.70 s · 0.6 GiB	9.70 s · 64 %	0.142 s	ppjj · 6.80 ms	
$u\bar{u} \to d\bar{d}g$ $2 \to 3$	1	**40.6** ×27	10.5° ×7.0	**3.17** ×2.1	**3.18** ×2.1	**1.52** fastest	reported 20.61 reported 5.35°
	128	**32.4** ×21	**1.72** ×1.1	**3.17** ×2.1	**3.18** ×2.1	**1.52** fastest	**fail** ×16 **not gated** 0.87
	gen.	9.94 s · 0.3 GiB	6.73 s · 0.6 GiB	9.87 s · 73 %	0.406 s	ppjjj · 4.73 ms	
$u\bar{u} \to u\bar{u}g$ $2 \to 3$	1	**68.9** ×24	18.5° ×6.3	**8.59** ×2.9	**5.55** ×1.9	**2.92** fastest	reported 18.12 reported 4.87°
	128	**54.5** ×19	**2.64** ×0.90	**8.59** ×2.9	**5.55** ×1.9	**2.92** fastest	**fail** ×14 **not gated** 0.69
	gen.	10.4 s · 0.3 GiB	7.07 s · 0.6 GiB	10.9 s · 75 %	0.462 s	ppjjj · 8.43 ms	
$u\bar{u} \to e^-e^+g$ $2 \to 3$	1	*not generable*	4.37° ×4.6	**0.948** fastest	**1.08** ×1.1	*no library*	— reported 4.61°
	128	*not generable*	**0.788** ×0.83	**0.948** fastest	**1.08** ×1.1	*no library*	— **not gated** 0.83
	gen.	CT coverage refused	10.1 s · 0.6 GiB	14.0 s · 79 %	0.701 s	—	
$u\bar{u} \to ggg$ $2 \to 3$	1	**215** ×17	48.1° ×3.9	**28.2** ×2.3	**25.3** ×2.0	**12.4** fastest	reported 13.38 reported 2.98°
	128	**179** ×14	**4.34** ×0.35	**28.2** ×2.3	**25.3** ×2.0	**12.4** fastest	**fail** ×11 **not gated** 0.27
	gen.	12.3 s · 0.4 GiB	6.62 s · 0.6 GiB	17.6 s · 76 %	0.634 s	ppjjj · 23.7 ms	
$gg \to ggg$ $2 \to 3$	1	**2 760** ×28	205° ×2.1	**247** ×2.5	**351** ×3.6	**97.0** fastest	reported 21.86 reported 1.62°
	128	**2 560** ×26	**20.1** ×0.21	**247** ×2.5	**351** ×3.6	**97.0** fastest	**fail** ×20 **pass** 0.16
	gen.	27.9 s · 0.6 GiB	8.10 s · 0.6 GiB	47.0 s · 78 %	6.63 s	ppjjj · 159 ms	
$u\bar{u} \to e^-e^+gg$ $2 \to 4$	1	*not generable*	44.8° ×4.4	**27.3** ×2.7	**10.2** fastest	*no library*	— reported 3.37°
	128	*not generable*	**4.21** ×0.41	**27.3** ×2.7	**10.2** fastest	*no library*	— **not gated** 0.32
	gen.	CT coverage refused	9.53 s · 0.6 GiB	24.8 s · 83 %	0.833 s	—	

Process	B	Python inputs[hist]	Native inputs, loop	MadLoop	Recola	OpenLoops	Gate
$gg \to gggg$ $2 \to 4$	1	**101 000** ×36	3 850° ×1.4	**38 700** ×14	**17 600** ×6.2	**2 840** fastest	reported 27.42 reported 1.04°
	128	*not timed*	**307** ×0.11	**38 700** ×14	**17 600** ×6.2	**2 840** fastest	— **pass** 0.08
	gen.	335 s · 4.1 GiB	14.0 s · 0.9 GiB	8 680 s · 51 %	2 130 s	ppjjjj · 4 950 ms	
$gg \to t\bar{t}gg$ $2 \to 4$	1	**7 310** ×15	817° ×1.7	**5 450** ×12	**1 970** ×4.2	**472** fastest	reported 11.92 reported 1.33°
	128	*not timed*	**79.8** ×0.17	**5 450** ×12	**1 970** ×4.2	**472** fastest	— **pass** 0.13
	gen.	62.5 s · 0.9 GiB	9.98 s · 0.6 GiB	218 s · 64 %	61.4 s	ppttjj · 821 ms	
$u\bar{u} \to gggg$ $2 \to 4$	1	**5 760** ×15	801° ×2.1	*not timed*	**908** ×2.4	**374** fastest	reported 11.85 reported 1.65°
	128	*not timed*	**65.3** ×0.17	*not timed*	**908** ×2.4	**374** fastest	— **pass** 0.13
	gen.	56.1 s · 0.7 GiB	9.86 s · 0.7 GiB	—	32.8 s	ppjjjj · 606 ms	
$u\bar{u} \to e^-e^+d\bar{d}g$ $2 \to 5$	1	*not timed*	*not timed*	**132** ×3.7	**35.3** fastest	**51.0** ×1.4	— —
	128	*not timed*	**13.9** ×0.39	**132** ×3.7	**35.3** fastest	**51.0** ×1.4	— **pass** 0.30
	gen.	—	12.1 s · 0.6 GiB	58.6 s	6.73 s	ppljjj · 108 ms	
$u\bar{u} \to e^-e^+ggg$ $2 \to 5$	1	*not timed*	*not timed*	**1 420** ×5.7	**251** fastest	**346** ×1.4	— —
	128	*not timed*	**47.9** ×0.19	**1 420** ×5.7	**251** fastest	**346** ×1.4	— **pass** 0.15
	gen.	—	11.1 s · 0.6 GiB	146 s	13.4 s	ppljjj · 727 ms	
$gg \to t\bar{t}e^-e^+g$ $2 \to 5$	1	*not timed*	*not timed*	**3 040** ×7.5	**620** ×1.5	**404** fastest	— —
	128	*not timed*	**51.1** ×0.13	**3 040** ×7.5	**620** ×1.5	**404** fastest	— **pass** 0.10
	gen.	—	11.8 s · 0.6 GiB	165 s	27.7 s	pplltj · 1 060 ms	
$u\bar{u} \to ggggg$ $2 \to 5$	1	*not timed*	*not timed*	**569 000** ≥×3.1	**≤ 182 000** (outside protocol, poles on: upper bound)	*no library*	— —
	128	*not timed*	**1 310** ≥×0.0072	**569 000** ≥×3.1	**≤ 182 000** (outside protocol, poles on: upper bound)	*no library*	— **pass** 0.0023
	gen.	—	88.2 s · 2.8 GiB	44 000 s	—	—	
$gg \to ggggg$ $2 \to 5$, no gate (no matched provider)	1	*not timed*	*not timed*	*not exported*	*not generated*	*no library*	no gate
	128	*not timed*	**4 990**	*not exported*	*not generated*	*no library*	no gate
	gen.	—	133 s · 4.4 GiB	—	—	—	

The providers run in production mode on the same stream: OpenLoops with `truepoles 0`, Recola without IR poles (one amplitude call per point), MadLoop at its production stability setting on exports built with MG5′s production flags. OpenLoops has no installed library with QCD corrections to $\ell^+\ell^-$ + 1 or 2 jets (`ppllj_ew` carries electroweak corrections only; existing provider builds only), so the V+jets gates bind on MadLoop and Recola. Providers have no batch API: one time per point serves every B. **Gate statistic:** the median per point on both sides — over pyAmpliCol's timed blocks of one B, and over the provider's timed points (the stream's first 16; 8 for MadLoop $gg \to gggg$, Recola $gg \to gggg$, OpenLoops $gg \to gggg$, Recola $u\bar{u} \to gggg$, MadLoop $u\bar{u} \to ggggg$ and Recola $u\bar{u} \to ggggg$). The provider times are the records of the measure-g1 campaign (23–24 October, on cores 129–345), not re-run since, and for the $2 \to 5$ rows the production-mode runs of package `measure-g1-7leg` (25 October, each provider with its production Collier cache: on for Recola's COLLIER 1.2.8 and OpenLoops, off for MadLoop's COLLIER 1.2.9 with its seven-leg cache defect; $u\bar{u} \to ggggg$ imported from the reference campaign's three points); pyAmpliCol ran on 21 cores between 132 and 215 (per run: `host_cpus` in `data/g1.json`) on 2026-11-06, at most one job per eight-core group sharing an L3 cache, each on the least busy core of a group whose least busy core was below 50 % busy, while the milestone's whole-suite gate and other users' processes loaded the host (`host` in `data/g1.json`: host load average 104-153 (other users' and other packages' processes)). The three $B = 128$ blocks of a process spread by at most 2.7 % (half the range over the median; $u\bar{u} \to d\bar{d}g$ 1.712, 1.720, 1.719 ms). The gate's uncertainty is pyAmpliCol's block spread (half the range of three blocks, the median's standard error from five) plus the provider's half-range over its repeated runs (8 %, the campaign's largest, for the single runs of Recola $gg \to gggg$, Recola $u\bar{u} \to gggg$, OpenLoops $u\bar{u} \to e^-e^+d\bar{d}g$, MadLoop $u\bar{u} \to e^-e^+d\bar{d}g$, OpenLoops $u\bar{u} \to e^-e^+ggg$, Recola $u\bar{u} \to e^-e^+ggg$, MadLoop $u\bar{u} \to e^-e^+ggg$, OpenLoops $gg \to t\bar{t}e^-e^+g$, Recola $gg \to t\bar{t}e^-e^+g$, MadLoop $gg \to t\bar{t}e^-e^+g$ and MadLoop $u\bar{u} \to ggggg$). A gate passes or fails only if the whole band does; otherwise it is **open** and the configuration is repeated once (plan §5.8 item 10). Every timed call returned all its points (failed points: 0); rescue is off by protocol (no checks, no escalation), so the rescue rate is 0.

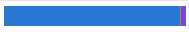
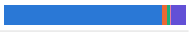
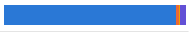
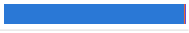
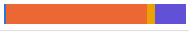
Table 13: **The gate at every batch size** (native inputs on `loop`, full colour): time over the binding limit and its uncertainty band. **The speed gate is taken at $B = 128$ only** (user 2026-10-28, MASTER_PLAN §5.8; `data/milestone.json` `gate_decisions.gate_batch`): only

that column has a verdict and a colour (green pass, red fail, amber open, grey not gated: below 5 ms per point at $B = 128$, user rule of 2026-10-24); every other batch size is **reported**: measured, its ratio and band shown in a white cell, no verdict. $B = 1$ is timed with the default 16 lanes (a one-point call pays a 16-lane chunk); the one-lane column is the development campaign's (`39824f9b`, history: the public path always runs 16 lanes in binary64), a configuration without a $B = 128$ timing and therefore without any verdict (gate missing: no B = 128 timing (timed at B = 1); G1 is taken at B = 128 only (user 2026-10-28)). This campaign timed $B = 128$ only (no verdict was open, so no repeat: `repeats` in `data/g1.json`); ° grey: the earlier run of `28a16255`, `39824f9b`, `4e9b5a9e` at that batch size, its median over the same binding limit with its time (no band; an earlier engine under the committed protocol; history).

Process	$B = 1$, **16 lanes**		$B = 1$, **1 lane** (39824f9b)		$B = 8$		$B = 16$		$B = 32$		$B = 128$ **(gate)**	
$u\bar{u} \to d\bar{d}g$	*5.35°*	10.5 ms	2.13	2.10–2.16	*1.46°*	2.87 ms	—		*1.15°*	2.26 ms	**0.87**	0.86–0.88
$u\bar{u} \to u\bar{u}g$	*4.87°*	18.5 ms	1.63	1.59–1.67	*1.16°*	4.40 ms	—		*0.87°*	3.31 ms	**0.69**	0.68–0.71
$u\bar{u} \to e^-e^+g$	*4.61°*	4.37 ms	1.61	1.60–1.62	*1.07°*	1.01 ms	—		*0.85°*	0.805 ms	**0.83**	0.83–0.83
$u\bar{u} \to ggg$	*2.98°*	48.1 ms	0.84	0.81–0.86	*0.68°*	11.0 ms	—		*0.74°*	11.9 ms	**0.27**	0.26–0.28
$gg \to ggg$	*1.62°*	205 ms	0.45	0.45–0.45	*0.42°*	52.7 ms	—		*0.34°*	42.9 ms	**0.16**	0.15–0.16
$u\bar{u} \to e^-e^+gg$	*3.37°*	44.8 ms	0.78	0.78–0.78	*0.62°*	8.20 ms	—		*0.42°*	5.54 ms	**0.32**	0.31–0.32
$gg \to gggg$	*1.04°*	3 850 ms	0.31	0.30–0.31	*0.30°*	1 120 ms	—		*0.26°*	977 ms	**0.08**	0.08–0.08
$gg \to t\bar{t}gg$	*1.33°*	817 ms	0.37	0.36–0.37	*0.35°*	215 ms	—		*0.28°*	172 ms	**0.13**	0.13–0.13
$u\bar{u} \to gggg$	*1.65°*	801 ms	0.43	0.43–0.44	*0.47°*	230 ms	—		*0.37°*	178 ms	**0.13**	0.13–0.14
$u\bar{u} \to e^-e^+d\bar{d}g$	—		—		—		*0.34°*	15.4 ms	—		**0.30**	0.28–0.33
$u\bar{u} \to e^-e^+ggg$	—		—		—		*0.18°*	57.2 ms	—		**0.15**	0.13–0.16
$gg \to t\bar{t}e^-e^+g$	—		—		—		*0.12°*	61.2 ms	—		**0.10**	0.09–0.11
$u\bar{u} \to ggggg$	—		—		—		*0.0030°*	1 730 ms	—		**0.0023**	0.0021–0.0025

6.1 Where the time goes

Table 14: **Stage split**, ms per point, full colour, at the largest batch timed (Python-input 2 → 4: $B = 1$). Trees/Born/CT: the per-point inputs, timed in Python before G-int (history: trees, counterterm program, Born and counterterm colour projections, scalars); on `loop` the executor forms them and, since lc-perf (`88657260`), times them itself (its `trees` stage: programs, closures, interferences). Numerator: seeds, fused factors and the dressing of the open-loop objects (the pure numerator evaluation). Scatter: closures into NinjaRS lanes. NinjaRS: the reduction without the masters; masters: the master integrals (canonical arguments, cache, oneloopmaster). Glue: point parameters, carrier conversion, assembly and, where the inputs are timed, the rest of the Python call (all rows of this table). MadLoop (instrumented copy; its points less one, 3 for $gg \to gggg$): Born/counterterm HELAS calls, loop-numerator coefficients, reduction plus masters, residual (stability test, colour sums). Bar: shares, colours as the column heads.

Process		Trees, Born, CT	Numerator	Scatter	NinjaRS	Masters	Glue	Total	Shares
$gg \to gg$	native, `loop`	0.191	0.665	0.0663	0.113	0.0372	0.219	**1.29**	
	Python^{hist}	128	0.851	0.0919	0.119	0.0381	2.96	**132**	
	MadLoop	0.0315	1.27	—	0.585		0.107	**1.99**	
$u\bar{u} \to d\bar{d}g$	native, `loop`	0.0768	0.647	0.397	0.142	0.109	0.342	**1.72**	
	Python^{hist}	28.1	0.749	0.462	0.137	0.111	2.66	**32.4**	
	MadLoop	0.0479	2.35	—	0.683		0.133	**3.22**	
$u\bar{u} \to u\bar{u}g$	native, `loop`	0.126	1.10	0.424	0.254	0.190	0.525	**2.64**	
	Python^{hist}	48.5	1.23	0.479	0.232	0.171	3.99	**54.5**	
	MadLoop	0.137	6.44	—	1.15		0.814	**8.54**	
$u\bar{u} \to e^-e^+g$	native, `loop`	0.0516	0.285	0.129	0.0355	0.0262	0.263	**0.788**	
	MadLoop	0.0184	0.754	—	0.133		0.0460	**0.951**	
$u\bar{u} \to ggg$	native, `loop`	0.389	2.32	0.411	0.456	0.211	0.576	**4.34**	
	Python^{hist}	168	2.92	0.476	0.459	0.225	5.75	**179**	
	MadLoop	0.346	18.5	—	2.54		2.87	**24.3**	
$gg \to ggg$	native, `loop`	5.57	9.88	0.518	2.27	0.524	1.14	**20.1**	
	Python^{hist}	2 530	11.7	0.519	2.03	0.483	10.5	**2 560**	
	MadLoop	2.42	184	—	11.3		39.5	**237**	
$u\bar{u} \to e^-e^+gg$	native, `loop`	0.161	2.26	0.589	0.435	0.237	0.524	**4.21**	
	MadLoop	0.289	22.7	—	2.38		2.11	**27.5**	
$gg \to gggg$	native, `loop`	93.7	158	2.33	38.4	6.05	8.29	**307**	
	Python^{hist}	97 300	3 750	4.08	38.6	5.81	137	**101 000**	
	MadLoop	88.9	18 600	—	209		17 700	**36 500**	
$gg \to t\bar{t}gg$	native, `loop`	11.0	34.9	0.888	18.9	6.63	7.49	**79.8**	
	Python^{hist}	6 490	751	3.55	20.7	6.58	29.3	**7 310**	

Process		Trees, Born, CT	Numerator	Scatter	NinjaRS	Masters	Glue	Total	Shares
	MadLoop	24.2	3 290	—	105		1 740	**5 160**	
$u\bar{u} \to gggg$	native, `loop`	9.37	37.2	1.74	9.04	2.98	4.78	**65.3**	
	Python[hist]	4 910	806	2.01	9.64	2.99	24.1	**5 760**	
$u\bar{u} \to e^-e^+d\bar{d}g$	native, `loop`	0.570	6.62	1.79	1.90	1.27	1.58	**13.9**	
$u\bar{u} \to e^-e^+ggg$	native, `loop`	1.56	30.1	2.05	7.93	2.75	3.53	**47.9**	
$gg \to t\bar{t}e^-e^+g$	native, `loop`	2.88	29.0	0.734	10.2	4.26	4.07	**51.1**	
$u\bar{u} \to ggggg$	native, `loop`	215	804	9.27	186	46.0	46.8	**1 310**	
$gg \to ggggg$	native, `loop`	1 970	2 580	11.0	344	42.0	28.4	**4 990**	

Verdict of gate G1 (processes from $2 \to 3$; $2 \to 2$ without a gate).

- **Native inputs on `loop`** (G1 at `ca602436`) at $B = 128$: **pass** for all 8 gated processes, $gg \to ggg$, $gg \to gggg$, $gg \to t\bar{t}gg$, $u\bar{u} \to gggg$, $u\bar{u} \to e^-e^+d\bar{d}g$, $u\bar{u} \to e^-e^+ggg$, $gg \to t\bar{t}e^-e^+g$ and $u\bar{u} \to ggggg$. **Not gated** (below 5 ms per point, user decision of 2026-10-24; ratios as measured): $u\bar{u} \to d\bar{d}g$ 1.72 ms (0.87 of its limit), $u\bar{u} \to u\bar{u}g$ 2.64 ms (0.69 of its limit), $u\bar{u} \to e^-e^+g$ 0.788 ms (0.83 of its limit), $u\bar{u} \to ggg$ 4.34 ms (0.27 of its limit) and $u\bar{u} \to e^-e^+gg$ 4.21 ms (0.32 of its limit). $B = 1$ (16 lanes) and the $2 \to 5$ rows' $B = 16$ are **reported** without a verdict (user 2026-10-28): $B = 1$ at ×1.0–5.3 of the limit (a one-point call pays a 16-lane chunk, Table 13); with one lane per call, in the development runs of `39824f9b` (history), ×0.31–2.1.
- **Python inputs** history (the engine before G-int, G1 at `1c971289`): failed every process it generated, at $B = 128$ at ×11–20 of the binding limit, the per-point Python inputs taking 87 %–99 % of its time (Table 14); at $B = 1$ at ×12–27, where $gg \to gggg$, $gg \to t\bar{t}gg$ and $u\bar{u} \to gggg$ were timed only. It did not generate $u\bar{u} \to e^-e^+g$ and $u\bar{u} \to e^-e^+gg$ (Furry zeros, WP3.5; Z fermion signs, WP3.2).
- **Fastest of all tools** (aspiration), native inputs at $B = 128$, within the uncertainty band: faster than every provider for $u\bar{u} \to u\bar{u}g$, $u\bar{u} \to e^-e^+g$, $u\bar{u} \to ggg$, $gg \to ggg$, $u\bar{u} \to e^-e^+gg$, $gg \to gggg$, $gg \to t\bar{t}gg$, $u\bar{u} \to gggg$, $u\bar{u} \to e^-e^+d\bar{d}g$, $u\bar{u} \to e^-e^+ggg$ and $gg \to t\bar{t}e^-e^+g$; faster than every provider with a production-mode time for $u\bar{u} \to ggggg$ (MadLoop; Recola has only an upper bound of its production time, ≤ 182 s); at parity with the fastest for none; slower than the fastest for $u\bar{u} \to d\bar{d}g$ (OpenLoops). NJet, not a matched provider, takes 4.59 s (NJet 3.1.2) for its two-evaluation full-colour $gg \to gggg$ test (Table 23).
- Seven-leg rows ($u\bar{u} \to e^-e^+d\bar{d}g$, $u\bar{u} \to e^-e^+ggg$, $gg \to t\bar{t}e^-e^+g$, $u\bar{u} \to ggggg$, $gg \to ggggg$) are generated by the WP3.1 generator, on `loop` since G-int, and timed by G1 since package `measure-g1-7leg` (full and leading colour at $B = 16$ and 128; re-timed at `ca602436`); $gg \to ggggg$ has no matched provider timing (addendum F, H): reported without a gate.

Table 15: **Levers**, measured on the gated rows of Table 14 ($2 \to 3$ and beyond; shares at $B = 128$). Gains are measured where a column says so; otherwise they are the stage's share, an upper bound.

Lever	Evidence	Expected gain	Owner
Native per-point inputs (WP3.1 I3–I6: trees, Born, counterterm scalars in the executor)	Python inputs 87 %–99 % of the Python-input time at $B = 128$ (history)	measured, Python ÷ native inputs at equal B: — at $B = 1$, ×19–127 at $B = 128$ ($2 \to 3$)	done: G-int, on `loop` since `4e9b5a9e`
Binary64 point inputs (the providers' input; the protocol passes the stream's exact rationals)	paired run on the loaded artifact of G-int (`4e9b5a9e`, history), $B = 128$: $u\bar{u} \to d\bar{d}g$ 1.68 against 2.29 ms, $u\bar{u} \to u\bar{u}g$ 2.94 against 3.39 ms	measured: ×1.37 and ×1.15 (information; not a G1 configuration)	none: moot under the 5 ms gate rule (user, 2026-10-24); the protocol keeps exact inputs
Compiled native trees (kernels as the tree engine's SymJIT planes, lanes batched)	native trees (with the Born and counterterm inputs; the executor's timer) 3 %–31 % of the time at $B = 128$; $gg \to gggg$ 31 %	up to the share; ×1.4 for the largest	in part on `loop`: SymJIT plane applications (`7db32901`), lane-chunked tree closures (`88657260`); next: the native tree program (A5, `native-dag`)
Lanes by batch in binary64 (`plan_export.lanes_for`: 16 lanes also for one point)	$B = 1$: one 16-lane chunk per call ($B = 1$ is reported, not gated: user 2026-10-28)	measured: — at $B = 1$ (16 lanes on `loop` ÷ 1 lane of `39824f9b`, different builds)	executor policy (WP3.1 or wave 4)
Finite-only result path (sum the virtual natively, no per-family lane export, Python assembly only for checks and observables)	glue 3 %–33 % of the native-input time; skipping the pole-only counterterm interactions natively (generation-scale-2's I5 record) saved 2–5 % of the total at $u\bar{u} \to gggg$ and $gg \to gggg$	up to the share (largest at $2 \to 3$ with quarks)	in part on `loop` (generation-scale-2); next `plan_export`, executor (wave 4)
Lane-major closures (no scatter, no second copy into the reducer's batch)	scatter 1 %–23 %	up to the share	executor, reducer adapter (wave 4)
AVX2 / AVX-512 kernels (runtime selection)	numerator 36 %–63 %	×1.0–1.4 on the dressing (gate G0)	wave 4 (WP4.6)
NinjaRS and masters: parallel reduction, transposed reduction, stage-2 sharing, a cheaper master cache key	NinjaRS + masters 8 %–32 %; masters 14 %–43 % of the reduction, oneloopmaster alone 1 %–12 %: the cache bookkeeping costs more than the integrals	one thread: up to the share; threads: DESIGN §5.6	NinjaRS requests (DESIGN §11.5)

Below 5 ms per point. $u\bar{u} \to d\bar{d}g$, $u\bar{u} \to u\bar{u}g$, $u\bar{u} \to e^-e^+g$, $u\bar{u} \to ggg$ and $u\bar{u} \to e^-e^+gg$ take 1.72, 2.64, 0.788, 4.34 and 4.21 ms per point at $B = 128$; by the user's decision of 2026-10-24 a process below 5 ms has no blocking gate and its slowdown against the providers need not be pursued: optimisation effort goes to higher multiplicities. Their ratios stay in the tables ($u\bar{u} \to d\bar{d}g$: 1.72 ms against a limit of 1.97 ms, OpenLoops 1.52 ms × 1.30, 0.87 of it; measured pass, band 0.86–0.88). The question whether G1 should pass the points as exact rationals (the protocol) or as the binary64 text the providers read is moot under this rule, and the protocol stays. For information only (not G1 configurations, no verdict formed from them): on `ca602436` the public `Runtime.evaluate` of the loaded artifact (poles off, core 221) takes 11.71 ms per point at $B = 1$ and 2.01 ms at $B = 128$ for $u\bar{u} \to d\bar{d}g$, against the campaign's 1.72 ms of the harness's executor at $B = 128$ (core 147) and 1.98 ms in its smoke run, on other cores and at other times (`public_path_check` in `data/g1.json`, the campaign's `STATE.md`; M-QCD's run, 2.20 ms on a busy host, was 1.83 ms at m1028a); at G-int (`4e9b5a9e`, committed protocol) a paired run (one core, ten alternating rounds) gave 1.68 ms on the binary64 text against 2.29 ms on the exact points (×0.73; $u\bar{u} \to u\bar{u}g$: 2.94 against 3.39 ms, ×0.87).

6.2 Leading colour

Table 16: **Leading colour** (LC), ms per point, same protocol as Table 12; in each cell the full-colour (FC) time of the same configuration and batch, and LC/FC. No provider ratio: the providers time full colour. LC uses the committed QCD colour continuation (`examples/loop/source_colour`); the V+jets rows are timed at LC with native inputs only. This campaign timed LC at $B = 128$; ° grey: the earlier run of `39824f9b` at $B = 1$ with its FC time (an earlier engine under the committed protocol; history).

Process	Python inputs[hist]				Native inputs, loop			
	$B = 1$		largest batch		$B = 1$		largest batch	
$gg \to gg$	*not timed*		*not timed*		*not timed*		**0.596**	FC 1.29 · 0.46
$u\bar{u} \to d\bar{d}g$	**40.5**	FC 40.6 · 1.00	**37.4**	FC 32.4 · 1.15	*5.25°*	FC 10.9 · 0.48	**1.12**	FC 1.72 · 0.65
$u\bar{u} \to u\bar{u}g$	*not timed*		*not timed*		*9.00°*	FC 19.1 · 0.47	**1.59**	FC 2.64 · 0.60
$u\bar{u} \to e^-e^+g$	*not timed*		*not timed*		*not timed*		**0.446**	FC 0.788 · 0.57
$u\bar{u} \to ggg$	**184**	FC 215 · 0.86	**171**	FC 179 · 0.96	*22.8°*	FC 50.6 · 0.45	**2.38**	FC 4.34 · 0.55
$gg \to ggg$	**2 350**	FC 2 760 · 0.85	**2 300** ($B = 32$)	FC 2 550 · 0.90	*91.1°*	FC 211 · 0.43	**7.59**	FC 20.1 · 0.38
$u\bar{u} \to e^-e^+gg$	*not timed*		*not timed*		*not timed*		**1.71**	FC 4.21 · 0.41
$gg \to gggg$	*not timed*		*not timed*		*2 070°*	FC 4 100 · 0.50	**126**	FC 307 · 0.41
$gg \to t\bar{t}gg$	*not timed*		*not timed*		*418°*	FC 840 · 0.50	**28.2**	FC 79.8 · 0.35
$u\bar{u} \to gggg$	*not timed*		*not timed*		*380°*	FC 860 · 0.44	**25.3**	FC 65.3 · 0.39
$u\bar{u} \to e^-e^+d\bar{d}g$	*not timed*		*not timed*		*not timed*		**5.75**	FC 13.9 · 0.41
$u\bar{u} \to e^-e^+ggg$	*not timed*		*not timed*		*not timed*		**9.42**	FC 47.9 · 0.20
$gg \to t\bar{t}e^-e^+g$	*not timed*		*not timed*		*not timed*		**10.1**	FC 51.1 · 0.20
$u\bar{u} \to ggggg$	*not timed*		*not timed*		*not timed*		**438**	FC 1 310 · 0.34
$gg \to ggggg$	*not timed*		*not timed*		*not timed*		**2 680**	FC 4 990 · 0.54

LC is cheaper in the loop stage (fewer families and skeletons) and, since lc-perf (`88657260`: LC-aware pruning of the closure blocks, native tree closures in lanes of points), also end to end: with native inputs at $B = 128$ LC/FC is 0.20–0.65 ($u\bar{u} \to e^-e^+ggg$ gains most, $u\bar{u} \to d\bar{d}g$ least), so LC is faster than full colour on every row; at G-int it was 0.55–1.28, slower than full colour for $u\bar{u} \to ggg$, $gg \to gggg$, $gg \to t\bar{t}gg$ and $u\bar{u} \to gggg$ (committed protocol; Section 6). The native trees, with the Born and counterterm inputs, take a larger share at LC ($u\bar{u} \to ggg$: 12 % at LC, 9 % at full colour; $gg \to gggg$: 36 % at LC, 31 % at full colour), the lever of compiled trees (Table 15). The $2 \to 5$ rows follow: LC/FC 0.20–0.54 at $B = 128$ ($u\bar{u} \to e^-e^+ggg$ 0.20, $gg \to t\bar{t}e^-e^+g$ 0.20, $u\bar{u} \to ggggg$ 0.34, $u\bar{u} \to e^-e^+d\bar{d}g$ 0.41, $gg \to ggggg$ 0.54). At seven gluons the stored LC plan of $gg \to ggggg$ keeps all 4 707 full-colour loop families, where the lower gluon multiplicities halve them (154 ⟶ 77 and 1 100 ⟶ 550 at five and six gluons, with the top). This is not a defect (DESIGN §7.9): a family is keyed by its propagators' momenta and masses, so a closed massless-quark loop shares its family with the gluon and ghost rings, and the prune drops a family only when none of its skeletons is live; with the top the families halve, under this generation's `particle_veto = [t]` they cannot, while the tries shrink alike. The LC plan generates in 1.0 min against 0.7 min at full colour (`LoopSlice.generate` alone; on `1d27426b` 40.6 against 37.9 min).

Identical-flavour quark lines: since identical-flavour-lc (`5cb5ea83`, milestone m1028a) the leading and next-to-leading colour of processes with identical-flavour quark lines (from $u\bar{u} \to u\bar{u}g$ up, any number of lines) are exact formal-N_c projections: the loop-request blocks and the counterterm program split by U(1)-subtraction N_c grade in the Python and Rust lanes, with order-free block grades (second genericity audit, G2-M1; validated against the legacy engine's LC and NLC, now frozen captures, and since MQ-NC by rule 7 against N_c scans and independent oracles, Section 5). G1′s LC of $u\bar{u} \to u\bar{u}g$ is a validated value (1.59 ms per point at $B = 128$).

6.3 Throughput, model recalibration and schedule M

Table 17: **Kernel throughput and the design recipe** (native inputs, full colour, largest batch). Work: the executor's modelled complex multiply-adds of the chosen kernel forms per point (one lane); throughput: work × lanes over the dressing time (fused-factor builds excluded). Native loop: seeds, factors, dressing, scatter and reduction; outside: everything else of the call (trees, glue). Recipe: DESIGN §1.3, finite only, where the design has a row; measured ÷ recipe for the total and for the NinjaRS stage (masters included) against its model R.

Process	Work, MCMAC/pt	Dressing, ms	G CMAC/s	Native loop, ms	Outside, ms	Total, ms	Total ÷ recipe	NinjaRS ÷ R
$gg \to gg$	0.932	0.526	1.77	0.881	0.410	**1.29**	0.8–2.6	0.96–1.09
$u\bar{u} \to d\bar{d}g$	0.556	0.352	1.58	1.30	0.424	**1.72**	no row	—
$u\bar{u} \to u\bar{u}g$	1.09	0.728	1.49	1.97	0.667	**2.64**	no row	—
$u\bar{u} \to e^-e^+g$	0.168	0.124	1.35	0.475	0.313	**0.788**	no row	—
$u\bar{u} \to ggg$	3.18	1.84	1.73	3.39	0.949	**4.34**	no row	—
$gg \to ggg$	14.3	8.91	1.61	13.2	6.87	**20.1**	no row	—
$u\bar{u} \to e^-e^+gg$	2.89	1.71	1.69	3.53	0.687	**4.21**	no row	—
$gg \to gggg$	268	152	1.76	205	102	**307**	0.8–2.2	0.87–1.02
$gg \to t\bar{t}gg$	56.0	32.4	1.73	61.4	18.5	**79.8**	0.5–1.5	0.95–1.33
$u\bar{u} \to gggg$	56.7	33.3	1.70	50.9	14.3	**65.3**	no row	—
$u\bar{u} \to e^-e^+d\bar{d}g$	8.80	5.35	1.65	11.6	2.36	**13.9**	0.8–2.6	0.97–1.74

Process	Work, MCMAC/pt	Dressing, ms	G CMAC/s	Native loop, ms	Outside, ms	Total, ms	Total ÷ recipe	NinjaRS ÷ R
$u\bar{u} \to e^-e^+ggg$	47.8	28.0	1.71	42.9	5.03	**47.9**	0.7–2.0	0.83–1.15
$gg \to t\bar{t}e^-e^+g$	49.3	27.9	1.77	44.2	6.92	**51.1**	0.7–1.9	1.03–1.42
$u\bar{u} \to ggggg$	1 170	677	1.72	1 050	261	**1 310**	0.5–1.6	0.79–1.07
$gg \to ggggg$	4 060	2 470	1.64	2 980	2 010	**4 990**	0.6–1.8	0.38–0.45

Recalibration. The kernels run 1.35–1.77 G complex MAC/s per core on the real plans, the rate gate G0 measured on its synthetic programs (1.29–1.37); the dressing term of the recipe holds. The NinjaRS stage, masters included, lands on the design's model R (measured ÷ R 0.38–1.74, Table 17). What the recipe misses is the **outer factor**: it assumed trees, seeds and counterterms add 5–10 % (DESIGN §1.3), but with the current native inputs the total is 1.18–2.3 times dressing plus reduction at full colour (up to 3 at LC). Table 18 re-estimates every design row with these measured terms.

Table 18: **Design estimates recalibrated by G1** (`tools/loop_v2_costmodel.py --g1`), ms per point, finite only, one thread. Recipe: DESIGN Table 1a's finite-only band (outer factor 1.05–1.10). G1: the recipe's dressing term, the NinjaRS model R × 0.385–1.737 (the band this campaign measures, Table 17) and the measured outer factor of the row's class representative at $B = 128$ with native inputs (same colour where measured). Measured: G1's own time where the row was timed. Gate: G1's binding limit at the row's largest batch where G1 timed the providers (1.3× the fastest production provider, or MadLoop's time where it binds), else 1.3× the best production provider of DESIGN Table 1b (protocol of 20 October), with the headroom gate ÷ upper estimate. The outer factor shrinks with compiled trees (Table 15); until then these replace the Table 1a targets.

Design row	Col.	Recipe	Outer	Class	G1 estimate	Measured	Gate	Headroom
$gg \to gg$	FC	0.502–1.52	1.60	$gg \to gg$	**0.629–2.40**	1.29	—	—
$gg \to gg$	LC	0.204–0.542	2.22	$gg \to gg$	**0.337–1.22**	0.596	—	—
$gg \to gggg$	FC	139–409	1.52	$gg \to gggg$	**161–624**	307	3 690	5.92
$gg \to gggg$	LC	55.7–146	1.70	$gg \to gggg$	**67.3–257**	126	—	—
$gg \to ggggg$	FC	2 710–7 920	1.69	$gg \to ggggg$	**3 470–13 400**	4 990	—	—
$gg \to ggggg$	LC	1 070–2 830	1.47	$gg \to ggggg$	**1 120–4 310**	2 680	—	—
$u\bar{u} \to ggggg$	FC	827–2 630	1.42	$u\bar{u} \to ggggg$	**931–3 700**	1 310	569 000	153.89
$u\bar{u} \to ggggg$	LC	206–526	1.51	$u\bar{u} \to ggggg$	**210–859**	438	—	—
$u\bar{u} \to e^-e^+ggg$	FC	24.5–68.1	1.18	$u\bar{u} \to e^-e^+ggg$	**20.8–84.3**	47.9	326	3.87
$u\bar{u} \to e^-e^+ggg$	LC	3.34–8.31	1.58	$u\bar{u} \to e^-e^+ggg$	**3.59–14.3**	9.42	—	—
$u\bar{u} \to e^-e^+d\bar{d}g$	FC	5.34–16.5	1.43	$u\bar{u} \to e^-e^+d\bar{d}g$	**5.66–24.9**	13.9	45.9	1.84
$u\bar{u} \to e^-e^+d\bar{d}g$	LC	1.19–3.14	1.87	$u\bar{u} \to e^-e^+d\bar{d}g$	**1.53–6.59**	5.75	—	—
$gg \to t\bar{t}gg$	FC	54.8–156	1.35	$gg \to t\bar{t}gg$	**54.5–219**	79.8	613	2.81
$gg \to t\bar{t}gg$	LC	14.6–40.1	1.65	$gg \to t\bar{t}gg$	**17.6–68.8**	28.2	—	—
$gg \to t\bar{t}e^-e^+g$	FC	27.4–76.1	1.18	$gg \to t\bar{t}e^-e^+g$	**23.5–94.1**	51.1	525	5.58
$gg \to t\bar{t}e^-e^+g$	LC	4.06–11.0	1.69	$gg \to t\bar{t}e^-e^+g$	**5.00–19.5**	10.1	—	—

Table 19: **Schedule M against B** (native inputs, one thread), pure-gluon classes, the only classes where M merges seed entries (DESIGN §5.6). The schedule-dependent stages are seeds, factors, dressing, scatter and the NinjaRS reduction ("the reduction included", §1.6 item 5); the trees, Born and counterterm inputs and the glue are the same work in both schedules. M: the development campaign's runs (`39824f9b`, not repeated); B: the configuration on `loop` (`ca602436`), a later engine (native tree closures and plane applications, frozen setup heap) on the same processes, so the ratios pair runs of different builds and campaigns (the gate tool pairs them; history for M): end to end at $B = 128$ they set the M time with the older, slower trees against this engine's B time and do not measure M; LC was re-timed at $B = 128$ only. M is adopted where M/B ≤ 0.85 on the schedule-dependent stages at every common batch size; ± is the timing uncertainty of the ratio.

Process	Col.	Stages M/B, $B = 1$	Stages M/B, $B = 128$	End to end, $B = 1$	End to end, $B = 128$	Adopt M
$gg \to ggg$	FC	—	**1.12** ± 0.05	—	2.53	no
$gg \to ggg$	LC	—	**1.11** ± 0.04	—	6.14	no
$gg \to gggg$	FC	—	**1.15** ± 0.05	—	2.95	no
$gg \to gggg$	LC	—	**1.04** ± 0.02	—	8.07	no

B against M (decision, provisional until the trees are compiled). On the schedule-dependent stages M passes the 15 % rule for no row, and not for $gg \to ggg$ at FC (1.12), $gg \to ggg$ at LC (1.11), $gg \to gggg$ at FC (1.15) and $gg \to gggg$ at LC (1.04). The same rows qualify as in the development campaign's paired runs, on which the decision was made (`data/g1.json` of `c3a7be09`). End to end the gain is diluted by the tree inputs, which both schedules share and which compiled trees will shrink (Table 15). With a quark line M's entries do not merge at all (M/B work 1.00, DESIGN §5.6; equal by test). **Decision (G1):** schedule M for pure-gluon classes at leading colour from $n = 6$; schedule B for every other class at $n \leq 6$, including pure-gluon full colour; M stays the planned default from $n \geq 8$. The seven-leg comparison of M with B is open (the schedule-M run of `measure-g1-7leg` stopped; G1 and Table 1 time schedule B).

Table 20: **Values at the timed points**: pyAmpliCol's binary64 finite part (native inputs on `loop`) against each provider's production-mode value on the common stream points (the providers read the binary64 decimal text of the exact points), largest deviation, metric $|a-b| \,/\, \max(|a|, |b|, |B|\ \alpha_s/2\pi)$; Recola's finite part converted with the universal double pole (it returns no poles in production). Worst point: the point of the row's largest deviation, evaluated again with one rotation and the reduction-direction check in binary64 and in

double-double (DD) by the stability runs of `e18a024b`, not repeated (their binary64 values equal the timed ones to rounding level, more at binary64-unstable points): E is the binary64 estimate, DD the DD value's deviation from the closest provider (named). Campaign: the first three points against the committed references (`tests/fixtures/loop/required`; $gg \to ggggg$: NJet's ten points), fixture metric and tolerance, asserted in binary64 where its estimate $E \le 10^{-9}$ and in DD elsewhere (user 2026-10-29; earlier records: binary64 throughout), native inputs (`loop`) and Python inputs (history); ‡ points above the tolerance that the user accepted (below the table).

Process	Pts	vs MadLoop	vs Recola	vs OpenLoops	Worst point: E	DD: closest provider	Campaign, native	Campaign, Python[nist]
$gg \to gg$	16	1.3×10^{-12}	2.9×10^{-13}	3.0×10^{-13}	pt 4: 3.8×10^{-12}	9.4×10^{-14} OpenLoops	—	—
$u\bar{u} \to d\bar{d}g$	16	4.3×10^{-13}	3.3×10^{-10}	7.9×10^{-13}	pt 1: 8.0×10^{-12}	3.2×10^{-13} MadLoop	3.6×10^{-10}	3.6×10^{-10}
$u\bar{u} \to u\bar{u}g$	16	1.2×10^{-11}	7.2×10^{-11}	1.4×10^{-12}	pt 1: 8.7×10^{-12}	2.4×10^{-13} MadLoop	8.6×10^{-11}	8.6×10^{-11}
$u\bar{u} \to e^-e^+g$	16	7.3×10^{-14}	7.5×10^{-14}	no library	—	—	—	—
$u\bar{u} \to ggg$	16	8.6×10^{-13}	1.5×10^{-12}	8.4×10^{-11}	pt 2: 1.3×10^{-12}	1.3×10^{-14} MadLoop	—	—
$gg \to ggg$	16	3.7×10^{-7}	3.3×10^{-7}	3.3×10^{-7}	pt 10: 1.3×10^{-5}	6.0×10^{-12} Recola	—	—
$u\bar{u} \to e^-e^+gg$	16	2.2×10^{-10}	1.1×10^{-10}	no library	pt 8: 1.7×10^{-9}	4.3×10^{-13} Recola	—	—
$gg \to gggg$	8	8.7×10^{-8}	1.5×10^{-7}	1.5×10^{-7}	pt 1: 1.8×10^{-6}	4.4×10^{-12} Recola	7.1×10^{-11}	8.7×10^{-8} **above 2.0×10^{-8}**
$gg \to t\bar{t}gg$	16	8.8×10^{-9}	2.7×10^{-9}	2.7×10^{-9}	pt 13: 8.9×10^{-8}	5.5×10^{-15} OpenLoops	1.1×10^{-11}	1.1×10^{-11}
$u\bar{u} \to gggg$	16	—	8.0×10^{-12}	2.3×10^{-10}	pt 9: 3.0×10^{-9}	1.9×10^{-13} OpenLoops	—	—
$u\bar{u} \to e^-e^+d\bar{d}g$	16	3.4×10^{-11}	4.5×10^{-11}	2.0×10^{-11}	—	—	4.4×10^{-13}	—
$u\bar{u} \to e^-e^+ggg$	16	2.2×10^{-11}	4.2×10^{-11}	4.3×10^{-11}	pt 8: 3.8×10^{-10}	7.7×10^{-14} OpenLoops	3.5×10^{-13}	—
$gg \to t\bar{t}e^-e^+g$	16	1.1×10^{-8}	6.5×10^{-9}	3.5×10^{-9}	pt 2: 4.3×10^{-8}	5.8×10^{-13} OpenLoops	1.0×10^{-8}	—
$u\bar{u} \to ggggg$	3	2.1×10^{-12}	2.9×10^{-10}	no library	pt 0: 1.7×10^{-11}	1.4×10^{-12} MadLoop	2.9×10^{-10}	—
$gg \to ggggg$	0	not generated	not generated	no library	—	—	7.1×10^{-9}	—

Reading the values. The Born agrees with every provider to 1.7×10^{-13} or better. For the finite part, the double-double (DD) value at each worst point (Table 20) separates the sources: it agrees with the closest provider to 6.0×10^{-12} or better everywhere. Where pyAmpliCol's binary64 estimate E exceeds the deviation — $gg \to ggg$, $gg \to gggg$ and $gg \to t\bar{t}gg$ (above 10^{-9}) and $u\bar{u} \to gggg$ (1.9×10^{-13} in DD against 2.3×10^{-10} in binary64 at point 9) — the deviation is pyAmpliCol's own binary64 loss, which E flags. The other deviations above 10^{-11} stay in DD and are the providers' own accuracy: Recola's at $u\bar{u} \to d\bar{d}g$ (3.3×10^{-10}), $u\bar{u} \to u\bar{u}g$ and $u\bar{u} \to e^-e^+gg$ (COLLIER's default target 10^{-8}, HANDOFF §7), OpenLoops' at $u\bar{u} \to ggg$ (8.4×10^{-11}) and MadLoop's at $u\bar{u} \to e^-e^+gg$ (1.1×10^{-10}); MadLoop's production values also stay off the DD values at the ill-conditioned gluon and top points (up to 2.4×10^{-7}). At the campaign points the check asserts binary64 where its estimate $E \le 10^{-9}$ and double-double elsewhere (user E1 ruling 2026-10-29 ~03:45, package g1-accuracy): $gg \to gggg$ point 1 (E = 1.8×10^{-6}) is asserted in DD, the check passing at 7.1×10^{-11} (tolerance 2×10^{-8} binary64 there 1.5×10^{-7}, recorded); its DD value agrees with Recola to 4.4×10^{-12}. ‡ $gg \to ggggg$ against NJet: points 4 and 8 exceed the tolerance 2×10^{-8} in binary64 (1.5×10^{-7} and 2.6×10^{-8} from NJet's double-double values, at most 7.1×10^{-9} over all its references), **accepted (user 2026-10-25): binary64 within the estimate, DD = NJet**: estimates E 1.5×10^{-6} and 3.4×10^{-7}, and pyAmpliCol's double-double values equal NJet's; the check keeps its tolerance and, asserting DD at these points ($E > 10^{-9}$, E1 ruling), passes.

7 Required processes: references and development validation

The plan requires the $2 \to 5$ processes of Table 1 (plan §5.7), compared with MadLoop and Recola, and with OpenLoops where a public library exists. A reference provider is run only if its generation completes in about 45 minutes, or in a time comparable to pyAmpliCol's own generation of the process (HANDOFF §0.1, addendum F). Otherwise pyAmpliCol is validated by Ward and gauge identities, Catani poles, LC/NLC/FC and DD/AP consistency, stability checks, lower multiplicities and existing references.

Table 21: **Provider generation on this host and the 45-minute rule** (HANDOFF §0.1): the rows of Table 1 where a provider generation takes more than 45 minutes or did not finish. MadLoop: graph generation + source export + compilation; Recola: process generation; OpenLoops: prebuilt public libraries (no generation time). Records: `data/g1.json` (G1 rows) and `data/measured.json` (`provider_generation`, the reference runs of the $2 \to 5$ rows), through `data/summary.json`. Every other row of Table 1 generates within 45 minutes with every provider (longest: MadLoop 11.9 min for $gg \to t\bar{t}e^-e^+g$ (EW)).

Process	MadLoop	Recola	OpenLoops	45-minute rule
$gg \to gggg$	2.4 h (285 + 8 354 + 37 s)	35.4 min	prebuilt ppjjjj	**MadLoop over 45 min; results exist and are used**
$u\bar{u} \to ggggg$	12.2 h (5 793 + 38 040 + 122 s)	2.8 h	no public library	**MadLoop, Recola over 45 min; results exist and are used**
$gg \to ggggg$	diagrams 17.4 h; unfinished (last state: exporting)	unfinished (last state: running, 2026-10-20)	no public library	**stopped 22 Oct; not to be restarted (HANDOFF section 0.1)**

Table 22: **Native validation of the required set** DEV (`work/v2-generation-scale` before its integration into `loop` at `4e9b5a9e`, DESIGN WP3.1: ninth increment (`6f5a64cc`; threads not recorded) for $u\bar{u} \to e^-e^+d\bar{d}g$, $gg \to gggg$, $gg \to ggggg$; eighth increment (`0ea83af2`; threads not recorded) for $u\bar{u} \to e^-e^+ggg$, $gg \to t\bar{t}e^-e^+g$ (QCD), $gg \to t\bar{t}gg$; seventh increment (`fef92189`; threads not recorded) for $u\bar{u} \to ggggg$): generation (`LoopSlice.generate`, then the plan export; peaks of each stage; shared host under load), then binary64 virtuals with native per-point inputs and the production counterterm route at the fixture points, against the references committed on that branch (`tests/fixtures/loop/required`; fixture tolerance 2×10^{-8}). Largest deviation over the references and points, metric of Table 1; the per-reference values are in `data/measured.json`. For $gg \to ggggg$ the references are NJet (double-double and two binary64 runs) and the Catani operator for the poles, top-free with $N_f = 5$; its estimate column is the rotation plus reduction-direction check per point.

Process	References	Born	Finite	$1/\varepsilon$	$1/\varepsilon^2$	Estimate E	Generation + export
$u\bar{u} \to e^-e^+d\bar{d}g$	ML, Rc, OL · 3 pts	1.9×10^{-14}	8.7×10^{-13}	2.3×10^{-13}	1.0×10^{-13}	1.2×10^{-12}, 2.8×10^{-12}, 1.1×10^{-12}	37.7 s 0.4 GiB
$u\bar{u} \to e^-e^+ggg$	ML, Rc, OL · 3 pts	2.3×10^{-14}	8.5×10^{-13}	4.6×10^{-13}	2.0×10^{-13}	—	36.7 s 0.5 GiB
$gg \to t\bar{t}e^-e^+g$ (QCD)	ML, Rc, OL · 3 pts	1.4×10^{-14}	5.7×10^{-9}	3.7×10^{-12}	5.3×10^{-12}	—	43.7 s 0.6 GiB
$gg \to t\bar{t}gg$	ML, Rc, OL · 3 pts	2.1×10^{-14}	1.1×10^{-11}	1.0×10^{-10}	1.6×10^{-11}	—	38.5 s 0.7 GiB
$gg \to gggg$	Rc, OL · 2 pts	4.7×10^{-14}	4.5×10^{-11}	1.2×10^{-11}	5.2×10^{-12}	1.5×10^{-10}, 1.4×10^{-10}	154 s 2.2 GiB
$u\bar{u} \to ggggg$	ML, Rc · 3 pts	1.7×10^{-13}	2.9×10^{-10}	6.1×10^{-10}	4.4×10^{-11}	—	517 s 6.1 GiB
$gg \to ggggg$	NJet DD, NJet R1, NJet R2, Catani · 1 pt	8.1×10^{-13}	3.8×10^{-9}	1.5×10^{-9}	1.1×10^{-12}	2.0×10^{-8}	51.5 min 28.7 GiB

The $gg \to t\bar{t}e^-e^+g$ (QCD) finite part deviates by 5.7×10^{-9} at one of its three points and by at most 3.7×10^{-12} at the others: the point's known binary64 conditioning (earlier record, DESIGN WP3.1 fourth increment: $\le 3.6\times10^{-12}$ at two points; third point 5.7×10^{-9} in binary64 (estimate 1.7×10^{-9}), MadLoop to 1.9×10^{-14} in DoubleFloat). The electroweak target is validated by the M-SM campaign (Section 4, Table 30), row TV07 (loop_qcd_qed_sm_Gmu, native inputs, 3 points, 124 checks): double-double = MadLoop QP, Recola and OpenLoops to 7.4×10^{-13}, binary64 from double-double 5.1×10^{-11}, poles = Catani to 6.4×10^{-26} (DD), gluon Ward identity at a stable top 3.2×10^{-28} (DD); its QCD twin TV07@qcd 1.4×10^{-12} (91 checks) G1 times and gates it since the final retime (`data/msm-g1.json`, Section 4; Table 1), and its generation and validation are the integrated configuration's, not a development record (so it is not in the table above). The QCD classes are M-QCD's (Section 5). The $u\bar{u} \to ggggg$ row comes from DESIGN.md WP3.1 seventh increment (fef92189), the newest record of that process: the later re-runs on the merged heads (eighth and ninth increments) do not include it. It uses the same production counterterm route; its check is in Table 24.

8 Seven and eight gluons: NJet and pyAmpliCol

8.1 NJet: Table 5 of arXiv:1209.0100, re-measured on this host

NJet is the independent reference for $gg \to ggggg$, where MadLoop and Recola cannot generate under the 45-minute rule and OpenLoops has no library. It is massless five-flavour QCD, and it is not one of the gate's matched providers. At eight gluons ($gg \to 6g$, M-FINAL item 1) NJet's leading-colour primitives (NGluon, on exactly rotated points) check pyAmpliCol's LC virtual in double-double to 6.3e−13 and in binary64 within 10 E (the $2 \to 6$ records); no provider times eight gluons.

Table 23: **NJet, s per point**, double precision, two evaluations each (NJet's scaling test). Paper: Xeon E3-1240 at 3.3 GHz, NJet 1.0. This host: one pinned core, idle at launch, of the shared host; 10 frozen points per channel, mean (spread ≤ 2 %); the multiplier is time ÷ the paper's. De-symmetrised sums equal the standard sums only after phase-space integration; they are timing figures, not point values.

Channel	Table 5	NJet 3.1.2		NJet 2.1.1		Table 5	NJet 3.1.2		NJet 2.1.1	
	standard full-colour sum					de-symmetrised sum				
$gg \to gg$	0.03	**0.00295**	×0.1	**0.00291**	×0.1	0.03	*no class*		*no class*	
$gg \to ggg$	0.22	*defective channel*		**0.141**	×0.6	0.07	*defective channel*		**0.0452**	×0.6
$gg \to gggg$	6.19	**4.59**	×0.7	**4.58**	×0.7	0.57	**0.39**	×0.7	**0.38**	×0.7
$gg \to ggggg$	171.31	**142**	×0.8	**246.7**	×1.4	3.07	**2.41**	×0.8	**4.19**	×1.4

At seven gluons one NJet 3.1.2 evaluation without the scaling test takes 71.2 s (standard) and 1.22 s (de-symmetrised; means over 10 points), at 69 MB; double-double costs 1 990 s per point. NJet 3.1.2's numerical five-gluon channel is 9/2 too large (NJet 2.1.1 is used there), and at the unrotated seven-gluon points NJet keeps only 4–6 digits in double precision (NaN at 3 of 10), so the reference values are taken at exactly rotated points.

8.2 pyAmpliCol at seven gluons DEV

On the WP3.1 branch before its integration (G-int; Table 22) the full $gg \to ggggg$ virtual, with the counterterms of the recursion and native per-point inputs, top-free with $N_f = 5$, agrees with NJet's double-double values to 3.8×10^{-9} in the finite part (per point: 3.8×10^{-9} over all references, point 0, ninth increment) and to 1.1×10^{-12} in the double pole; each point's rotation plus reduction-direction estimate bounds its difference. Generation (ninth increment) took 44.2 min at 14.7 GiB and the plan export 7.3 min at 28.7 GiB (peak RSS of each stage; 160 MB of plan JSON; the whole validation run, point evaluations included, peaked at 38.6 GiB), inside addendum H's budget (hours, below 150 GB). DESIGN WP3.1 attributes the cost: 325 s before the counterterm recursion, 1,122 s with its counterterm program, 2,766 s once poles-uvir's always-generated companion merged (eighth increment), and 2,652 s after the ct-simplify merge (ninth increment, this record), 1,843 s of it the companion's seed table, which built generic seeds. Since then native-seeds' in-crate colour reductions and generation-scale-2's lean companion (sharing the slice's certificates) brought the companion's seeds to 4.7 s (DESIGN WP3.1), and G1's stored process on `loop` (`ca602436`, full colour) generates in 0.7 min plus 1.5 min of export at a 4.4 GiB peak (one pinned core, cold caches; Table 12). A point with its rotation and reduction-direction checks took 82.3 s, a single-point development evaluation (threads not recorded): not a batched one-thread time and not comparable with NJet's 71.2 s per evaluation on one pinned core. The batched one-thread timing is G1's (Table 1, Section 6), and so is the check of the timed configuration against NJet's ten points (Table 24).

8.3 Validation without a matched provider

Table 24: **Checks of the seven-leg QCD results.** First and third rows: G1's checks of the timed configurations (engine `ca602436` on `loop`, native per-point inputs, full colour; binary64 finite part and Born, and the double pole of G1's physical rows against the universal one); second row DEV on `work/v2-generation-scale`; Ward row WP2.5 on `loop`; deviation metric of Table 1. NJet values: fixture `gg_ggggg_njet.json` (NJet 3.1.2 double-double at exactly rotated points, scaling-test estimates ≤ 1.3×10^{-24}; points 0–2 are G1's stream points 0–2 at NJet's scale).

Check	Result	Status
$gg \to ggggg$ full virtual against NJet (top-free, $N_f = 5$), 10 points (G1)	finite ≤ 1.5×10^{-7} against NJet's double-double values (median 3.4×10^{-9}; NJet's binary64 values ≤ 7.1×10^{-9}), Born ≤ 9.3×10^{-13}; every point within its estimate E; asserted in DD at points 0, 1, 4, 5, 6, 7, 8, 9 ($E > 10^{-9}$, E1 ruling); above the tolerance: point 4 1.5×10^{-7} ($E = 1.5\times10^{-6}$) and point 8 2.6×10^{-8} ($E = 3.4\times10^{-7}$), where the double-double values equal NJet's; double pole against the universal one ≤ 1.2×10^{-11} (points 0, 1, 2)	passed (tolerance 2×10^{-8})
Top-free full virtual against NJet at lower multiplicity	$gg \to gg$: ≤ 2.1×10^{-14} at 10 points, full virtual; $gg \to ggg$: ≤ 2.3×10^{-13} at the stable points; the flagged point 3×10^{-14} in DoubleFloat; $gg \to gggg$: 1.0×10^{-11} at point 0; point 1 1.6×10^{-7} in binary64 (estimate 2.7×10^{-6}), 4.7×10^{-15} in DoubleFloat (`test_loop_v2_topfree.py`)	passed
$u\bar{u} \to ggggg$ full virtual against MadLoop and Recola, 3 points (G1, the timed stream points)	finite: MadLoop ≤ 2.1×10^{-12}, Recola ≤ 2.9×10^{-10} (Recola's own: pyAmpliCol's double-double value lies 2.9×10^{-10} from it and 1.4×10^{-12} from MadLoop); Born ≤ 1.7×10^{-13}; double pole against the universal one ≤ 2.5×10^{-13}; the providers agree to 6.2×10^{-10} (reference merge)	passed (tolerance 2×10^{-8})
Ward identity, one gluon's $\varepsilon \to p$, at seven gluons	MQ-WARD (`d3081568`): the generic gauge-Ward check on the native per-point input path, top-free $gg \to ggggg$, legs 1, 4 and 7: virtual ≤ 2.9×10^{-10} (binary64) and 2.8×10^{-26} (DD) of the physical scale, Born ≤ 9.8×10^{-16} / 5.3×10^{-32} (nightly `test_loop_v2_ward.py`; `data/milestone.json` `status_notes`)	**passed (MQ-WARD) at seven gluons (addendum H)**

9 Capabilities at this milestone

Table 25: **Capability matrix** of the v2 engine on `loop` at ca602436, with branch work marked.

Feature	Status	Boundary
Recursive engine	supported	Trees, families and loop numerators from recursions; no loop diagram is enumerated (FeynKit graphs are test oracles only). Any number of coloured fermion lines, identical or distinct flavours, QCD and EW (fermion-lines, db1ed5cb: the fold of mixed colour-connection parities before the evaluation reuse, Python and Rust lanes; $u\bar{u} \to d\bar{d}s\bar{s}$, $u\bar{u} \to u\bar{u}d\bar{d}$, $u\bar{u} \to u\bar{u}u\bar{u}$ and $u\bar{u} \to u\bar{u}e^+e^-$ equal MadLoop in quadruple precision and Recola). Since ew-leptons (EWI-5, b82d8280) any number of lepton and colour-singlet fermion lines, through the loop too, with one reading-sign rule and one fermion classifier taken from the tree engine (Bhabha against MadLoop; genericity audit 2 G2-M5, G2-M6); since ew-epsilon (6da37352) one root-closure reading sign on every fermionless vertex (G2-M7). Listed, typed refusals replace every silent-wrong-value path of the second genericity audit (fail-closed, fb54246c; visible in the dry run). Counterterms: see below.
Default engine	supported	v2 is the only one-loop engine (MQ-LEGACY-REMOVE, 539dffdd; user 2026-10-28 ~16:10 Q3): the per-diagram legacy engine, its plan v3 artifacts and its precision-rescue options are removed (the built-in rescue stack of M-FINAL-b replaces them in the Python API, see Stability); `loop.engine`, `--loop-engine` and a non-default `[evaluation.stability]` are refused with the removal reason; tests that used the legacy engine as an oracle keep an independent one (frozen captures or providers).
Shipped models, public route	supported	The wheel ships MG5 3.7.0's `loop_sm` and its restrictions with a `[loop.model]` profile (MQ-MODEL, 321efdb4): every QCD class generates through `Generator` / TOML / CLI without a `tools/` helper or hand-written declaration, a multiparticle label one entry per subprocess, bitwise the tools route (Section 5); MadLoop syntax mapped (`[virt=QCD]` → `loop_coupling_orders`, `QED=n` → `max_coupling_orders`, `/ x` → `particle_veto`), decay chains and `$` refused typed (decays are served since M-FINAL-b); `pyamplicol self-test` evaluates a `loop_sm` virtual against MadLoop. **Since M-SM** the wheel also ships `loop_qcd_qed_sm` and `loop_qcd_qed_sm_Gmu`, UFOmodeldb 6ca4b120 plus the user's issue-1 patch (R_2 of R2_G0tt split by loop origin, exactly unitary CKM, the G_μ bottom-mass charge counterterm; PROVENANCE and `MANIFEST.sha256` record the local delta; MS-MODEL 2ae8b88b, MS-UFODB 0e5930c2, MS-UFO-PATCH 3eb4db9f), with digest-keyed profiles and the structural Goldstone-flag repair: every EW class generates through the public route, bitwise the tools view (Section 4); a multiparticle request generates its members with a Born and records the loop-induced members as skipped (`loop-induced`, Q2); the self-test also evaluates an EW virtual from the shipped prepared artifact (= MadLoop to 5.4e-15).
Full, NLC and LC colour	supported	**Full colour is the default** of loop requests (user 2026-10-28 Q1); LC and NLC are opt-in through the shipped `loop_sm` N_c declaration (MQ-NC, d8d9349f: every counterterm vertex, resolved by structural digest; formal-N_c projections equal to N_c scans and checked against independent oracles, rule 7; an uncovered vertex refused, `colour-declaration-coverage`). **LC helicity and colour-flow selection** (MQ-SELECT, e4964fc3): `helicities=` / `color_flows=` at generation (recursion pruned to the selection) and at evaluation, tree-engine IDs; a complete set of selections sums to the inclusive LC value; NLC and full colour keep their API (selectors refused there). Exact seed tables in N_c, built in Rust since native-seeds (7e3841f3), with native colour cells, seed records and dipole orbits since native-cells (5ebcebea; the Python stages are the oracles); LC/NLC by generated classes and truncated cells. LC is faster than full colour end to end on every G1 row at $B = 128$, the $2 \to 5$ rows included (LC/FC 0.20–0.65, Table 16); the top-free seven-gluon LC plan keeps every full-colour family by construction (Section 6.2). Since identical-flavour-lc (5cb5ea83) LC and NLC of identical-flavour quark lines (from $u\bar{u} \to u\bar{u}g$ up, any number of lines) are exact formal-N_c projections: blocks and counterterm program split by U(1)-subtraction N_c grade in both lanes, order-free grades (G2-M1; the limitation of m1026a is removed, Section 6.2); a block that still mixes grades is refused (`colour-accuracy-nc-exact`, fail-closed). **EW LC/NLC is not supported, by design** (user 2026-10-30 Q4: the typed `colour-declaration-coverage` refusal is permanent); electroweak requests run at full colour. **Since M-FINAL-b:** `strict-lc` / `strict-nlc` and N_c-graded colour output (PR #32), the exact U(1)-subtraction weight $1/N_c$ (PR #34; binary64 quark-line values move by rounding only, model compiler 17: older loop artifacts are refused at load and regenerated); stored LC selections are masks on one template per process with per-call row liveness (mf-masks, mf-masks-rows: no executor or native heap per selection, bitwise the per-selection oracle).
FFT colour contraction	supported	Certified symmetric-group FFT of the seed contraction in the native executor at full colour and NLC, binary64, double-double and MPFR (MQ-FFT, da6980b0): generic over the plan's certified pairs, policy `auto | direct | fft` (the algorithm run is reported; a forced FFT without a certified pair is refused, `forced-seed-fft-unserved`), public `color.contraction` route; FFT = direct to rounding up to $gg \to ggggg$; the seed stage ×0.46–0.70 of direct at S_5, S_6, slower below, the whole evaluation unchanged within noise (Table 11).
Colour and spin correlations	supported	Directed colour histories at LC/NLC/FC, spin sources, resolved helicities, flavour aliases, Ward sources (WP2.5). Since generation-scale-2 the directed metrics are native and stored with the process (correlated and resolved calls of a stored process regenerate nothing; one seed column per self-adjoint dipole): declaring the NLO dipoles costs +10 % of the six-gluon generation (38.5 against 35.0 s); at seven gluons 784 against 419 s, whose metric tables go to `native-cells`' next run. LC/NLC correlations of identical-flavour quark lines use the graded blocks of the colour row. **Independent oracles since M-FINAL-b** (MF-CORR, 821ae808): Recola's colour, spin and spin-colour correlations on four processes equal the one-loop `evaluate_correlated_many` and the tree `Runtime.evaluate_correlated_many` (Born ≤ 6.5×10^{-14}, Laurent coefficients ≤ 4.1×10^{-10} of the row scale, Table 7); a spin-correlated leg without a four-component vector source is refused at generation, as on the tree route (MF-SMALL2); correlations with an LC selection are refused (`selected-correlators`).
Per-order partials	supported	Every squared coupling order at tree level and at one loop, mixed Born orders, per-correction loop orders.
Poles at runtime	supported	`poles = "all" | "none" | "uv" | "ir" | "split"` (WP2.10, on loop since 24 October): finite-only mode with the finite part bitwise unchanged and the pole work skipped (natively per counterterm interaction since generation-scale-2's I5 record); `"uv"` and `"ir"`

Feature	Status	Boundary
		with the user's semantics of 23 October, from a lean stored companion of the certified-scaleless families (bitwise at $gg \to ggg$ and $u\bar{u} \to gggg$, equal to rounding at six and seven gluons) and in the SDK v4 too; the IR poles sum one Catani operator per correction, QED included (ew-poles, 8b52179b; a correction without a prediction gives none); since e6e7543a a mixed QCD×EW sector is predicted only where the requested corrections complete its loops, and joint requests predict it exactly; since MQ-EWJOIN a sector whose QCD loop selection omits IR-divergent loops (gluon exchange between quark lines joined by a W/Z in the loop, `separated-exchange`) is not predicted and checked against MadLoop per split order. Photon-only QED corrections on `loop_qcd_qed_sm` are refused (`veto-unlabelled-counterterms`: the veto reaches counterterms with empty or nominal loop labels; full EW on that model matches Catani; EWI-7). The UV residue comes from an interim local module until the NinjaRS re-pin; a `"uv"` estimate needs a rotation check. **EW pole split** (MS-POLES, `42b2794f`): on EW requests `"uv"` and `"split"` come from the model's pole regulators (`epsUV`, `epsIR`) under both mass schemes, at the cost of one more regulated pass of the tree program; the conventions record the regulators (= MadLoop on DY01 and TT01).
Stability estimates	supported	**Recommended checks** (MQ-STAB, `a2aa0d84`): `loops.v2.stability.RECOMMENDED_CHECKS`, one generic exact rotation and the rescaled reduction check, the O1 estimate of the campaign and G1 (calibration: binary64 misses 0 % / 0.05–0.10 % of points, DD none; true error ≤ 10 E at every campaign point; R3, DECISIONS 2026-10-29 ~18:25). Rotation checks (any number, 0 = estimate absent), reduction direction, exact rescaling, in any combination; no automatic escalation by default. **Rescue stack** (MF-RESCUE, `2234ee33`; user Q9: only in pyAmpliCol's Python native evaluate API, never in an SDK, the CLI or the config): `stability=VirtualRescuePolicy(...)` per call or as the `Runtime.load(..., stability=)` default, `DEFAULT_RESCUE` double → double-double → arbitrary precision (57 digits) to a target of 10^{-5}, each level with an Euler rotation (`euler:0.1,0.2,0.3`) and the rescaled check; per-lane `VirtualEvaluation.rescue` records and `rescue_counts`, typed refusals; the legacy `VirtualStabilityPolicy` is removed. A non-finite counterterm coupling fails only its own points (MF-SMALL, `isolate_scale_failures`). Generation's post-build validation re-checks a binary64 miss alone against a 32-digit reference: double-double must agree at the same tolerance, binary64 Borns at the tolerance and binary64 virtuals within the larger of the tolerance and their own estimate (one exact rotation and the reduction direction; an estimate above 10^{-9} or absent refuses), else generation refuses; an accepted point is logged (onshell-consistency, `4fc72e39`). The rotation and reduction-direction checks flag the fast binary64 lanes' known misses (precision row): use them for such kinematics or model units.
Precision carriers	partial	Binary64, DoubleFloat, arbitrary precision and exact rationals in the native executor; since G-int the native per-point inputs run in binary64, DoubleFloat and arbitrary precision (DESIGN §5.5). Since `7db32901` binary64 runs SymJIT plane applications and needs no Symbolica licence; DoubleFloat and arbitrary precision run under pyAmpliCol's library key. Since precision-ap (`a8409830`) a master-integral invariant of one external leg is that leg's squared mass at the point's parameters: DoubleFloat and arbitrary-precision virtuals were O(1) wrong at binary64-rounded massive momenta and now agree with a MadLoop quadruple-precision run within 10^{-15} (and DoubleFloat massive spinors along $-z$ are fixed); since generation-scale-2 derived masses (e.g. M_W of `loop_sm`) are parameter-program outputs, bound exactly (genericity M2). Since `onshell-threshold` (`6e123fdb`, G2-M3) binary64 lanes at the fast engine's absolute thresholds (light masses, invariants near a mass shell) go to the generic binary64 engine. **The fast binary64 lanes' light-like spinors** near the $-z$ axis (off by up to ~10^{-5} before, with silent percent-level errors at model units ~ 10^4 from GeV) are taken from the larger light-cone component since NinjaRS #10 (in NinjaRS main since mf-symbolica: the five xfails are plain cases, large unit factors keep the virtual). **Tree runtime** (lsimon125's PRs #26, #29, #36, #39, LP-PRECISION): double-double and arbitrary-precision contracted-colour replay, exact correlated weights, couplings, parameters and masses at working precision, the G_μ UFO boost fix. **Binary64 evaluation of a stored artifact never calls Symbolica** (user 2026-10-30), complex-mass artifacts included: a cut-function argument that becomes real when complex-mass widths are exactly zero takes its real part at run time (MS-MASTERS, `71e2041e`; binary64 = DD = AP(32) on the CM01 zero-width fuzz, = MadLoop where its double and quadruple precision agree).
Native executor, schedules	supported	Rust, 16 lanes (also for a one-point call, a G1 lever), threads; native tree closures in lanes of points with runtime AVX-512/AVX2 widths (lc-perf); in finite-only mode the counterterm interactions that only feed poles are skipped natively (generation-scale-2); G1 decision on the schedule-dependent stages (Table 19, provisional until compiled trees): M for pure-gluon LC from $n = 6$, B for every other class at $n \leq 6$; F as oracle. The NinjaRS reduction runs serially. Development report: stage times, modelled kernel work, master-integral time. **Memory and speed rounds of M-FINAL** (seven-leg table, values bitwise): tree/CT fill and batches within fixed state, row and chunk-input budgets (mf-8g-2), stored members read on demand at load with a compact Rust plan and tree program (mf-8g-2b), binary stored forms and native CT DAG maps (mf-8g-1), lighter generation (mf-8g-1b, mf-8g-1c), row-sized interference blocks (mf-perf-opt).
Plan v4 artifacts and API	partial	Published and loaded through Python, TOML and CLI; the artifact stores the native process (`point_inputs = "native-v2"`). Since `7db32901` (addendum G) its SymJIT -O2 applications are stored as IR and recompiled at load for the loading CPU (AVX and FMA3 required), no host string is embedded, and binary64 values are bitwise equal after relocation to another machine, user and path; since M-FINAL-b SymJIT 2.28.3 (MF-UP: siravan/symjit #26, #30, #31) makes the stored runtime members byte-reproducible between generations, and the parameter program's evaluator bytes no longer depend on how the model was obtained (MF-SMALL2: the runtime-bound scalar predicates in their own id block). **Permutation aliases and process-group members** (MF-MAP, `9ad0c77f`): a one-loop alias or a permuted member of a plan v4 process is served from the stored target at permuted momenta, colour and spin correlations composed, LC anchors relabelled;

Feature	Status	Boundary
		numerical process-group analysis stays refused (process-group-analysis) and helicity/ colour-flow selection is refused for mapped processes. **Virtual C ABI version 2** (MF-CLEAN, 25d39f6b): the legacy agreement check and the SDK policy argument are removed from the Rust SDK, the C ABI, the C/C++/Fortran wrappers and UMAMI. Since model-determinism (4e9262b5) the compiled model and the artifact's digest no longer depend on the process history; since native-lowering (79f2a5bd) the plan JSON is written by the Rust plan model (Python writer as oracle). Since generation-scale-2 (11542fd5) no call of a stored process regenerates (every pole selection with the stored companion, correlated, resolved, flavour and catalogue calls), a process the native export cannot store fails its generation with a typed refusal (genericity M1; the plan-only, regenerated-at-load artifacts are gone), and the Rust/C/C++/Fortran SDK v4 evaluates these artifacts with the "uv"/"ir" selections, bitwise the Python API. Since sdk-parity (f6619f34) one parameter rule holds in the Python API, V2Runtime, the Rust SDK and the C/C++/Fortran clients: ASCII decimal values checked before names, unknown, generation-fixed (naming who fixed them) and scale parameters refused with typed errors (G2-M2). The release natives need glibc $\geq$ 2.38, a lower glibc giving the typed NativePlatformError at load (glibc-floor, 5301e24a). Since M-SM external-photon and complex-mass artifacts are native-v2 too (ew-native, MS-CMS-N: complex pole masses through the native trees and executor, bitwise after relocation). **Since M-FINAL-b:** the SDK v4 extras (mf-sdk: resolved, colour/spin-correlated, double-double and arbitrary precision through the Rust SDK, the C ABI, the C++/Fortran wrappers, UMAMI and the packaged Rust wrapper); every stored saved evaluator is canonical (mf-symbolica: referenced symbols only, numbered by name), so artifact bytes no longer depend on the producing Symbolica (Community or PyPI 3.0.3) or the process history. Open: the plane check's outcome depends on the parameter order (G2-M4: a listed plane-mismatch refusal, never a wrong value).
Counterterms (model UV, R_2)	supported	Plain tree vertices of the recursion inserted exactly once (WP2.2, a18b7599; addendum J, b4d76cc3), admitted by coupling orders and the public particle_veto; no witness table or loop matching in production. **Stock loop_sm compatibility rule** (MQ-MODEL): its four-G^0 vertex V_1001 names the three-slot constant SSS1 where the four-slot SSSS1 (same exact value) is meant; model preparation moves such references by structure, only on vertices a request reaches, and records each repair (pyamplicol_source_repairs, source_compatibility_diagnostics in the dry run); any other slot mismatch is refused (lorentz-slot-arity). Vertices whose couplings vanish at the restriction leave the loop catalogue (fail-closed-2, loop-vertex-outside-tree-model otherwise). An unlowerable counterterm vertex that no tree of the request reaches is inert and recorded in the artifact; a reachable one is refused. Site closure is the test oracle. Per point they are formed natively (G-int), the Python stages kept as oracle; generation of the counterterm route has output-identical savings (ct-perf, 201bece5: compile memos, a content-addressed lowered-model cache), and since A4 (1b13005a) the native DAG builder is the production route for the loop blocks and the counterterm program (its counterterm emission byte-identical to the Python oracle; 11 listed features keep a recorded per-request Python fallback, any other native failure fails closed), at parity with the site oracle. In EW requests ('t Hooft–Feynman gauge, EWI-4b, 8b19dae8) the Goldstone–vector mixing counterterms are lowered as vertex insertions. The UFO's EW counterterms (ew-ct, 392d3f4a) are lowered through the model's own helper bodies with MadLoop's branch side for exactly real arguments: every Laurent piece of MadLoop's model library equals pyAmpliCol's to 3.7×10^{-15} in binary64 and 6.0×10^{-31} in double-double (against MadLoop's quadruple precision). Since EWI-3R (59d47cb3) the lowering runs at scale: native, parallel kernel-equivalence proofs (Python the oracle), one reference pair per Laurent coefficient, shared Horner schemes; about 26 s cold at the target instead of 259 s, 8.4 % of the G_μ $gg \to t\bar{t}$ (EW) generation (criterion < 20 %, user 2026-10-28); a non-polynomial coupling term or one below ε^{-2} is refused when reachable. Lorentz Epsilon is lowered (EWI-7b, 6da37352; sign = MadLoop to 10^{-12}), and the Goldstone–vector mixing sign, from the tree engine's momenta, matches MadLoop's counterterm selection. **UFO-complete R_2 insertion** on the EW UFOs (user Q7): every R_2 vertex a tree of the request reaches is inserted. Compatibility note: MadLoop never inserts R2_G0tt_b on the unpatched models (its label [W+, b] equals no loop's propagator set; mg5amcnlo#475, up to 8.4 % on $t\bar{t} \to zh$); on the shipped patched models MadLoop's exports call every R_2 coupling pyAmpliCol inserts (MS-VALIDATE).
Decays $1 \to n$	supported	**Since M-FINAL-b** (MF-DEC-1 eddb4b49, MF-DEC-2 28b24a2a): one initial leg is admitted generically ($1 \to n$, $n \geq 2$) at tree level and one loop, QCD and EW, through the same recursion; the decaying particle at its real mass with zero own width (the existing width guards; internal CMS widths allowed), MadLoop's convention (amputated loops, its absorptive self-energy dropped, δZ from the UFO counterterms); the rest-frame massive Dirac spinor fixed (was $2 \times \lvert M\rvert^2$ at $\lvert p\rvert = 0$). XX01–XX08 = MadLoop's quadruple precision (Table 5: 8 of 8 pass, DD ≤ 9.6×10⁻¹⁴), $Z \to u\bar{u}$ = the massless form factor, Ward on $t \to bW^+g$, LC/NLC by the N_c scan, public route and native inputs. Refused typed: loop-induced decays (loop-induced), chains and $ (madloop-syntax).
UFO genericity: HC_NLO_X0	partial	**Since M-FINAL-b** (MF-GM-HC, e2f0c78e): the Higgs characterisation UFO at NLO, HC01–HC04 and HC06 = MadLoop 3.7.0 through generic fixes (UFO namespace head, Feynman-gauge refusal only where massive vectors run in the loops, positional counterterm keys, a colour-singlet contact split through an adjoint auxiliary); **since M-FINAL** HC05 ($gg \to x_0gg$) too (MF-HC05, e3d4ef4d): a counterterm colour the tree vocabulary refuses (its five-point $ffff$ R_2 contact) is projected exactly at $N_c = 3$ onto products of at most two primitive adjoint tensors (zero residual required, else refused); all six ≤ 2.6×10⁻¹² (Table 6). **PR #38** (LP-MODELS): compiled-mode HEFT quark tree amplitudes (built-in SM-HEFT and HC_NLO_X0) were 4× too large, now = MG5 and the MadLoop Borns. Open, typed: $u\bar{u} \to x_0ggg$ (the U(1)-leak check), single-trace and C-odd five-point R_2 colours.
Electroweak, complex mass	supported	**M-SM** (Section 4, Table 9): QCD, QED and joint [virt=QCD QED] corrections on the shipped loop_qcd_qed_sm and loop_qcd_qed_sm_Gmu reproduce MadLoop in every class of m-sm/ PLAN.md §1.1 (69 of 71 rows pass, LL02 and SC02 partial: F7, accepted). EW native inputs (ew-native, EWI-N), the EW ladder (ew-validate, EWI-7: ghosts of vetoed vector bosons,

Feature	Status	Boundary
		Goldstone-side root mixing), **per-leg photon tagging** (MS-PHOTON: `loop.external_photons = "final"` by default, `"all"`, `"none"` or a leg list, as MG5's `!a!`; $\alpha(0)$ and the conversion counterterm on the tagged legs only), the **complex-mass scheme** (MS-CMS: MG5's transform per (mass, width) pair, coloured widths included; MS-CMS-N: native complex pole masses; = MadLoop CMS on CM01–CM05 and CM07–CM10), the EW pole split (MS-POLES), the EW generation stages in Rust (MS-GEN: catalogue segment polynomials and the parameter program's renaming behind PyO3, the Python stages the oracles, outputs byte-identical, EW generation 2.6–3.2× faster) and Symbolica-free binary64 (MS-MASTERS); since M-FINAL-b the EW export's kernel-string reading and parameter-program renaming run in Rust too (MF-RUST, `b36af289`: Python stages the oracles, byte-identical; EW $2 \to 3$ export 62.3 → 32.4 s). Before M-SM on `loop`: the QED Catani operator, EW fermion lines, 't Hooft–Feynman gauge for trees, loops and counterterms, on-shell photons at $\alpha(0)$, the UFO's EW counterterms lowered at scale, Lorentz `Epsilon`, lepton lines through the loop (EWI-0 to EWI-7b). Refused, typed: the Missing box of Section 1; EW LC/NLC never.
Required set	supported	All eight generate and agree with their references in all coefficients with native inputs: seven in Table 22 (validated on the WP3.1 branch before its integration), the EW target in the M-SM campaign (TV07 against MadLoop, Recola and OpenLoops, Section 4); $gg \to ggggg$ against NJet's double-double values (its binary64 points 4 and 8, above the check's 2×10^{-8}, accepted by the user: within the estimate, DD = NJet); the generator is on `loop` since G-int, and G1 times all seven in batches at full and leading colour (Section 6; the $2 \to 5$ rows since `measure-g1-7leg`; all re-timed at `ca602436`). The seven-gluon Ward check (addendum H): passed (MQ-WARD). The G1 accuracy check asserts binary64 where $E \leq 10^{-9}$ and DD elsewhere (user E1 ruling, `g1-accuracy`); R3 for QCD is recorded (DECISIONS 2026-10-29 ~18:25: two providers per process within their agreement, $gg \to ggggg$ against NJet alone, addendum F). G1 times and gates the EW target since the final retime (Section 4). The provider reference trees no longer read the old working root (reference-relocate, `48aeb123`; provider values bitwise, and the toolkit refuses provider files outside the configured reference root).
Performance gates	partial	G1 on `loop` (Section 6, native per-point inputs, `ca602436`) at $B = 128$: 8 of 8 gated processes pass; 5 below 5 ms per point are not gated (user, 2026-10-24); $gg \to ggggg$ without a matched provider: no gate; the gate is taken at $B = 128$ only (user 2026-10-28): $B = 1$ and the $2 \to 5$ rows' $B = 16$ are reported without a verdict (Table 13). The Python per-point inputs before G-int failed every gate they generated (history). The EW row $gg \to t\bar{t}e^-e^+g$ (EW) is gated since the final retime: ×0.13 of Recola's time at $B = 128$, pass. Re-timed at M-LIT-a0 (the final retime deferred from M-FINAL, on `ca602436`). Missing: the provider protocol-mode columns.
$2 \to 6$	supported	M-FINAL item 1 (user ruling 2026-10-31): $gg \to 6g$ ($N_f = 5$, no top), $u\bar{u} \to 6g$ and $u\bar{u} \to u\bar{u}gggg$, each at LC and full colour, generate, store (every step below 150 GB), load and evaluate on one engine after the optimisation rounds (`2f3f76d9`, an ancestor of the pin; Table 1, records and checks; the N = 8 reduction adapter is tested, MF-8G-0 `7cdd32e9`); double pole, Catani (FC), NGluon (LC), double-double and the quark rows' Ward pass, the 8-gluon Ward is recorded but not asserted (beyond the binary64 scope) and its full-colour and double-double Wards were not run (user, 2026-11-04). Internal checks only: no provider exists.
Built-in SM (tree)	supported	**Since M-FINAL** (MF-NORM `049a9389`, MF-LIMITS-PY `d958f2ea`): equal to the UFO `sm` in the compiled, eager, on-the-fly and recurrence lanes (the defects corrected and their history: Section 1, built-in SM fixes), $1 \to n$ decays; checked by the default-tier sweep `tests/integration/test_builtin_sm_ufo_sweep.py` (every $2 \to 2$ class, the $1 \to 2$ decays, $2 \to 3$ / $1 \to 3$ samples) against the UFO `sm` to 10^{-10}. Typed refusals as for the UFO `sm`: four-quark processes without gluon exchange in the recurrence lane, `dirac-weyl-closure`. A one-loop request needs an NLO UFO (the built-in SM is tree only).
Dependencies	supported	No fork (user 2026-11-02, MF-SYMBOLICA `e0a8375c`; `dependencies/manifest.json`): Symbolica = symbolica-dev `community` `36cbac3a` + the consolidated PR symbolica-dev/symbolica#163 (`20a9da94`), NinjaRS main `b3a29c85`, oneloopmaster main `128ea664`; Python through the Community build of `20a9da94` or PyPI symbolica 3.0.3.
Hard-coded limits	partial	Mechanical limits lifted (MF-LIMITS-PY `d958f2ea`, MF-LIMITS-RUST `72fe20dc`; Section 1, limits box): one rank admission from the reducer, one leg boundary 1–64, checked narrowing, exact rational charges, uncapped mappings and fixpoints, tuning bounds as settings. **Not supported** (capability generalisations, typed refusals, no milestone planned): rank $> N + 2$, multi-fermion loop vertices, coloured n-point EFT contacts, nested singlet contacts, octet/sextet fermions, an LC/NLC menu, binary64-only paths, C-API precision.

10 Architecture of the v2 engine

10.1 The computed quantity

For Born order a, loop order b and squared coupling order $s = a + b$,

$$V_{s,k} = \mathcal{N} \sum_{a+b=s} 2\,\mathrm{Re}\big[B_a^\dagger G_{BL} L_{b,k}\big], \quad k \in \{0,-1,-2\}, \quad V = C_\varepsilon \left(V_0 + \frac{V_{-1}}{\varepsilon} + \frac{V_{-2}}{\varepsilon^2}\right), \quad C_\varepsilon = \frac{(4\pi)^\varepsilon}{\Gamma(1-\varepsilon)},$$

where G_{BL} is the colour metric between Born and loop colour structures and $\mathcal{N}$ holds the initial-state averages and identical-particle factors (MadLoop/BLHA convention; couplings included, nothing divided by the Born). Directed colour and spin correlations replace $2\,\mathrm{Re}(\ldots)$ by the two cross terms $B_A^\dagger G L_B + L_A^\dagger G B_B$.

10.2 Pipeline

For every denominator family f, sector s and observable column o, v2 builds the four-dimensional numerator $N_{f,s,o}(q) = \sum_\alpha c_\alpha q^\alpha$ that NinjaRS reduces: Born-weighted and helicity-summed, produced by three recursions that share all work (Figure 2).

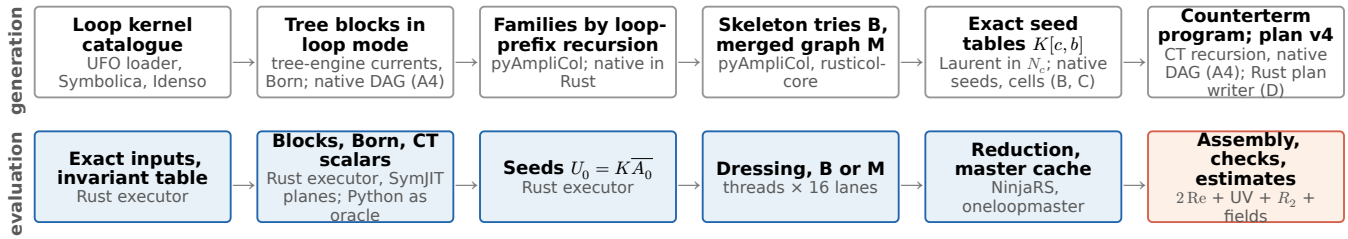


Figure 2: **v2 pipeline.** Upper lane: generation, once per process, into a plan v4 artifact. Lower lane: evaluation per batch of points. Blue: native batched numerics in the Rust executor (the per-point inputs since G-int); orange: Python per batch; white: generation. The owning code or dependency is named under each step.

Trees. The tree engine's colour-ordered Berends–Giele recursion computes, once per point, every off-shell current over every leg subset a loop vertex can absorb (the *blocks*), the Born colour components A_0 and, with WP2.2, every tree with one counterterm vertex. A loop-aware reachability test keeps exactly the currents some family position or root needs. Blocks with transverse sources use the exact transverse Yang–Mills kernels; loop segments, counterterm currents and Ward sources use the full ones.

Families and skeletons. A Recola-type recursion over loop prefixes (ordered leg sets, propagator species and masses) yields the denominator families, keyed by the canonical cyclic sequence of (leg mask, mass) and equal to Recola's tensor-integral identities. Per family it yields the colour-ordered *skeletons*: rings of blocks, each block summing all trees on its leg subset (Figure 3). Skeletons are never stored as a list; they are the leaves of the schedule tries. Only certified structural zeros are dropped (scaleless rings, external wave-function corrections, VEV tadpoles, colour zeros); nothing is restricted to 1PI.

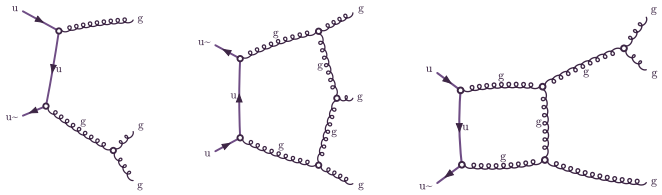


Figure 3: Three graphs of $u\bar{u} \to ggg$ drawn by FeynKit's renderer from its own generation (an oracle, not a production input). In v2 the Born (left) is a closure of blocks; the pentagon (middle) is a ring of five single-leg blocks; the box with a tree attachment (right) is a ring of four blocks, one of them the two-leg current, computed once per point and shared by every skeleton that contains that leg subset.

Loop numerators and schedules. An open-loop recursion dresses skeletons one segment at a time with polynomial-valued objects, $O' = O \cdot Y + \mathrm{raise}\big(O \cdot Z_\mu, q^\mu\big)$, in the NinjaRS monomial layout ($\binom{r+4}{4}$ coefficients at rank r, no rank cap). The seed is the colour-summed Born interference, and each block's helicities are summed as it is attached. One generator feeds three schedules (Table 26).

Table 26: **Three schedules on one generator** (DESIGN §2.2, §5.6).

	F: forward	B: suffix trie (default)	M: merged, primitive-seeded
Born weighting	at closure	at the leaves (one seed per skeleton class)	at the cut, one seed per colour ordering
Shared	prefixes across families	all prefixes with equal family and remainder	prefixes across families and compositions
Colour certificate	none needed	none needed	certified classes: adjoint rings, traces, quark rings
Role	oracle	default; generic fallback	pure-gluon LC from $n = 6$ (G1); planned default from $n \geq 8$

Colour. Colour lives only in exact seed tables $K[c,b]$, Laurent polynomials in N_c between seed classes c (the colour vector obtained by closing the cut of a skeleton's colour flow) and Born basis elements b; per point the seeds are $U_0 = \sum_b K[c,b]\overline{A_0[b]}$, with a symmetric-group FFT where certified. LC and NLC differ from full colour only in which classes are generated and how cells are truncated: at LC each colour ordering is seeded by its own Born amplitude and skeletons below the leading N_c degree are never generated. Certificates (adjoint rings, fundamental traces, quark rings) are proved at loop level, never inherited from the tree level.

Reduction and masters. Families close directly into NinjaRS `TensorNumerator` lanes: one reduction per (family, sector, column), families with equal (propagators, rank) batched as lanes of one call, scalar masters from oneloopmaster cached per event. A second reduction in a relabelled frame reuses the numerators and the cached masters as the cheap stability check.

Counterterms. The model's UV and R_2 vertices, two-point ones included, are ordinary vertices of the tree recursion, and each contribution contains exactly one of them; external field factors are separate sources. Admission is the requested coupling orders, the exactly-once constraint and a public `particle_veto`: a coupling term whose loop particles include a vetoed particle is dropped, and a vetoed particle is removed from every tree vertex, loop line and propagator of the generation, so loops and counterterms stay consistent. Nothing is matched to the loop (user, 23 October; addendum J). On `loop` the counterterms come from the WP2.2 recursion (one counterterm order in the tree recursion, the default route, whose $gg \to ggggg$ program leaves no rule to site closure) with this admission (`ct-simplify`, b4d76cc3): no witness table and no loop matching. Since A4 (`1b13005a`) the native DAG builder emits this program (byte-identical to the Python builder, which stays as oracle) and builds the loop blocks. The UFO's electroweak counterterms (`392d3f4a`) enter the same way, lowered through the model's own helper bodies with MadLoop's branch side for exactly real arguments, and in 't Hooft–Feynman gauge the Goldstone–vector mixing counterterms as vertex insertions (`8b19dae8`). A counterterm vertex that the model cannot lower and that no tree of the request reaches (the tree engine's own reachability, with the program's order limits and `particle_veto`) is inert and recorded in the artifact; a reachable one is refused. `particle_veto = [t]` replaces the internal top-free switch of Table 24; site closure and the witness machinery remain only as test oracles.

Precision, stability and assembly. Every kernel is generic over binary64, DoubleFloat, arbitrary-precision floats and exact rationals. A failed lane fails its point and never branches a kernel. The assembly takes $2\,\mathrm{Re}$ of $i/(16\pi^2)$ times the reduced family sum, plus local UV, R_2 and field terms, per sector and column, with optional rotation, reduction-direction and rescaling estimates and the finite-only mode. The public API (TOML, CLI or Python; v2 is the only one-loop engine since MQ-LEGACY-REMOVE) stores plan v4 as a relocatable PACBIN member (zlib JSON of exact kernel term lists, written by the Rust plan model since `79f2a5bd`, no compiled code, digest checked at load) and the native process (`point_inputs = "native-v2"`: the tree program with the Born, counterterm and field components, their saved exact evaluators and the binary64 SymJIT applications of every tree kernel and of the parameter program, stored as IR and recompiled at load) with the companion of the certified-scaleless families and the observables' tables, so no call of a stored process regenerates it (generation-scale-2), on another machine, user and path too; a process the native export cannot store is refused at generation. The Rust/C/C++/Fortran SDK v4 loads the same artifact.

Table 27: **Asymptotic classes per event** for n gluons, full colour (DESIGN §1.4): $F_n \approx (n-1)!\left(\frac{1}{\ln 2}\right)^n/2$ families, 2^n helicities. The required set ($n \le 7$) is far below the crossover of schedule B with the unitarity class ($n \approx 30$–50); schedule M is planned as the default from $n \ge 8$.

Stage or code	Class
Tree blocks	$\Theta(n \cdot n!)$ currents; $\Theta(n! \cdot 2^n)$ current-helicity rows
Dressing, schedule B	$\Theta((n-1)! \cdot 4^n)$ seed injection; trie dressing dominates for $n \le 8$
Dressing, schedule M	$\Theta(n! \cdot 2^n \cdot n^4)$
NinjaRS	$\Theta(F_n)$ reductions, $\Theta(F_n \cdot n^8)$ work, $\Theta(4^n)$ distinct masters after caching
MadLoop (per diagram, per helicity)	$\Theta(D_n \cdot 2^n \cdot n^4)$, D_n super-exponential ($gg \to gggg$: 26 975 loop diagrams)
Unitarity with exact helicity sums (NJet)	$\Theta((n-1)! \cdot 2^n \cdot n^{6\ldots9})$

Table 28: **Ownership of the mathematics.** pyAmpliCol reuses its dependencies for algebra, tensors, graphs and integrals and owns the recursions and the assembly.

Component	Responsibility
ufo_model_loader	Lossless UFO loop-model import: Laurent-valued parameters and couplings, UV/R_2 categories, loop groups
Symbolica	Expressions, coefficient extraction in q, evaluators; exact rationals and arbitrary precision
Spenso, Idenso (Community)	Lorentz and colour tensors behind the kernel catalogue and the colour metrics
FeynKit (Community)	Test oracles: graph and family counts, overall factors, drawings (Figure 3)
NinjaRS, oneloopmaster	Integrand reduction with Laurent expansion; scalar masters in all carriers
pyAmpliCol tree engine	Blocks, Born, coupling orders and colour flow (Python lane; Rust recurrence)
pyAmpliCol loop v2	Families, skeletons, schedules, seed tables, counterterm program, plan v4, native executor, API

11 Sources and reproducibility

- **Build:** `typst compile benchmarks.typ benchmarks.pdf` in `docs/performance_reports/loop/` (Typst 0.15).
- **Status, in-flight work, decisions:** `data/milestone.json`, from the commits and integration merges of ca602436 since 7266c160 (CI `936e81b3`, 1.0.2 `ec99ad04`, #40 `5fee3347`, CI `b81b77b3`, SDK `053a0778`, CI `36e60f6c`, CI `ca602436`; commit subjects "Integrate work/v2-…"); `workspace/reports/engine-v2/slice-status.md` (integration blocks, DECISIONS 2026-10-23 to 29, the milestone pushes m1025a of 2026-10-25 ~13:45 UTC, m1025c of 2026-10-26 ~06:35 UTC (after the m1025b attempt at `db13c260`, not pushed) and m1026a of 2026-10-27 ~21:00 UTC (`8335e9e8`), and milestone m1028a (launched 2026-10-28 ~09:00 UTC; pin `28a16255`, native `loop-native-v2-slice-5cb5ea83`; the user's rules of 2026-10-28: the speed gate at $B = 128$ only (~08:20 item 2, `gate_decisions.gate_batch`), the counterterm-lowering share below 20 % (~08:20 item 1) and the glibc floor (~07:40; 2.38 since glibc-floor); the known limitations from the coordinator's milestone brief and `known_limitations` in `data/g1.json`; the identical-flavour LC limitation of m1026a, DECISIONS 2026-10-26 ~21:10 item 1, is removed by identical-flavour-lc), milestone M-QCD (DECISIONS 2026-10-29 ~17:20: pin `e18a024b`, native `loop-native-v2-slice-d0fef168`; the user's milestone plan of 2026-10-28 ~12:50 and its rulings ~15:25, ~15:35/40 and ~16:10, the user's rulings of 2026-10-29 ~03:45 (E1) and ~10:20 (precision), the coordinator's ~11:25 (F7, F11) and R3 for QCD ~18:25; `workspace/reports/engine-v2/m-qcd/PLAN.md` §1, §5) and milestone M-SM (DECISIONS 2026-10-31 ~07:45: pin `71e2041e`, native `loop-native-v2-slice-71e2041e`; the plan `workspace/reports/engine-v2/m-sm/PLAN.md` §1, §5 and its §6a user answers of 2026-10-30 ~08:10 (Q1 per-leg photons, Q2 multiparticle members, Q3 decays in M-FINAL, Q4 EW LC/NLC never, Q5 both EW UFOs, Q7 the UFO-complete R_2); the user's rules of 2026-10-30 ~08:40 (UFOmodeldb as ground truth), ~08:50 (M-FINAL order), ~12:45 (the issue-1 patch, with the coordinator's ~13:40) and ~17:35 (binary64 without Symbolica); the accepted F7 and the SC03 τ-mass relaxation; mg5amcnlo#475 and UFOmodeldb issue 1); the user's gate decision of 2026-10-24, 16:05 UTC, is recorded in `data/milestone.json` `gate_decisions` and no longer applies after the re-measure; the user's rules of 2026-10-24, 22:35 UTC — uncoloured generation, green-to-red run times, summary pages with only Table 1, Figure 1 and legends, and the gate rule "below 5 ms: not gated (user 2026-10-24)", recorded as `gate_decisions.threshold_rule` and applied by `scripts/build_summary.py` and this report); milestone M-FINAL-b (DECISIONS 2026-11-01 ~23:35 and 2026-11-02 ~00:20, ~01:05, ~03:08: pin `f3d9e361` = `617b59cd` + the test-only mf-map-fix, native `loop-native-v2-slice-b36af289`; the plan `workspace/reports/engine-v2/m-final/PLAN.md` §2, §3 and the user's rulings of 2026-10-31 ~11:00 (M-FINAL order, rescue only in the Python API, decays at the real mass with zero own width, EW LC/NLC never), ~14:20 (direct R_2 dropped), Q11 (one final retime) and 2026-11-01 ~23:15 (8-leg runs stopped for optimisation rounds); the G1 regression spot check at `617b59cd` against `71e2041e`: `workspace/tmp/g1-m1101-mfb/logs/compare.txt`); milestone M-FINAL, closed (pin `c8885cc5`, native `loop-native-v2-slice-e0a8375c`; USER 2026-11-04 ~08:00 'close M-FINAL now' and the lean close of 2026-11-05 ~07:30, DECISIONS ~07:35: no local gate sweep, no retime, no closing audit, global tests by CI on `loop`; DECISIONS 2026-11-04 ~17:40 (no Ward/Catani/campaign re-runs; status changes from package evidence only), ~22:00 (no pedantic tests), ~19:15 (mechanical limits only), ~08:35 (8-leg stop lifted) and 2026-11-03 ~13:05 (model compiler 17); the integration blocks of the packages in `data/milestone.json` `integrated`; the 2 ⟶ 6 report `workspace/reports/engine-v2/m-final/2to6-report.md`, the seven-leg round table `m-final/round-table-7leg.md`, mf-8g-3′s report `m-final/8g3-report.md` §9 (A24) and the Symbolica re-pin report `m-final/symbolica-repin-report.md`; the built-in SM fixes: mf-norm's package report (its table for the collaborator) with the integration block of slice-status UPDATE 2026-11-05 ~08:05 and DECISIONS ~06:10); milestone M-LIT-a0 (pin `ca602436`, native `loop-native-v2-slice-ca602436`; the user's green light for M-LIT of 2026-11-06 ~07:05 and DECISIONS ~07:20, ~07:32; `workspace/reports/engine-v2/m-lit/PLAN.md` §3.2 LIT-RETIME, §3.7 "Pushes", §5.1, and `m-final/PLAN.md` §2.4 MF-RETIME; the user's request of 2026-11-06 ~06:00 for the loop-induced summary table; the retime's records `workspace/reports/engine-v2/measure-g1/records/mlita0`); `HANDOFF_24_10_2026.md` §6; MASTER_PLAN 2026-10-23/24 entries; REQUIREMENTS_ADDENDUM A, D, F, G, H, I, J; DESIGN §5.5, §3.3, §5.7 (G1), §6.1, §6.2, §7.2, §7.5, §7.9, §12.1, §12.4, WP3.1, WP3.2; the 24 October pushes (59badc7a and b4d76cc3, gates in slice-status.md), the G-int push (79df367a, 25 October), the m1025a push (bae3b9ea, 25 October), the m1025c push (36f43250, 26 October) and the m1026a push (8335e9e8, 27 October). WP3.2 increments and dependencies: `workspace/reports/engine-v2/wp32/PLAN.md`; the EW counterterm comparison with MadLoop and measurement M1: `wp32/M1-ew-ct.md` and DESIGN §6.1; the genericity audits: `workspace/reports/engine-v2/genericity-audit/REPORT.md` and `genericity-audit-2/REPORT.md` (at `11542fd5`, with the first audit's follow-up); precision-ap's MadLoop comparison: `workspace/reports/engine-v2/precision-ap/`; reference-relocate's inventory and evidence: `workspace/reports/engine-v2/reference-relocate/`; model-determinism: DESIGN WP2.5; native-lowering: DESIGN §12.1 and WP3.1 (Rust migration D); onshell-consistency: DESIGN §11.1 and `docs/user/configuration.md`; native-seeds and native-cells: DESIGN §7.2, §7.9 and `workspace/reports/engine-v2/wp31/RUST-MIGRATION.md`; fermion-lines: DESIGN §3.3 and the fixtures `tests/fixtures/loop/required/uux_{ddxssx,uuxddx,uuxuux,uuxepem}.json`; generation-scale-2: DESIGN WP3.1 (its status) and §12.1; EWI-4b and EWI-6b: DESIGN §12.4 and `wp32/PLAN.md`; the packages of m1028a: their integration blocks in slice-status.md (ew-ct EWI-3R `59d47cb3`, ew-epsilon EWI-7b `6da37352`, fail-closed `fb54246c`, sdk-parity `f6619f34`, ew-leptons EWI-5 `b82d8280` with its R2_G0tt_b evidence (DECISIONS 2026-10-28 ~05:20 item 2 and ~08:10 item 1), identical-flavour-lc `5cb5ea83`, gate-b128 `28a16255`), DESIGN §5.7 ("Gate batch"), §12.4 and the listed generation refusals, `docs/user/configuration.md` and `docs/user/ufo-model-coverage.md`; the packages of M-QCD: their integration blocks in slice-status.md (Section 1), `docs/user/loop-sm-and-madloop.md` (the `loop_sm` repair rule) and the listed refusals of `src/pyamplicol/loops/v2/refusals.py` and `processes/refusals.py`.
- **Section 1 (M-SM), Section 4, Table 8, Table 9, Table 30:** `data/msm.json` (`scripts/build_mqcd.py` with `--manifest tests/fixtures/loop/msm/manifest.json`, from the campaign's records, rebuilt byte-identically), `data/msm-findings.json` (MS-VALIDATE's triage, unchanged at the pin), `data/msm-public.json` (the public-route test file's JUnit outcomes and cases), `data/msm-selftest.json` (wheel self-test), combined with the manifest by `scripts/build_msm_matrix.py` into `data/msm-matrix.json`; the G1 EW row `data/msm-g1.json` and `data/msm-g1-values.json` (`python -m tools.loop_benchmark gate` over the EW records; Table 1 through `data/summary.json`); the campaign's records, logs and scripts in `workspace/reports/engine-v2/m-sm/records/m1031-msm/` (a copy of `workspace/tmp/campaign-m1031-msm`: `results/`, the JUnit and logs of the 32 jobs, `wheel/`, `STATE.md`). Re-run at M-FINAL-b (`617b59cd`): `data/msm.json`, `msm-public.json` and `msm-selftest.json` from `workspace/reports/engine-v2/m-sm/records/m1101-mfb/` (a copy of `workspace/tmp/campaign-m1101-mfb`: `results/`, the JUnit and logs of its 33 jobs, `wheel/out/`, `compare.py`, `msm-table.md`, `STATE.md`). The M-QCD re-run at

M-FINAL-b: data/mqcd.json, records in workspace/reports/engine-v2/m-qcd/records/m1101-mfb-mqcd/ (c1/, compare.json against the M-SM re-run's file, which came from records/m1031-msm-mqcd/). At M-FINAL data/mqcd.json is rebuilt from those records with mf-8g-3's A24 nightly record (workspace/tmp/mf-8g-3/a24/results/A24.nightly.json, catani native) as the later directory, and F11 of data/mqcd-findings.json is resolved (no new run).

- **Section 2, Table 5, Table 29:** data/dec.json (scripts/build_mqcd.py with --manifest tests/fixtures/loop/dec/manifest.json and the decay checks' JUnit, re-run at 617b59cd; dec-compare.json against MF-DEC-2's file: no status change), combined with the manifest by scripts/build_dec_matrix.py into data/dec-matrix.json; records in workspace/reports/engine-v2/m-final/records/m1101-mfb-new/ (results/, junit/, logs/, scripts/, STATE.md).
- **Table 6, Table 7:** data/hc.json and data/corr.json, written by scripts/build_hc.py and scripts/build_corr.py from the whole-file jobs of tests/unit/test_loop_v2_hc.py and test_loop_v2_corr_recola.py (their artifacts reloaded, JUnit) in the same records; the rescue and mapping counts of Section 1: that campaign's junit/rescue.xml and mf-map-fix's run at f3d9e361 (slice-status UPDATE 2026-11-02 ~01:05). HC05's row: the mf-hc05 record (workspace/tmp/mf-hc05/STATE.md, slice-status UPDATE 2026-11-03 ~21:25), entered by hand (record).
- **Table 1 (2 → 6 rows), Table 3:** data/2to6.json, written by scripts/build_2to6.py from the 106 mf-2to6 records it lists in sources (workspace/tmp/mf-2to6/runs/, workspace/artifacts/2to6/*-2f3f76d9/generation.json, processes.json; arguments in workspace/tmp/mf-2to6/build_args.sh); reproduced byte for byte; Table 1 through data/summary.json preliminary. **Table 4 (rounds at seven legs):** data/rounds-7leg.json, transcribed from workspace/reports/engine-v2/m-final/round-table-7leg.md (records workspace/tmp/mf-8g-table/rec/).
- **Section 1 (M-QCD), Section 5, Table 11, Table 31:** data/mqcd.json (scripts/build_mqcd.py from the campaign's records), data/mqcd-findings.json (MQ-VALIDATE's triage, re-checked at the pin), data/mqcd-public.json (the public-route representatives), data/mqcd-fft-cost.json and data/mqcd-fft-worst.json (FFT cost and worst points), data/mqcd-selftest.json (wheel self-test), combined by scripts/build_mqcd_matrix.py into data/mqcd-matrix.json; the campaign's records, logs and scripts in workspace/reports/engine-v2/m-qcd/records/m1029-mqcd/ (a copy of workspace/tmp/campaign-m1029-mqcd: c1/ results, JUnit and logs, x/ the FFT and stock-loop_sm test logs, public/, fft/, wheel/, STATE.md).
- **Table 1, Figure 1:** data/summary.json, written by scripts/build_summary.py from data/g1.json, the EW row of data/msm-g1.json, data/measured.json, data/milestone.json, tools/loop_benchmark/processes.json and the preliminary rows of data/2to6.json by its main-configuration rule (the most recent measured configuration of each cell; a newer check that only repeats part of an older comparison does not replace it; revision ancestry from data/milestone.json); every entry names its configuration and source (README).
- **Table 21:** HANDOFF §0.1 (addendum F); the provider generation records of Table 1 (data/summary.json): data/g1.json for the G1 rows and the reference runs' records (workspace/reference/runs/required/<process>/, collected into data/measured.json) for the $2 \to 5$ rows; MadLoop's production-build times of the $2 \to 5$ rows: the same runs' madloop-production-result.json. **Table 22, Section 7, Section 8:** the native validations on work/v2-generation-scale, records workspace/tmp/generation-scale/gint2/after/, gint/val1/ and i7/val/ (the earlier seven-gluon points: gint/val1/), collected into data/measured.json; the earlier $gg \to t\bar{t}e^-e^+g$ record: DESIGN WP3.1 fourth increment. **Conventions:** MadLoop standalone, Recola, OpenLoops 2.
- **Section 6 (Table 12 to Table 20):** data/g1.json and the per-point values data/g1-values.json, written by python -m tools.loop_benchmark gate (README and docstrings of tools/loop_benchmark); streams tools/loop_benchmark/streams/; run records archived in workspace/reports/engine-v2/measure-g1/records/ (first campaign runs/, fix round runs2/, the re-measure at G-int (4e9b5a9e) gint/ with its paired binary64-input runs confirm/<process>/ab.json, milestone m1025a's at 1d27426b m1025a/ (with the frozen-heap launcher gc_launcher.py and its protocol experiment gc/), milestone m1026a's at 11542fd5 m1026a/ (a copy of workspace/tmp/g1-m1026a/, with the cold/warm generation A/B ab/ and in killed/ the first full-colour seven-gluon run, no gate input), milestone m1028a's at 28a16255 m1028a/ (a copy of the campaign workspace/tmp/g1-m1028a/: runs/ with every row's timing, generation, accuracy check and stability rows, round 2 of the $2 \to 2$ to $2 \to 4$ rows (m1028a-r2-*, deciding) and the seven-gluon stores' records, logs/ with /usr/bin/time -v and host-load samples (hostload.txt, hostload-r2.txt), the public-path confirmation confirm/, the same-core old/new A/B and bisect runs ab/ with logs/ab-summary.txt, the gate inputs out/inputs-final.txt and labels out/meta-final.json, the campaign's scripts scripts/), milestone M-QCD's at e18a024b m1029-mqcd/ (a copy of workspace/tmp/g1-m1029-mqcd: runs/, logs/, bin/, the gate inputs inputs-final.txt, labels meta-final.json, compare.json, gate.log, STATE.md; the seven-gluon stores are not copied), milestone M-SM's at 71e2041e m1031-msm/ (a copy of workspace/tmp/g1-m1031-msm: the same files, the EW gate's msm-g1.json, msm-g1-values.json, msm-gate.log and the public-path check confirm/) and the final retime's (M-LIT-a0, lit-retime) at ca602436 mlita0/ (a copy of workspace/tmp/lit-retime: runs/, logs/, bin/, the gate inputs and labels, compare.json, the EW gate's files and confirm/; the seven-gluon stores are not copied); provider records reused from the first campaign and the fix round, not re-run); the earlier $2 \to 5$ records and the provider runs of those rows in workspace/reports/engine-v2/measure-g1-7leg/records/ (runs-f2/, runs-f4/), MadLoop exports of the G1-only processes in workspace/reference/runs/g1/; recalibration tools/loop_v2_costmodel.py --g1 data/g1.json. Builds: loop-native-v2-slice-ca602436 (loop, ca602436); history: loop-native-measure-g1-genscale-39824f9b (work/v2-generation-scale 889af02c plus the timers of scratch commit 39824f9b, tag g1-native-inputs-39824f9b: earlier runs, one lane, schedule M, stability) and loop-native-measure-g1-033edf15 (Python inputs); the earlier runs of M-QCD: loop-native-v2-slice-d0fef168, of m1028a: loop-native-v2-slice-5cb5ea83, of m1026a: loop-native-v2-slice-11542fd5, of m1025a: loop-native-v2-slice-7db32901, of G-int: loop-native-v2-slice-4e9b5a9e, release with one codegen unit and thin LTO; the harness is the committed tools/loop_benchmark of ca602436 for every row (the heap freeze in the harness; m1025a's launcher and measure-g1-7leg's run worktree 2265a9f7 are history). Schedule M at seven legs: DESIGN §5.6.
- **Table 23** and NJet prose: workspace/reports/engine-v2/njet-benchmark.md §2, §5, §6 (raw workspace/reference/runs/njet/timing/; the seven-gluon single evaluations are read from those records into data/measured.json); NJet, arXiv:1209.0100 Table 5.

- **Table 24:** `data/measured.json` (`scripts/collect_milestone.py <workspace>`); top-free values at lower multiplicity: DESIGN WP3.1 and `test_loop_v2_topfree.py`; the seven-gluon generation history: DESIGN WP3.1 eighth and ninth increments. Memory in GiB = 2^{30} bytes.
- **Table 25, Section 10:** `docs/development/loop_v2/DESIGN.md` §1–§3, §5.5, §5.6, §6, §10, §12, REQUIREMENTS_ADDENDUM J (counterterm rule) and the integration records above.
- **History:** the legacy per-diagram engine's snapshot of 2026-10-18 (`data/measurements.json`, `provider-cold.json`, `legacy-labels.json`, `colour-measurements.json`) stays for the DESIGN citations and is no longer shown; its table is in the earlier report revisions (README).

Appendix: the decay campaign rows

Table 29: **Rows of the decay campaign** (data/dec.json, written by scripts/build_mqcd.py with the decay manifest tests/fixtures/loop/dec/manifest.json from the campaign's records at 617b59cd, native loop-native-v2-slice-b36af289 through data/dec-matrix.json rows). Per row: gauge and the Catani route (python: the tools' operator; engine: the runtime's poles = "ir"), tiers (d default, n nightly) with points · checks · failures, binary64 against the MadLoop reference per coefficient (Born, finite, ε^{-1}, ε^{-2}; information: rule 2 asserts binary64 against DD), the coefficients whose binary64 estimate is out of scope (DD asserted there), the extras and the reference fixture.

Row	Process	Gauge · Catani	Tiers	Born	Finite	ε^{-1}	ε^{-2}	Out of scope	Extras	Reference
XX01	$t \to bw^+$	unitary · python	d 2 · 42 · 0; n 2 · 52 · 0	7.2e−16	4.5e−15	1.4e−14	5.8e−16	—	uv, ap	t_bwp_qcd.json
XX02	$t \to bw^+$	feynman · engine	n 4 · 134 · 0	9.0e−16	5.1e−15	2.0e−13	2.2e−15	—		t_bwp_ew.json
XX03	$z \to e^-e^+$	feynman · python	n 4 · 82 · 0	1.0e−15	3.6e−15	2.3e−15	6.1e−16	—		z_emep_ew.json
XX04	$z \to u\bar{u}$	feynman · python	n 4 · 134 · 0	1.1e−15	7.6e−15	3.9e−13	1.3e−15	—		z_uux_ew.json
XX05	$w^+ \to e^+$ ve	feynman · python	n 4 · 82 · 0	1.3e−15	1.9e−15	5.3e−14	1.3e−15	—		wp_epve_ew.json
XX06	$h \to b\bar{b}$	unitary · python	n 4 · 84 · 0	7.7e−16	2.9e−15	1.2e−15	—	4 × double pole (QCD2 QED2)	uv	h_bbx_4fs.json
XX07	$t \to bw^+g$	unitary · python	n 4 · 92 · 0	6.5e−15	1.4e−13	6.2e−14	5.9e−15	—	uv, ward	t_bwpg_qcd.json
XX08	$z \to u\bar{u}g$	feynman · python	n 4 · 134 · 0	1.2e−13	3.6e−12	5.2e−11	4.9e−14	—		z_uuxg_ew.json

Appendix: the M-SM campaign rows

Table 30: **Rows of the M-SM validation campaign** (data/msm.json, written by scripts/build_mqcd.py with its M-SM manifest from the campaign's records at 617b59cd, native loop-native-v2-slice-b36af289 through data/msm-matrix.json rows). Per row: class, process, model and restriction, tiers (d default, n nightly, c campaign) with their points, per-point inputs (n native), the reference (ML MadLoop in quadruple precision, "ML d" its double-precision controls, Rc Recola, OL OpenLoops), the largest deviations over the row and its gating variants (DD from the reference, rule 1; binary64 from DD, rule 2; poles from Catani in DD, rule 4, "≤" a relaxed tolerance; — not formed or no predicted sector), the extras and gating variants, the findings that name the row (data/msm-findings.json) and the status.

Row	Class	Process	Model	Points	Reference	DD	b64/DD	Catani	Extras · variants	Findings	Status
LL01	W-lep	$e^-e^+ \to m^-m^+$	G_μ	n 3	ML d	1.9e−14	5.0e−14	2.0e−30	ap, uv		pass
LL02	W-lep	$e^-e^+ \to m^-m^+$	G_μ, m_μ	n 3	ML d	1.1e−14	6.1e−15	—		F7	partial
LL03	W-lep	$e^-e^+ \to e^-e^+$	qcd_qed_sm	n 3	ML	4.0e−15	6.9e−15	4.5e−31			pass
LL05	W-lep	$e^-e^+ \to w^+w^-$	G_μ	n 4	ML	1.6e−13	3.7e−13	5.7e−29			pass
LL08	W-lep	$e^-e^+ \to m^-m^+a$	G_μ a0	n 4	ML	4.6e−13	2.3e−12	3.5e−28	a0		pass
LL09	W-lep	$e^-e^+ \to t\bar{t}$	G_μ	n 4	ML	2.8e−15	7.9e−15	5.1e−31	uv		pass
LL10	W-lep	$e^-e^+ \to t\bar{t}h$	G_μ	n 4	ML, Rc	2.0e−14	1.7e−12	4.9e−29			pass
DY01	W-DY	$u\bar{u} \to e^-e^+$	G_μ	n 4	ML	1.2e−15	1.9e−14	2.7e−31	uv, hel, sp · mwmz		pass
DY02	W-DY	$u\bar{d} \to e^+\ \mathrm{ve}$	G_μ	n 4	ML	6.9e−15	4.7e−12	1.8e−29			pass
DY03	W-DY	$u\bar{u} \to e^-e^+$	qcd_qed_sm	n 4	ML	1.9e−15	3.1e−14	5.0e−31			pass
DY04	W-DY	$e^-u \to e^-u$	G_μ	n 4	ML	2.3e−15	2.2e−14	8.0e−31			pass
DY05	W-DY	$u\bar{u} \to e^-e^+g$	G_μ	n 3	ML	9.1e−13	4.8e−13	5.8e−30	ward, joint, ap	F4	pass
DY07	W-DY	$u\bar{u} \to e^-e^+m^-m^+$	G_μ	n 4	ML	8.1e−15	4.2e−12	1.1e−28			pass
DY08	W-DY	$u\bar{u} \to e^-e^+gg$	G_μ	n 4	ML	2.5e−14	1.5e−12	1.8e−29			pass
VJ01	W-V	$u\bar{u} \to zg$	G_μ	n 4	ML	3.6e−15	1.2e−13	3.3e−29	ward		pass
VJ02	W-V	$u\bar{d} \to w^+g$	G_μ	d 2 · n 2	ML	4.0e−15	7.8e−15	2.1e−30	uv		pass
VJ03	W-V	$u\bar{u} \to ag$	G_μ a0	n 3	ML	2.8e−14	9.5e−13	3.3e−30	a0, ward	F4	pass
VV01	W-V	$u\bar{u} \to w^+w^-$	G_μ	n 4	ML	6.7e−13	5.3e−13	6.3e−28	sp · 4tev		pass
VV02	W-V	$u\bar{u} \to zz$	G_μ	n 4	ML	2.0e−14	2.0e−12	5.8e−30	hel	F1	pass
VV03	W-V	$u\bar{d} \to w^+z$	G_μ	n 4	ML	2.4e−13	5.9e−12	4.5e−29			pass
VV05	W-V	$u\bar{d} \to w^+a$	G_μ a0	n 4	ML	1.9e−13	2.5e−11	2.3e−28	a0, ward, sp · zero	F3	pass
VV06	W-V	$u\bar{u} \to aa$	G_μ a0	n 3	ML	8.8e−13	1.3e−12	2.9e−30	a0	F4	pass
VV07	W-V	$u\bar{u} \to w^+w^-g$	G_μ	n 4	ML	5.3e−11	9.2e−11	1.1e−26	ward		pass
V301	W-V	$u\bar{u} \to w^+w^-z$	G_μ	n 4	ML	5.9e−14	4.8e−14	2.6e−28		F1	pass
VH01	W-V	$u\bar{u} \to zh$	G_μ	n 4	ML	9.7e−15	2.0e−13	1.6e−30		F1	pass
VH02	W-V	$u\bar{d} \to w^+h$	G_μ	n 4	ML	1.7e−14	1.8e−12	5.8e−29		F1	pass
VH03	W-V	$ud \to udh$	G_μ	n 4	ML	4.7e−13	8.4e−12	2.3e−28		F1	pass
TT01	W-top	$gg \to t\bar{t}$	G_μ	d 2 · n 1	ML, Rc, OL	5.9e−15	8.2e−15	6.6e−31	joint, ap, uv		pass
TT02	W-top	$u\bar{u} \to t\bar{t}$	G_μ	n 4	ML	2.7e−14	5.2e−14	1.0e−29	so, uv		pass
TT03	W-top	$gg \to t\bar{t}g$	G_μ	n 4	ML	1.4e−11	3.4e−11	2.1e−26	ward	F3	pass
TT04	W-top	$t\bar{t} \to t\bar{t}$	G_μ	n 4 · n	ML, Rc	1.4e−14	2.8e−12	1.8e−30			pass
TT05	W-top	$gt \to tt\bar{t}$	G_μ	n 4	ML, Rc	3.9e−12	1.0e−11	1.1e−29 ≤ 1×10⁻¹⁴	uv		pass
TT06	W-top	$gb \to t\bar{t}b$	G_μ with_b_mass	n 4	ML, Rc	1.2e−11	4.2e−16	4.2e−30 ≤ 1×10⁻¹³		F10	pass
TV01	W-top	$gg \to t\bar{t}h$	G_μ	n 4	ML	3.6e−13	7.3e−12	7.4e−29			pass
TV02	W-top	$gg \to t\bar{t}z$	G_μ	n 4	ML	2.1e−13	3.2e−12	6.6e−29			pass
TV03	W-top	$u\bar{d} \to t\bar{t}w^+$	G_μ	n 4	ML	3.8e−14	1.5e−11	8.7e−29	so	F1	pass
TV06	W-top	$gg \to t\bar{t}e^-e^+$	G_μ	n 3	ML, Rc, OL	1.2e−12	1.9e−10	1.8e−28			pass
TV07	W-top	$gg \to t\bar{t}e^-e^+g$	G_μ	c 3 · n	ML, Rc, OL	1.4e−12	5.1e−11	6.4e−26	ward, joint · qcd	F3	pass
ST01	W-singletop	$ub \to dt$	G_μ	n 4	ML	3.4e−14	1.9e−14	7.4e−31	uv		pass
ST02	W-singletop	$u\bar{d} \to t\bar{b}$	G_μ	n 3	ML, Rc	4.2e−13	4.7e−14	1.9e−30			pass
ST03	W-singletop	$gb \to tw^-$	G_μ	n 4	ML	1.7e−15	2.0e−14	1.5e−29	ward	F2, F3	pass
ST04	W-singletop	$u\bar{d} \to t\bar{b}$	G_μ with_b_mass	n 4	ML, Rc	3.0e−12	2.3e−11	4.8e−27		F10	pass
JJ01	W-jets	$u\bar{u} \to d\bar{d}$	G_μ	n 4	ML	6.7e−15	8.4e−14	9.6e−31	so, uv		pass
JJ02	W-jets	$u\bar{u} \to u\bar{u}$	G_μ	n 4	ML	2.8e−14	6.4e−13	7.8e−29			pass
JJ03	W-jets	$u\bar{u} \to gg$	G_μ	n 4	ML	6.7e−15	6.5e−13	1.8e−30			pass
JJ04	W-jets	$u\bar{u} \to d\bar{d}g$	G_μ	n 4	ML	6.9e−13	7.6e−12	2.0e−28	so		pass
PI01	W-photon	$aa \to e^-e^+$	G_μ	n 4	ML	9.3e−15	1.7e−13	1.9e−30	a0 · all		pass
PI02	W-photon	$aa \to w^+w^-$	G_μ	n 4	ML	3.2e−15	4.9e−14	3.3e−30			pass
PI03	W-photon	$ag \to u\bar{u}$	G_μ	n 4	ML	1.9e−15	2.9e−13	1.7e−31	uv		pass
PI04	W-photon	$au \to au$	G_μ a0	n 4	ML	9.7e−15	2.7e−13	1.0e−29	a0		pass
CM01	W-CMS	$e^-e^+ \to m^-m^+$	G_μ	n 4	ML d	1.2e−13	1.0e−13	2.0e−30	ap, sp · zero, small, source		pass
CM02	W-CMS	$u\bar{u} \to e^-e^+$	G_μ	n 4	ML	8.8e−15	1.5e−13	2.2e−29	sp		pass
CM03	W-CMS	$u\bar{d} \to e^+\ \mathrm{ve}$	G_μ	n 4	ML	2.2e−14	3.9e−12	3.1e−30	sp		pass
CM04	W-CMS	$gg \to t\bar{t}$	G_μ	n 4	ML	1.3e−15	1.5e−14	9.1e−31	uv	F8	pass
CM05	W-CMS	$u\bar{u} \to e^-e^+g$	G_μ	n 4	ML	1.1e−13	5.1e−13	4.8e−30	ward		pass
CM07	W-CMS	$gg \to t\bar{t}e^-e^+g$	G_μ	c 3 · n	ML	6.9e−13	5.6e−11	7.2e−26	ward	F3, F8	pass
CM08	W-CMS	$u\bar{d} \to e^+\ \mathrm{ve}\ b\bar{b}$	G_μ	n 4	ML	1.3e−14	1.8e−13	—			pass
CM09	W-CMS	$u\bar{d} \to e^+\ \mathrm{ve}\ a$	G_μ a0	n 4	ML	2.7e−14	1.6e−12	2.4e−29	a0, ward		pass
CM10	W-CMS	$u\bar{d} \to w^+b\bar{b}$	loop_sm no_b_mass	n 4	ML	1.6e−13	1.8e−10	—			pass
MP01	W-multi	$pp \to e^-e^+$	G_μ	variants	ML	4.2e−14	1.5e−12	2.9e−30	P2, P3, P4, P5	F2	pass
MP02	W-multi	$pp \to t\bar{t}$	G_μ	variants	ML	2.7e−14	1.6e−13	1.0e−29	P0, P2, P3, P4, P5, P6	F2	pass

Row	Class	Process	Model	Points	Reference	DD	b64/DD	Catani	Extras · variants	Findings	Status
SC01	W-scheme	$e^-e^+ \to u\bar{u}$	G_μ	n 4	ML d	6.8e−15	7.9e−15	1.9e−31			pass
SC02	W-scheme	$u\bar{u} \to m^-m^+$	G_μ, m_μ	n 4	ML	4.9e−9	1.6e−14	—		F5, F7	partial
SC03	W-scheme	$u\bar{d} \to e^+\ \mathrm{ve}$	G_μ ckm	n 4	ML	3.0e−12	5.4e−12	1.1e−6 ≤ 1×10^{-5}	uv	F6	pass
SC04	W-scheme	$b\bar{b} \to w^+w^-$	G_μ with_b_mass	n 4	ML	2.3e−10	3.1e−14	4.8e−29		F1, F9, F10	pass
SC05	W-scheme	$u\bar{u} \to e^-e^+gg$	loop_sm	n 3	fixture	8.0e−13	2.8e−12	4.1e−14 ≤ 1×10^{-12}			pass
SC06	W-scheme	$w^+w^- \to w^+w^-$	G_μ	n 4 · n	ML	6.3e−10	1.1e−10	1.1e−26		F1	pass
SC07	W-scheme	$hh \to hh$	G_μ	n 4 · n	ML	8.2e−16	2.6e−12	1.6e−32		F1	pass
SC08	W-scheme	$t\bar{t} \to w^+w^-$	G_μ	n 4 · n	ML	6.3e−13	1.2e−12	5.0e−29		F2	pass
SC09	W-scheme	$t\bar{t} \to zh$	G_μ	n 4 · n	ML, Rc	3.0e−15	2.0e−11	1.3e−30			pass
SC10	W-scheme	$zz \to hh$	G_μ	n 4	ML, Rc	4.3e−13	7.0e−13	3.7e−32			pass

Appendix: the M-QCD campaign rows

Table 31: **Rows of the M-QCD validation campaign** (data/mqcd.json, written by scripts/build_mqcd.py from the campaign's records; 617b59cd (M-FINAL-b pin, native loop-native-v2-slice-b36af289). Per row: class, process and loop_sm restriction, tiers (d default, n nightly) with their points, the reference providers of the row and its gating variants (ML MadLoop, Rc Recola, OL OpenLoops, NJet; MadLoop in quadruple precision from MQ-REF), the largest deviations over the row and its gating variants (DD from the reference, rule 1; binary64 from DD where asserted, rule 2; poles from Catani in DD, rule 4; — none formed), the extras and gating variants (status where not pass), and the row's status; the failures and their findings are in data/mqcd.json and data/mqcd-findings.json.

Row	Class	Process	Restriction	Points	Reference	DD	b64/DD	Catani	Extras · variants	Status
A01	Q-massless	$gg \to gg$	no_b_mass	d 2 · n 1	ML, Rc, OL	1.1e−15	2.9e−14	1.2e−31	nc, ap · sectors	pass
A02	Q-massless	$u\bar{u} \to gg$	no_b_mass	n 3	ML	1.2e−14	4.2e−13	1.4e−31	madloop	pass
A03	Q-massless	$gg \to u\bar{u}$	no_b_mass	n 4	ML	6.6e−15	4.2e−14	1.6e−31		pass
A04	Q-massless	$ug \to ug$	no_b_mass	d 2 · n 2	ML	1.0e−15	7.3e−15	2.2e−30		pass
A05	Q-massless	$u\bar{u} \to d\bar{d}$	no_b_mass	n 4	ML	1.1e−15	6.8e−15	1.9e−31	uv	pass
A06	Q-massless	$u\bar{u} \to u\bar{u}$	no_b_mass	n 3	ML, Rc, OL	2.8e−14	5.8e−13	1.1e−28	sectors	pass
A07	Q-massless	$uu \to uu$	no_b_mass	n 4	ML	2.6e−15	2.9e−14	1.6e−30		pass
A08	Q-massless	$gg \to ggg$	no_b_mass	n 4	ML	3.4e−14	4.5e−13	6.1e−30	nc, ward, hel	pass
A09	Q-massless	$u\bar{u} \to ggg$	no_b_mass	d 2 · n 2	ML, Rc, OL	7.8e−13	4.5e−11	1.2e−22	nc, sp · soft, collinear fail	fail
A12	Q-massless	$u\bar{u} \to d\bar{d}g$	no_b_mass	d 2 · n 1	ML, Rc, OL	5.0e−14	4.5e−13	7.7e−30	nc · sectors	pass
A13	Q-massless	$u\bar{u} \to u\bar{u}g$	no_b_mass	n 3	ML, Rc, OL	6.3e−14	7.9e−13	5.3e−30	nc, ap · sectors	pass
A14	Q-massless	$gg \to gggg$	no_b_mass	n 3	ML, Rc, OL	3.3e−14	4.6e−11	1.4e−23	ward	pass
A15	Q-massless	$u\bar{u} \to gggg$	no_b_mass	n 4	ML	3.4e−14	4.9e−12	1.7e−29	nc	pass
A18	Q-massless	$gg \to u\bar{u}u\bar{u}$	no_b_mass	n 4	ML, Rc, OL	1.2e−11	1.1e−10	1.4e−27	nc	pass
A19	Q-massless	$u\bar{u} \to d\bar{d}s\bar{s}$	no_b_mass	n 3	ML, Rc, OL	8.4e−16	2.0e−13	8.1e−30	sectors	pass
A20	Q-massless	$u\bar{u} \to u\bar{u}d\bar{d}$	no_b_mass	n 3	ML, Rc, OL	1.9e−14	2.4e−12	2.5e−30	sectors	pass
A21	Q-massless	$u\bar{u} \to u\bar{u}u\bar{u}$	no_b_mass	n 3	ML, Rc, OL	3.7e−14	4.1e−12	2.3e−29	nc · sectors	pass
A22	Q-massless	$u\bar{u} \to d\bar{d}s\bar{s}g$	no_b_mass	n 3	ML, Rc	7.0e−13	4.8e−13	1.3e−28		pass
A23	Q-massless	$u\bar{u} \to ggggg$	no_b_mass	n 3	ML, Rc	3.2e−12	1.3e−12	2.7e−29		pass
A24	Q-massless	$gg \to ggggg$	no_b_mass	n 2	NJet	8.0e−17	3.2e−12	3.5e−26	uv	pass
A25	Q-massless	$ud \to ud$	no_b_mass	n 4	ML	2.1e−15	5.8e−14	9.2e−30		pass
A26	Q-massless	$\bar{u}\bar{u} \to \bar{u}\bar{u}$	no_b_mass	n 4	ML	2.6e−15	8.4e−14	1.6e−30		pass
B01	Q-top	$gg \to t\bar{t}$	no_b_mass	d 2 · n 1	ML, Rc, OL	1.7e−15	1.2e−10	3.0e−26	ap, uv, hel, sp · threshold, mt170, sectors	pass
B02	Q-top	$u\bar{u} \to t\bar{t}$	no_b_mass	n 3	ML	3.9e−14	8.7e−14	4.0e−30	madloop	pass
B03	Q-top	$gg \to t\bar{t}g$	no_b_mass	n 4	ML	3.4e−13	3.4e−11	7.5e−27	nc, ward	pass
B05	Q-top	$gg \to t\bar{t}gg$	no_b_mass	n 3	ML, Rc, OL	1.3e−13	3.2e−11	1.4e−27	sectors	pass
B06	Q-top	$u\bar{u} \to t\bar{t}d\bar{d}$	no_b_mass	n 4	ML	1.5e−14	1.5e−12	1.8e−28		pass
B07	Q-top	$gg \to t\bar{t}t\bar{t}$	no_b_mass	n 4	ML	2.5e−14	1.4e−13	1.4e−29		pass
B08	Q-heavy	$gg \to t\bar{t}$	default	n 4	ML	1.9e−15	3.1e−14	1.5e−30		pass
C01	Q-heavy	$gg \to b\bar{b}$	default	d 2 · n 2	ML, Rc	9.3e−16	6.2e−15	8.1e−29	nc, ap, ward, sp · params	pass
C02	Q-heavy	$gg \to gg$	default	n 4	ML	5.0e−14	9.0e−12	1.7e−29	uv	pass
C03	Q-heavy	$b\bar{b} \to h$	default	n 4	ML	3.9e−16	4.3e−16	9.9e−31		pass
C04	Q-heavy	$b\bar{b} \to hg$	default	n 4	ML	5.9e−16	1.9e−15	3.2e−30		pass
C05	Q-heavy	$gg \to b\bar{b}h$	default	n 4	ML	5.0e−15	2.3e−10	1.6e−29		pass
C06	Q-heavy	$gg \to t\bar{t}b\bar{b}$	default	n 4	ML	1.9e−15	1.5e−13	1.9e−29		pass
C07	Q-singletop	$ug \to dt\bar{b}$	default	n 4	ML	8.0e−15	1.6e−14	—		pass
C08	Q-heavy	$gg \to c\bar{c}$	c_mass	n 4	ML, Rc	1.7e−15	4.5e−15	4.9e−29		pass
C09	Q-heavy	$gs \to w^- c$	c_mass	n 4	ML, Rc	4.6e−16	1.5e−14	1.1e−28		pass
D01	Q-EWBorn	$u\bar{u} \to e^- e^+$	no_b_mass	d 2 · n 2	ML	5.3e−16	4.9e−15	1.0e−31	ap, ff	pass
D02	Q-EWBorn	$u\bar{u} \to e^- e^+$	no_b_mass	n 4	ML	7.1e−16	4.3e−15	2.1e−31		pass
D05	Q-EWBorn	$u\bar{u} \to e^- e^+ g$	no_b_mass	d 2 · n 2	ML	2.3e−14	2.5e−14	1.6e−31	nc, ward, uv, hel	pass
D06	Q-EWBorn	$u\bar{u} \to e^- e^+ gg$	no_b_mass	n 3	ML, Rc	8.6e−14	1.1e−12	2.7e−31	madloop	pass
D07	Q-EWBorn	$u\bar{u} \to e^- e^+ ggg$	no_b_mass	n 3	ML, Rc, OL	1.3e−14	5.2e−13	2.4e−30	sectors	pass
D08	Q-EWBorn	$u\bar{u} \to e^- e^+ d\bar{d}g$	no_b_mass	n 3	ML, Rc, OL	1.1e−14	8.9e−13	1.4e−30	sectors	pass
D09	Q-EWBorn	$u\bar{u} \to e^- e^+ m^- m^+$	no_b_mass	n 3	ML	8.9e−16	2.7e−14	5.3e−31	sectors	pass
D10	Q-EWBorn	$u\bar{u} \to u\bar{u}e^+ e^-$	no_b_mass	n 3	ML, Rc	3.1e−15	1.7e−12	3.0e−30	sectors	pass
D11	Q-EWBorn	$u\bar{u} \to zg$	no_b_mass	n 4	ML	7.1e−15	3.4e−14	2.5e−30		pass
D12	Q-EWBorn	$u\bar{d} \to e^+ \text{ ve } g$	no_b_mass	n 4	ML	1.3e−15	1.2e−14	1.4e−31		pass
D13	Q-EWBorn	$ug \to au$	no_b_mass	n 4	ML	1.2e−15	1.2e−14	4.9e−31		pass
D15	Q-EWBorn	$u\bar{u} \to zz$	no_b_mass	n 4	ML	3.2e−15	2.8e−13	1.5e−30		pass
D16	Q-EWBorn	$u\bar{u} \to w^+ w^-$	no_b_mass	n 4	ML	1.2e−13	1.4e−12	3.2e−28		pass
D17	Q-EWBorn	$b\bar{b} \to w^+ w^-$	default	n 4	ML	4.6e−14	6.6e−13	6.1e−29	5F	pass
D18	Q-EWBorn	$u\bar{u} \to zh$	no_b_mass	n 4	ML	8.5e−16	2.0e−12	1.1e−29		pass
D19	Q-EWBorn	$ud \to udh$	no_b_mass	n 4	ML	5.9e−16	3.3e−15	—		pass
D20	Q-EWBorn	$e^+ e^- \to u\bar{u}g$	no_b_mass	n 4	ML	7.3e−14	2.0e−12	7.0e−30		pass
D21	Q-EWBorn	$e^+ e^- \to t\bar{t}$	no_b_mass	n 3	ML	1.3e−14	1.8e−14	1.2e−31		pass
D23	Q-EWBorn	$ag \to u\bar{u}$	no_b_mass	n 4	ML	4.3e−16	3.1e−13	3.2e−31		pass
D24	Q-EWBorn	$e^- u \to e^- u$	no_b_mass	n 4	ML	4.8e−16	4.1e−15	8.0e−32		pass
D25	Q-EWBorn	$u\bar{d} \to w^+ a$	no_b_mass	n 4	ML	4.8e−15	4.4e−13	1.3e−28	sp · zero	pass
D26	Q-heavy	$u\bar{u} \to e^- e^+ g$	default	n 4	ML	2.3e−14	8.5e−14	1.5e−31		pass
D27	Q-EWBorn	$gu \to e^+ \text{ ve } d$	no_b_mass	n 4	ML	1.3e−15	4.3e−14	6.5e−31		pass
D28	Q-EWBorn	$u\bar{u} \to w^+ w^- g$	default	n 4	ML	1.8e−12	1.4e−11	1.6e−28	ward	pass

Row	Class	Process	Restriction	Points	Reference	DD	b64/DD	Catani	Extras · variants	Status
D29	Q-EWBorn	$b\bar{b} \to w^+w^-g$	default	n 4	ML	1.7e−13	8.8e−12	7.3e−28	ward · 5F	pass
E01	Q-ttX	$gg \to t\bar{t}h$	no_b_mass	n 4	ML	3.6e−13	7.4e−12	6.7e−30	nc, ward, uv	pass
E02	Q-ttX	$gg \to t\bar{t}z$	no_b_mass	n 4	ML	7.2e−14	3.2e−12	5.8e−30		pass
E05	Q-ttX	$e^+e^- \to t\bar{t}h$	no_b_mass	n 3	ML	5.8e−15	6.6e−15	2.4e−31	ap	pass
E06	Q-ttX	$gg \to t\bar{t}e^-e^+g$	no_b_mass	n 3	ML, Rc, OL	6.7e−14	4.4e−12	5.7e−27	sectors	pass
F01	Q-singletop	$ub \to dt$	default	n 4	ML	1.2e−15	8.7e−16	1.3e−31	nc, uv, ap · 5F	pass
F02	Q-singletop	$bu \to dt$	default	n 4	ML	1.2e−15	1.8e−15	1.1e−31	5F	pass
F03	Q-singletop	$u\bar{d} \to t\bar{b}$	no_b_mass	d 2 · n 2	ML	1.7e−15	4.1e−16	8.6e−32		pass
F04	Q-singletop	$gb \to tw^-$	default	n 4	ML	6.2e−16	3.2e−15	3.6e−30	ward · 5F	pass
G01	Q-mixed	$u\bar{u} \to d\bar{d}$	none	d 2 · n 1	ML	4.6e−15	2.4e−15	5.0e−32	uv, ap	pass
G02	Q-mixed	$u\bar{u} \to t\bar{t}$	no_b_mass	n 4	ML	2.0e−15	3.1e−14	3.7e−30		pass
H01	Q-CKM	$u\bar{s} \to w^+g$	zeromass_ckm	n 4	ML	1.4e−15	7.2e−14	4.7e−29	uv, ap	pass
H02	Q-CKM	$gs \to tw^-$	zeromass_ckm	n 4	ML	1.2e−15	7.6e−15	3.2e−30		pass
H03	Q-CKM	$gd \to tw^-$	ckm	n 4	ML	1.2e−15	5.6e−15	4.6e−30		pass
M01	Q-multi	$pp \to t\bar{t}$	no_b_mass		ML	2.0e−15	4.1e−14	4.0e−30	P0, public-P0, P2, public-P2	pass
R01	Q-EWBorn	$u\bar{u} \to m^-m^+$	lepton_masses	n 4	ML	3.1e−16	1.2e−15	3.9e−32		pass